

CS 341: Algorithms

Douglas R. Stinson

David R. Cheriton School of Computer Science
University of Waterloo

Fall, 2015

- 1 Course Information
- 2 Introduction
- 3 Divide-and-Conquer Algorithms
- 4 Greedy Algorithms
- 5 Dynamic Programming Algorithms
- 6 Graph Algorithms
- 7 Intractability and Undecidability

Table of Contents

1 Course Information

Course mechanics

- Sections and Instructors:
 - ▶ Section 1, Doug Stinson, M W 11:30–12:50, MC 1056
 - ▶ Section 2, Doug Stinson, M W 1:00–2:20, MC 1056
 - ▶ Section 3, Kevin Lanctot, M W 10:00–11:20, RCH 305
- Scheduled Office Hours:
 - ▶ Doug Stinson, T 1:30–2:30
 - ▶ Kevin Lanctot, W 12:00–1:00

Course mechanics

- **Come to class!** Not all the material will be on the slides or in the text.
- You will need an account in the student.cs environment
- The course website can be found at
`https://www.student.cs.uwaterloo.ca/~cs341/`
 - ▶ Syllabus, calendar, policies, etc. can be found there (or on Learn).

Learn and Piazza

- Slides, assignments, solutions, grades, etc. will be available on **Learn**.
- However, discussion related to the course will take place on **Piazza** (piazza.com).
 - ▶ General course questions, announcements
 - ▶ Assignment-related questions
 - ▶ You will be getting an invitation via email to join Piazza in the first week of classes.
- Keep up with the information posted on the course website, Learn and Piazza

Courtesy

- Please silence cell phones and other mobile devices before coming to class.
- Questions are encouraged, but **please refrain from talking** in class – it is distracting to your classmates who are trying to listen to the lectures and to your professor who is trying to think, talk and write at the same time.
- Carefully consider whether **using your laptop in class** will help you learn the material and follow the lectures.
- Do not play games, tweet, watch youtube videos, update your facebook page or use your laptop in any other way that will **distract your classmates**.

Course syllabus

- You are expected to be familiar with the contents of the course syllabus
- Available on the course home page (or on Learn)
- If you haven't read it, read it after this lecture

Plagiarism and academic offenses

- We take academic offences very seriously
- There is a good discussion of plagiarism online:
 - ▶ http://www.math.uwaterloo.ca/navigation/Current/cheating_policy.shtml
- Read this and understand it
 - ▶ Ignorance is no excuse!
 - ▶ Questions should be brought to instructor
- Plagiarism applies to both text and code
- You are free (even encouraged) to exchange ideas, but **no sharing code or text**

Plagiarism (2)

- Common mistakes
 - ▶ Excess collaboration with other students
 - ★ Share ideas, but no design or code!
 - ▶ Using solutions from other sources (like for previous offerings of this course, maybe written by yourself)
- Possible penalties
 - ▶ First offense (for assignments; exams are harsher)
 - ★ 0% for that assignment, -5% on final grade
 - ▶ Second offense
 - ★ Expulsion is possible
- More information linked to from course syllabus

Grading scheme for CS 341

- Midterm (25%)
 - ▶ Friday, Oct. 23, 2015, 6:30–8:00 PM
- Assignments (25%)
 - ▶ There will be five assignments.
 - ▶ Work alone
 - ▶ See syllabus for reappraisal policies, academic integrity policy, and other details
- Final (50%)
- For medical conditions, you need to submit a Verification of Illness form.

Assignments

- All three sections will have the same assignments, midterm and final exam.
- Assignments will be due **at 4:00 PM** on the due date
- Late submissions will be accepted up to **24 hours after due date**
- There will be a **penalty of 25%** for accepted late submissions
- Multiple assignments can be submitted late, including the last one
- **No assistance will be given after the due date**
- You need to notify your instructor **before the due date** of a severe, long-lasting problem that prevents you from doing an assignment
- The same TA will mark the same question across all sections.

Assignment due dates

- Assignment 1: due Thursday October 1
- Assignment 2: due Thursday October 15
- Assignment 3: due Thursday November 5
- Assignment 4: due Thursday November 19
- Assignment 5: due Thursday December 3

Required textbook

- **Introduction to Algorithms, Third Edition**, by Cormen, Leiserson, Rivest and Stein, MIT Press, 2009.
- You are expected to know
 - ▶ entire textbook sections, as listed on course website or Learn
 - ▶ all the material presented in class

Table of Contents

2 Introduction

- Algorithm Design and Analysis
- The Maximum Problem
- The Max-Min Problem
- The 3SUM Problem
- Definitions and Terminology
- Order Terminology
- Formulae
- Algorithm Analysis Techniques

Analysis of algorithms

In this course, we study the **design** and **analysis** of algorithms. “Analysis” refers to mathematical techniques for establishing both the **correctness** and **efficiency** of algorithms.

Correctness: We often want a formal proof of correctness of an algorithm we design. This might be accomplished through the use of **loop invariants** and mathematical induction.

Analysis of algorithms (cont.)

Efficiency: Given an algorithm A , we want to know how efficient it is. This includes several possible criteria:

- What is the **asymptotic complexity** of algorithm A ?
- What is the **exact number** of specified computations done by A ?
- How does the **average-case** complexity of A compare to the **worst-case** complexity?
- Is A the most efficient algorithm to solve the given problem? (For example, can we find a **lower bound** on the complexity of **any** algorithm to solve the given problem?)
- Are there problems that cannot be solved efficiently? This topic is addressed in the theory of **NP-completeness**.
- Are there problems that cannot be solved by **any** algorithm? Such problems are termed **undecidable**.

Design of algorithms

“Design” refers to **general strategies** for creating new algorithms. If we have good design strategies, then it will be easier to end up with correct and efficient algorithms. Also, we want to avoid using **ad hoc** algorithms that are hard to analyze and understand.

Here are some useful design strategies, many of which we will study:

divide-and conquer

greedy

dynamic programming

depth-first and breadth-first search

local search (not studied in this course)

linear programming (not studied in this course)

The “Maximum” problem

Problem

Maximum

Instance: *an array A of n integers,*

$$A = [A[1], \dots, A[n]].$$

Find: *the maximum element in A .*

The **Maximum** problem has an obvious simple solution.

Algorithm: *FindMaximum*($A = [A[1], \dots, A[n]]$)

$max \leftarrow A[1]$

for $i \leftarrow 2$ **to** n

do $\begin{cases} \text{if } A[i] > max \\ \text{then } max \leftarrow A[i] \end{cases}$

return (max)

Correctness of FindMaximum

How can we formally prove that *FindMaximum* is correct?

Claim: At the end of iteration i ($i = 2, \dots, n$), the current value of max is the maximum element in $[A[1], \dots, A[i]]$.

The claim can be proven by induction. The base case, when $i = 2$, is obvious.

Now we make an induction assumption that the claim is true for $i = j$, where $2 \leq j \leq n - 1$, and we prove that the claim is true for $i = j + 1$ (fill in the details!).

When $j = n$ we are done and the correctness of *FindMaximum* is proven.

Analysis of FindMaximum

It is obvious that the complexity of *FindMaximum* is $\Theta(n)$.

More precisely, we can observe that the number of comparisons of array elements done by *FindMaximum* is **exactly** $n - 1$.

It turns out that *FindMaximum* is **optimal** with respect to the number of comparisons of array elements.

That is, any algorithm that correctly solves the **Maximum** problem for an array of n elements requires **at least** $n - 1$ comparisons of array elements.

How can we prove this assertion?

The “Max-Min” problem

Problem

Max-Min

Instance: an array A of n integers, $A = [A[1], \dots, A[n]]$.

Find: the maximum and the minimum element in A .

The **Max-Min** problem also has an obvious simple solution.

Algorithm: *FindMaximumAndMinimum*($A = [A[1], \dots, A[n]]$)

$max \leftarrow A[1]$

$min \leftarrow A[1]$

for $i \leftarrow 2$ **to** n

do $\left\{ \begin{array}{l} \text{if } A[i] > max \\ \quad \text{then } max \leftarrow A[i] \\ \text{if } A[i] < min \\ \quad \text{then } min \leftarrow A[i] \end{array} \right.$

return (max, min)

Analysis of *FindMaximumAndMinimum*

Exercise: Give a formal proof by induction that *FindMaximumAndMinimum* is correct.

The complexity of *FindMaximumAndMinimum* is $\Theta(n)$

More precisely, *FindMaximumAndMinimum* requires $2n - 2$ comparisons of array elements given an array of size n .

The **complexity** is optimal (why?), but there are algorithms to solve the **Max-Min** problem which require fewer comparisons of array elements than *FindMaximumAndMinimum*.

Note: An algorithm requiring fewer comparisons of array elements **may or may not be faster** than *FindMaximumAndMinimum*.

A more significant improvement

With some ingenuity, we can actually reduce the number of comparisons of array elements by (roughly) 25%.

Suppose n is even and we consider the elements **two at a time**. Initially, we compare the first two elements and initialize maximum and minimum values. (**One comparison** is required here.)

Then, each time we compare a new pair of elements, we subsequently compare the larger of the two elements to the current maximum and the smaller of the two to the current minimum. (**Three comparisons** are done here to process two array elements.)

This yields an algorithm requiring a total of $3n/2 - 2$ comparisons.

An improved algorithm

Algorithm: *ImprovedFindMaximumAndMinimum*(A)

comment: assume n is even

```

if  $A[1] > A[2]$  then  $\begin{cases} \max \leftarrow A[1] \\ \min \leftarrow A[2] \end{cases}$ 

    else  $\begin{cases} \max \leftarrow A[2] \\ \min \leftarrow A[1] \end{cases}$ 

for  $i \leftarrow 2$  to  $n/2$ 
    do  $\begin{cases} \text{if } A[2i-1] > A[2i] \\ \text{then } \begin{cases} \text{if } A[2i-1] > \max \text{ then } \max \leftarrow A[2i-1] \\ \text{if } A[2i] < \min \text{ then } \min \leftarrow A[2i] \end{cases} \\ \text{else } \begin{cases} \text{if } A[2i] > \max \text{ then } \max \leftarrow A[2i] \\ \text{if } A[2i-1] < \min \text{ then } \min \leftarrow A[2i-1] \end{cases} \end{cases}$ 

return  $(\max, \min)$ 
  
```

Optimality of the previous algorithm

It is possible to **prove** that any algorithm that solves the **Max-Min** problem requires at least $3n/2 - 2$ comparisons of array elements in the worst case.

Therefore the algorithm *ImprovedFindMaximumAndMinimum* is in fact **optimal** with respect to the number of comparisons of array elements required.

The “3SUM” problem

Problem

3SUM

Instance: *an array A of n distinct integers, $A = [A[1], \dots, A[n]]$.*

Question: *do there exist three elements in A that sum to 0?*

The **3SUM** problem also has an obvious algorithm to solve it.

Algorithm: *Trivial3SUM*($A = [A[1], \dots, A[n]]$)

for $i \leftarrow 1$ **to** $n - 2$

do $\left\{ \begin{array}{l} \text{for } j \leftarrow i + 1 \text{ to } n - 1 \\ \text{do } \left\{ \begin{array}{l} \text{for } k \leftarrow j + 1 \text{ to } n \\ \text{do } \left\{ \begin{array}{l} \text{if } A[i] + A[j] + A[k] = 0 \\ \text{then output } (i, j, k) \end{array} \right. \end{array} \right. \end{array} \right.$

The complexity of *Trivial3SUM* is $O(n^3)$.

A possible improvement

Instead of having three nested loops, suppose we have **two** nested loops (with indices i and j , say) and then we **search** for an $A[k]$ for which $A[i] + A[j] + A[k] = 0$.

If we try all possible k -values, then we basically have the previous algorithm.

What can we do to make the search more efficient?

What effect does this have on the complexity of the resulting algorithm?

An improved algorithm for the “3SUM” problem

Algorithm: *Improved3SUM*($A = [A[1], \dots, A[n]]$)

sort A in increasing order

for $i \leftarrow 1$ **to** $n - 2$

do $\left\{ \begin{array}{l} \textbf{for } j \leftarrow i + 1 \textbf{ to } n - 1 \\ \textbf{do } \left\{ \begin{array}{l} \text{perform a binary search for the value } A[k] = -A[i] - A[j] \\ \text{if the search is successful, } \textbf{output } (i, j, k) \end{array} \right. \end{array} \right.$

The complexity of *Improved3SUM* is $O(n \log n + n^2 \log n) = O(n^2 \log n)$.

A further improvement

In *Improved3SUM*, we **pre-sorted** the array A , which enabled us to do binary searches.

There is a better way to make use of the sorted array, however ...

Namely, for a given $A[i]$, we **simultaneously scan from both ends** of A looking for $A[j] + A[k] = -A[i]$.

We start with $j = i + 1$ and $k = n$.

At any stage of the algorithm, we either **increment** j or **decrement** k (or both, if $A[i] + A[j] + A[k] = 0$).

Does this remind you of a familiar algorithm you have seen in CS 240?

The resulting algorithm will have complexity $O(n \log n + n^2) = O(n^2)$.

A quadratic time algorithm for the “3SUM” problem

Algorithm: *Quadratic3SUM*($A = [A[1], \dots, A[n]]$)

sort A in increasing order

for $i \leftarrow 1$ **to** $n - 2$

do $\left\{ \begin{array}{l} j \leftarrow i + 1 \\ k \leftarrow n \\ \textbf{while } j < k \\ \textbf{do} \left\{ \begin{array}{l} S \leftarrow A[i] + A[j] + A[k] \\ \textbf{if } S < 0 \textbf{ then } j \leftarrow j + 1 \\ \textbf{else if } S > 0 \textbf{ then } k \leftarrow k - 1 \\ \textbf{else } \left\{ \begin{array}{l} \textbf{output } (i, j, k) \\ j \leftarrow j + 1 \\ k \leftarrow k - 1 \end{array} \right. \end{array} \right.$

Problems

Problem: Given a problem instance I for a problem P , carry out a particular computational task.

Problem Instance: **Input** for the specified problem.

Problem Solution: **Output** (correct answer) for the specified problem.

Size of a problem instance: $\text{Size}(I)$ is a positive integer which is a measure of the size of the instance I .

Algorithms and Programs

Algorithm: An algorithm is a step-by-step process (e.g., described in **pseudocode**) for carrying out a series of computations, given some appropriate input.

Algorithm solving a problem: An Algorithm **A** **solves** a problem **P** if, for every instance I of **P**, **A** finds a valid solution for the instance I in finite time.

Program: A program is an **implementation** of an algorithm using a specified computer language.

Running Time

Running Time of a Program: $T_{\mathbf{M}}(I)$ denotes the running time (in seconds) of a program \mathbf{M} on a problem instance I .

Worst-case Running Time as a Function of Input Size: $T_{\mathbf{M}}(n)$ denotes the **maximum** running time of program \mathbf{M} on instances of size n :

$$T_{\mathbf{M}}(n) = \max\{T_{\mathbf{M}}(I) : \text{Size}(I) = n\}.$$

Average-case Running Time as a Function of Input Size: $T_{\mathbf{M}}^{avg}(n)$ denotes the **average** running time of program \mathbf{M} over all instances of size n :

$$T_{\mathbf{M}}^{avg}(n) = \frac{1}{|\{I : \text{Size}(I) = n\}|} \sum_{\{I : \text{Size}(I) = n\}} T_{\mathbf{M}}(I).$$

Complexity

Worst-case complexity of an algorithm: Let $f : \mathbb{Z}^+ \rightarrow \mathbb{R}$. An algorithm A has **worst-case complexity** $f(n)$ if there exists a program M implementing the algorithm A such that $T_M(n) \in \Theta(f(n))$.

Average-case complexity of an algorithm: Let $f : \mathbb{Z}^+ \rightarrow \mathbb{R}$. An algorithm A has **average-case complexity** $f(n)$ if there exists a program M implementing the algorithm A such that $T_M^{avg}(n) \in \Theta(f(n))$.

Running Time vs Complexity

Running time can only be determined by implementing a program and running it on a specific computer.

Running time is influenced by many factors, including the programming language, processor, operating system, etc.

Complexity (AKA **growth rate**) can be analyzed by high-level mathematical analysis. It is **independent** of the above-mentioned factors affecting running time.

Complexity is a less precise measure than running time since it is asymptotic and it incorporates unspecified constant factors and unspecified lower order terms.

However, if algorithm *A* has lower complexity than algorithm *B*, then a program implementing algorithm *A* will be faster than a program implementing algorithm *B* for **sufficiently large inputs**.

Order Notation

O -notation:

$f(n) \in O(g(n))$ if **there exist** constants $c > 0$ and $n_0 > 0$ such that $0 \leq f(n) \leq c g(n)$ for all $n \geq n_0$.

Here the complexity of f is **not higher** than the complexity of g .

Ω -notation:

$f(n) \in \Omega(g(n))$ if **there exist** constants $c > 0$ and $n_0 > 0$ such that $0 \leq c g(n) \leq f(n)$ for all $n \geq n_0$.

Here the complexity of f is **not lower** than the complexity of g .

Θ -notation:

$f(n) \in \Theta(g(n))$ if **there exist** constants $c_1, c_2 > 0$ and $n_0 > 0$ such that $0 \leq c_1 g(n) \leq f(n) \leq c_2 g(n)$ for all $n \geq n_0$.

Here f and g have the **same complexity**.

Order Notation (cont.)

o -notation:

$f(n) \in o(g(n))$ if **for all** constants $c > 0$, there exists a constant $n_0 > 0$ such that $0 \leq f(n) \leq c g(n)$ for all $n \geq n_0$.

Here f has **lower complexity** than g .

ω -notation:

$f(n) \in \omega(g(n))$ if **for all** constants $c > 0$, there exists a constant $n_0 > 0$ such that $0 \leq c g(n) \leq f(n)$ for all $n \geq n_0$.

Here f has **higher complexity** than g .

Techniques for Order Notation

Suppose that $f(n) > 0$ and $g(n) > 0$ for all $n \geq n_0$. Suppose that

$$L = \lim_{n \rightarrow \infty} \frac{f(n)}{g(n)}.$$

Then

$$f(n) \in \begin{cases} o(g(n)) & \text{if } L = 0 \\ \Theta(g(n)) & \text{if } 0 < L < \infty \\ \omega(g(n)) & \text{if } L = \infty. \end{cases}$$

Relationships between Order Notations

$$f(n) \in \Theta(g(n)) \Leftrightarrow g(n) \in \Theta(f(n))$$

$$f(n) \in O(g(n)) \Leftrightarrow g(n) \in \Omega(f(n))$$

$$f(n) \in o(g(n)) \Leftrightarrow g(n) \in \omega(f(n))$$

$$f(n) \in \Theta(g(n)) \Leftrightarrow f(n) \in O(g(n)) \text{ and } f(n) \in \Omega(g(n))$$

$$f(n) \in o(g(n)) \Rightarrow f(n) \in O(g(n))$$

$$f(n) \in \omega(g(n)) \Rightarrow f(n) \in \Omega(g(n))$$

Algebra of Order Notations

“Maximum” rules: Suppose that $f(n) > 0$ and $g(n) > 0$ for all $n \geq n_0$.
Then:

$$O(f(n) + g(n)) = O(\max\{f(n), g(n)\})$$

$$\Theta(f(n) + g(n)) = \Theta(\max\{f(n), g(n)\})$$

$$\Omega(f(n) + g(n)) = \Omega(\max\{f(n), g(n)\})$$

“Summation” rules:

$$O\left(\sum_{i \in I} f(i)\right) = \sum_{i \in I} O(f(i))$$

$$\Theta\left(\sum_{i \in I} f(i)\right) = \sum_{i \in I} \Theta(f(i))$$

$$\Omega\left(\sum_{i \in I} f(i)\right) = \sum_{i \in I} \Omega(f(i))$$

Sequences

Arithmetic sequence:

$$\sum_{i=0}^{n-1} (a + di) = na + \frac{dn(n-1)}{2} \in \Theta(n^2).$$

Geometric sequence:

$$\sum_{i=0}^{n-1} ar^i = \begin{cases} a \frac{r^n - 1}{r - 1} \in \Theta(r^n) & \text{if } r > 1 \\ na \in \Theta(n) & \text{if } r = 1 \\ a \frac{1 - r^n}{1 - r} \in \Theta(1) & \text{if } 0 < r < 1. \end{cases}$$

Arithmetic-geometric sequence:

$$\sum_{i=0}^{n-1} (a + di)r^i = \frac{a}{1 - r} - \frac{(a + (n-1)d)r^n}{1 - r} + \frac{dr(1 - r^{n-1})}{(1 - r)^2}$$

provided that $r \neq 1$.

Sequences (cont.)

Harmonic sequence:

$$H_n = \sum_{i=1}^n \frac{1}{i} \in \Theta(\log n)$$

More precisely, it is possible to prove that

$$\lim_{n \rightarrow \infty} (H_n - \ln n) = \gamma,$$

where $\gamma \approx 0.57721$ is **Euler's constant**.

Miscellaneous Formulae

$$\log_b xy = \log_b x + \log_b y$$

$$\log_b x/y = \log_b x - \log_b y$$

$$\log_b 1/x = -\log_b x$$

$$\log_b x^y = y \log_b x$$

$$\log_b a = \frac{1}{\log_a b}$$

$$\log_b a = \frac{\log_c a}{\log_c b}$$

$$a^{\log_b c} = c^{\log_b a}$$

$$n! \in \Theta(n^{n+1/2}e^{-n})$$

$$\log n! \in \Theta(n \log n)$$

Techniques for Algorithm Analysis

Two general strategies are as follows:

- Use Θ -bounds **throughout the analysis** and thereby obtain a Θ -bound for the complexity of the algorithm.
- Prove a O -bound and a **matching** Ω -bound **separately** to get a Θ -bound. Sometimes this technique is easier because arguments for O -bounds may use simpler upper bounds (and arguments for Ω -bounds may use simpler lower bounds) than arguments for Θ -bounds do.

Techniques for Loop Analysis

Identify **elementary operations** that require constant time (denoted $\Theta(1)$ time).

The complexity of a loop is expressed as the **sum** of the complexities of each iteration of the loop.

Analyze independent loops **separately**, and then **add** the results: use “maximum rules” and simplify whenever possible.

If loops are nested, start with the **innermost loop** and proceed outwards. In general, this kind of analysis requires evaluation of **nested summations**.

Example of Loop Analysis

Algorithm: *LoopAnalysis1*($n : integer$)

```

(1)  $sum \leftarrow 0$ 
(2) for  $i \leftarrow 1$  to  $n$ 
    do  $\left\{ \begin{array}{l} \textbf{for } j \leftarrow 1 \textbf{ to } i \\ \textbf{do } \left\{ \begin{array}{l} sum \leftarrow sum + (i - j)^2 \\ sum \leftarrow sum / i \end{array} \right. \end{array} \right.$ 
(3) return ( $sum$ )
  
```

Θ -bound analysis

(1) $\Theta(1)$

(2) Complexity of inner **for** loop: $\Theta(i)$

Complexity of outer **for** loop: $\sum_{i=1}^n \Theta(i) = \Theta(n^2)$

Note: $\sum_{i=1}^n i = n(n+1)/2$

(3) $\Theta(1)$

total $\Theta(n^2)$

Example of Loop Analysis (cont.)

Proving separate O - and Ω -bounds

We focus on the two nested **for** loops (i.e., (2)).

The total number of iterations is $\sum_{i=1}^n i$, with $\Theta(1)$ time per iteration.

Upper bound:

$$\sum_{i=1}^n O(i) \leq \sum_{i=1}^n O(n) = O(n^2).$$

Lower bound:

$$\sum_{i=1}^n \Omega(i) \geq \sum_{i=n/2}^n \Omega(i) \geq \sum_{i=n/2}^n \Omega(n/2) = \Omega(n^2/4) = \Omega(n^2).$$

Since the upper and lower bounds **match**, the complexity is $\Theta(n^2)$.

Another Example of Loop Analysis

Algorithm: *LoopAnalysis2*($A : \text{array}; n : \text{integer}$)

$max \leftarrow 0$

for $i \leftarrow 1$ **to** n

do { **for** $j \leftarrow i$ **to** n

do { $sum \leftarrow 0$

for $k \leftarrow i$ **to** j

do { $sum \leftarrow sum + A[k]$

if $sum > max$

then $max \leftarrow sum$

return (max)

Yet Another Example of Loop Analysis

Algorithm: *LoopAnalysis3*($n : \text{integer}$)

$sum \leftarrow 0$

for $i \leftarrow 1$ **to** n

do $\left\{ \begin{array}{l} j \leftarrow i \\ \textbf{while } j \geq 1 \\ \quad \textbf{do } \left\{ \begin{array}{l} sum \leftarrow sum + i/j \\ j \leftarrow \left\lfloor \frac{j}{2} \right\rfloor \end{array} \right. \end{array} \right.$

return (sum)

Table of Contents

3 Divide-and-Conquer Algorithms

- Recurrence Relations
- Master Theorem
- Divide-and-Conquer Design Strategy
- Mergesort
- Divide-and-Conquer Recurrence Relations
- Max-Min Problem
- Non-dominated Points
- Closest Pair
- Multiprecision Multiplication
- Matrix Multiplication
- Selection and Median

Recurrence Relations

Suppose a_1, a_2, \dots , is an infinite sequence of real numbers.

A **recurrence relation** is a formula that expresses a general term a_n in terms of one or more previous terms a_1, \dots, a_{n-1} .

A recurrence relation will also specify one or more **initial values** starting at a_1 .

Solving a recurrence relation means finding a formula for a_n that does **not** involve any previous terms a_1, \dots, a_{n-1} .

There are many methods of solving recurrence relations. Two important methods are **guess-and-check** and the **recursion tree method**.

We will make extensive use of the recursion tree method. However, we first take a quick look at the guess-and-check method.

Guess-and-check Method

- step 1** Tabulate some values a_1, a_2, \dots using the recurrence relation.
- step 2** Guess that the solution a_n has a specific form, involving undetermined constants.
- step 3** Use a_1, a_2, \dots to determine specific values for the unspecified constants.
- step 4** Use **induction** to prove your guess for a_n is correct.

Example of the Guess-and-check Method

Suppose we have the recurrence $T(n) = T(n-1) + 6n - 5$, $T(0) = 4$.

We compute a few values: $T(1) = 5$, $T(2) = 12$, $T(3) = 25$, $T(4) = 44$.

If we are sufficiently perspicacious, we might guess that $T(n)$ is a **quadratic function**, e.g., $T(n) = an^2 + bn + c$.

Next, we use $T(0) = 4$, $T(1) = 5$, $T(2) = 12$ to compute a , b and c by solving three equations in three unknowns.

We get $a = 3$, $b = -2$, $c = 4$.

Now we can use induction to prove that $T(n) = 3n^2 - 2n + 4$ for all $n \geq 0$.

Recursion Tree Method

The following recurrence relation arises in the analysis of *Mergesort*:

$$T(n) = \begin{cases} 2T\left(\frac{n}{2}\right) + cn & \text{if } n > 1 \text{ is a power of 2} \\ d & \text{if } n = 1, \end{cases}$$

where c and d are constants.

We can solve this recurrence relation when n is a power of two, by constructing a **recursion tree**, as follows:

- step 1** Start with a **one-node tree**, say N , having the value $T(n)$.
- step 2** Grow **two children** of N . These children, say N_1 and N_2 , have the value $T(n/2)$, and the value of N is replaced by cn .
- step 3** Repeat this process recursively, terminating when a node receives the value $T(1) = d$.
- step 4** Sum the values on each level of the tree, and then compute the **sum of all these sums**; the result is $T(n)$.

Master Theorem

The **Master Theorem** provides a formula for the solution of many recurrence relations typically encountered in the analysis of algorithms.

The following is a simplified version of the **Master Theorem**:

Theorem

Suppose that $a \geq 1$ and $b > 1$. Consider the recurrence

$$T(n) = a T\left(\frac{n}{b}\right) + \Theta(n^y), \quad (1)$$

where n is a power of b . Denote $x = \log_b a$. Then

$$T(n) \in \begin{cases} \Theta(n^x) & \text{if } y < x \\ \Theta(n^x \log n) & \text{if } y = x \\ \Theta(n^y) & \text{if } y > x. \end{cases}$$

Proof of the Master Theorem (simplified version)

Suppose that $a \geq 1$ and $b \geq 2$ are integers and

$$T(n) = a T\left(\frac{n}{b}\right) + c n^y, \quad T(1) = d.$$

Let $n = b^j$.

level	# nodes	value at each node	value of the level
j	1	$c n^y$	$c n^y$
$j - 1$	a	$c (n/b)^y$	$c a (n/b)^y$
$j - 2$	a^2	$c (n/b^2)^y$	$c a^2 (n/b^2)^y$
\vdots	\vdots	\vdots	\vdots
1	a^{j-1}	$c (n/b^{j-1})^y$	$c a^{j-1} (n/b^{j-1})^y$
0	a^j	d	$d a^j$

Computing $T(n)$

Summing the values at all levels of the recursion tree, we have that

$$T(n) = d a^j + c n^y \sum_{i=0}^{j-1} \left(\frac{a}{b^y} \right)^i.$$

Recall that $b^x = a$ and $n = b^j$. Hence $a^j = (b^x)^j = (b^j)^x = n^x$.

The formula for $T(n)$ is a **geometric sequence** with ratio $r = a/b^y = b^{x-y}$:

$$T(n) = d n^x + c n^y \sum_{i=0}^{j-1} r^i.$$

There are **three cases**, depending on whether $r > 1$, $r = 1$ or $r < 1$.

Complexity of $T(n)$

case	r	y, x	complexity of $T(n)$
heavy leaves	$r > 1$	$y < x$	$T(n) \in \Theta(n^x)$
balanced	$r = 1$	$y = x$	$T(n) \in \Theta(n^x \log n)$
heavy top	$r < 1$	$y > x$	$T(n) \in \Theta(n^y)$

heavy leaves means that the value of the recursion tree is dominated by the values of the leaf nodes.

balanced means that the values of the levels of the recursion tree are constant (except for the last level).

heavy top means that the value of the recursion tree is dominated by the value of the root node.

Master Theorem (modified general version)

Theorem

Suppose that $a \geq 1$ and $b > 1$. Consider the recurrence

$$T(n) = a T\left(\frac{n}{b}\right) + f(n),$$

where n is a power of b . Denote $x = \log_b a$. Then

$$T(n) \in \begin{cases} \Theta(n^x) & \text{if } f(n) \in O(n^{x-\epsilon}) \text{ for some } \epsilon > 0 \\ \Theta(n^x \log n) & \text{if } f(n) \in \Theta(n^x) \\ \Theta(f(n)) & \text{if } f(n)/n^{x+\epsilon} \text{ is an increasing function of } n \\ & \text{for some } \epsilon > 0. \end{cases}$$

The Divide-and-Conquer Design Strategy

divide: Given a problem instance I , construct one or more smaller problem instances, denoted I_1, \dots, I_a (these are called **subproblems**). Usually, we want the size of these subproblems to be small compared to the size of I , e.g., half the size.

conquer: For $1 \leq j \leq a$, solve instance I_j recursively, obtaining solutions S_1, \dots, S_a .

combine: Given S_1, \dots, S_a , use an appropriate **combining** function to find the solution S to the problem instance I , i.e.,
 $S \leftarrow \text{Combine}(S_1, \dots, S_a)$.

Example: Design of Mergesort

Here, a problem instance consists of an array A of n integers, which we want to sort in increasing order. The size of the problem instance is n .

divide: Split A into two subarrays: A_L consists of the first $\lceil \frac{n}{2} \rceil$ elements in A and A_R consists of the last $\lfloor \frac{n}{2} \rfloor$ elements in A .

conquer: Run *Mergesort* on A_L and A_R .

combine: After A_L and A_R have been sorted, use a function *Merge* to merge A_L and A_R into a single sorted array. Recall that this can be done in time $\Theta(n)$ with a single pass through A_L and A_R . We simply keep track of the “current” element of A_L and A_R , always copying the smaller one into the sorted array.

Mergesort

Algorithm: *Mergesort*($A : \text{array}; n : \text{integer}$)

if $n = 1$

then $S \leftarrow A$

else $\left\{ \begin{array}{l} n_L \leftarrow \lceil \frac{n}{2} \rceil \\ n_R \leftarrow \lfloor \frac{n}{2} \rfloor \\ A_L \leftarrow [A[1], \dots, A[n_L]] \\ A_R \leftarrow [A[n_L + 1], \dots, A[n]] \\ S_L \leftarrow \textit{Mergesort}(A_L, n_L) \\ S_R \leftarrow \textit{Mergesort}(A_R, n_R) \\ S \leftarrow \textit{Merge}(S_L, n_L, S_R, n_R) \end{array} \right.$

return (S, n)

Analysis of Mergesort

Let $T(n)$ denote the time to run *Mergesort* on an array of length n .

divide takes time $\Theta(n)$

conquer takes time $T(\lceil \frac{n}{2} \rceil) + T(\lfloor \frac{n}{2} \rfloor)$

combine takes time $\Theta(n)$

Recurrence relation:

$$T(n) = \begin{cases} T(\lceil \frac{n}{2} \rceil) + T(\lfloor \frac{n}{2} \rfloor) + \Theta(n) & \text{if } n > 1 \\ \Theta(1) & \text{if } n = 1. \end{cases}$$

Sloppy and Exact Recurrence Relations

It is simpler to replace the $\Theta(n)$ term by cn , where c is an unspecified constant. The resulting recurrence relation is called the **exact recurrence**.

$$T(n) = \begin{cases} T(\lceil \frac{n}{2} \rceil) + T(\lfloor \frac{n}{2} \rfloor) + cn & \text{if } n > 1 \\ d & \text{if } n = 1. \end{cases}$$

If we then remove the floors and ceilings, we obtain the so-called **sloppy recurrence**:

$$T(n) = \begin{cases} 2T(\frac{n}{2}) + cn & \text{if } n > 1 \\ d & \text{if } n = 1. \end{cases}$$

The exact and sloppy recurrences are **identical** when n is a power of two. Further, the sloppy recurrence makes sense **only** when n is a power of two.

Complexity of the Solution to the Exact Recurrence

The **Master Theorem** provides the **exact** solution of the recurrence when $n = 2^j$ (it is in fact a **proof** for these values of n).

We can express this solution (for powers of 2) as a function of n , using Θ -notation.

It can be shown that the resulting function of n will in fact yield the **complexity** of the solution of the exact recurrence for **all values** of n .

This derivation of the complexity of $T(n)$ is **not a proof**, however. If a rigorous mathematical proof is required, then it is necessary to use **induction** along with the **exact recurrence**.

The Max-Min Problem

Let's design a divide-and-conquer algorithm for the **Max-Min** problem.

Divide: Suppose we split A into two equal-sized subarrays, A_L and A_R .

Conquer: We find the maximum and minimum elements in each subarray recursively, obtaining max_L , min_L , max_R and min_R .

Combine: Then we can easily “combine” the solutions to the two subproblems to solve the original problem instance:

$$max \leftarrow \max\{max_L, max_R\}$$

and

$$min \leftarrow \min\{min_L, min_R\}$$

The Max-Min Problem (cont.)

The recurrence relation describing the complexity of the running time is $T(n) = 2T(n/2) + \Theta(1)$.

The **Master Theorem** shows that the $T(n) \in \Theta(n)$.

However, we can also count the **exact number** of comparisons done by the algorithm, obtaining the (sloppy) recurrence

$$C(n) = 2C(n/2) + 2, \quad C(2) = 1.$$

For n a power of 2, the solution to this recurrence relation is $C(n) = 3n/2 - 2$, so the divide-and-conquer algorithm is **optimal** for these values of n (see slide 26).

Non-dominated Points

Given two points $(x_1, y_1), (x_2, y_2)$ in the Euclidean plane, we say that (x_1, y_1) **dominates** (x_2, y_2) if $x_1 \geq x_2$ and $y_1 \geq y_2$.

Problem

Non-dominated Points

Instance: A set S of n points in the Euclidean plane, say $S = \{S[1], \dots, S[n]\}$.

Question: Find all the **non-dominated points** in S , i.e., all the points that are not dominated by any other point in S .

Non-dominated Points has a trivial $\Theta(n^2)$ algorithm to solve it, based on comparing all pairs of points in S . Can we do better?

Problem Decomposition

Observe that the non-dominated points form a **staircase** such that all the other points are “under” the staircase.

Suppose we **pre-sort** the points in S with respect to their x -co-ordinates. This takes time $\Theta(n \log n)$.

Divide: Let the first $n/2$ points be denoted S_1 and let the last $n/2$ points be denoted S_2 .

Conquer: Recursively solve the subproblems defined by the two instances S_1 and S_2 .

Combine: Given the non-dominated points in S_1 and the non-dominated points in S_2 , how do we find the non-dominated points in S ?

Observe that **no point in S_1 dominates a point in S_2** .

Therefore we only need to eliminate the points in S_1 that are dominated by a point in S_2 . This can be done in time $O(n)$.

Non-dominated Points

Algorithm: *Non-dominated*(S)

comment: the n points in S are pre-sorted WRT their x -co-ordinates

if $n = 1$ **then return** ($\{S[1]\}$)

else $\left\{ \begin{array}{l} \{Q[1], \dots, Q[\ell]\} \leftarrow \textit{Non-dominated}(\{S[1], \dots, S[\lfloor n/2 \rfloor]\}) \\ \{(R[1], \dots, R[m])\} \leftarrow \textit{Non-dominated}(\{S[\lfloor n/2 \rfloor + 1], \dots, S[n]\}) \\ i \leftarrow 1 \\ \textbf{while } i \leq \ell \textbf{ and } Q[i].y > R[1].y \\ \quad \textbf{do } i \leftarrow i + 1 \end{array} \right.$

return ($\{Q[1], \dots, Q[i - 1], R[1], \dots, R[m]\}$)

Closest Pair

Problem

Closest Pair

Instance: a set Q of n distinct points in the Euclidean plane,

$$Q = \{Q[1], \dots, Q[n]\}.$$

Find: Two distinct points $Q[i] = (x, y)$, $Q[j] = (x', y')$ such that the Euclidean distance

$$\sqrt{(x' - x)^2 + (y' - y)^2}$$

is minimized.

Closest Pair: Problem Decomposition

Suppose we presort the points in Q with respect to their x -coordinates (this takes time $\Theta(n \log n)$).

Then we can easily find the vertical line that partitions the set of points Q into two sets of size $n/2$: this line has equation $x = Q[m].x$, where $m = n/2$.

The set Q is global with respect to the recursive procedure *ClosestPair1*.

At any given point in the recursion, we are examining a subarray $(Q[\ell], \dots, Q[r])$, and $m = \lfloor (\ell + r)/2 \rfloor$.

We call *ClosestPair1*(1, n) to solve the given problem instance.

Closest Pair: Solution 1

Algorithm: *ClosestPair1*(ℓ, r)

if $\ell = r$ **then** $\delta \leftarrow \infty$

else $\left\{ \begin{array}{l} m \leftarrow \lfloor (\ell + r)/2 \rfloor \\ \delta_L \leftarrow \textit{ClosestPair1}(\ell, m) \\ \delta_R \leftarrow \textit{ClosestPair1}(m + 1, r) \\ \delta \leftarrow \min\{\delta_L, \delta_R\} \\ R \leftarrow \textit{SelectCandidates}(\ell, r, \delta, Q[m].x) \\ R \leftarrow \textit{SortY}(R) \\ \delta \leftarrow \textit{CheckStrip}(R, \delta) \end{array} \right.$

return (δ)

Selecting Candidates from the Vertical Strip

Algorithm: *SelectCandidates*(ℓ, r, δ, x_{mid})

$j \leftarrow 0$

for $i \leftarrow \ell$ **to** r

do $\left\{ \begin{array}{l} \text{if } |Q[i].x - x_{mid}| \leq \delta \\ \text{then } \left\{ \begin{array}{l} j \leftarrow j + 1 \\ R[j] \leftarrow Q[i] \end{array} \right. \end{array} \right.$

return (R)

Checking the Vertical Strip

Algorithm: *CheckStrip*(R, δ)

$t \leftarrow \text{size}(R)$

$\delta' \leftarrow \delta$

for $j \leftarrow 1$ **to** $t - 1$

do $\left\{ \begin{array}{l} \textbf{for } k \leftarrow j + 1 \textbf{ to } \min\{t, j + 7\} \\ \textbf{do } \left\{ \begin{array}{l} x \leftarrow R[j].x \\ x' \leftarrow R[k].x \\ y \leftarrow R[j].y \\ y' \leftarrow R[k].y \\ \delta' \leftarrow \min \left\{ \delta', \sqrt{(x' - x)^2 + (y' - y)^2} \right\} \end{array} \right. \end{array} \right.$

return (δ')

Closest Pair: Solution 2

Algorithm: *ClosestPair2*(ℓ, r)

if $\ell = r$ **then** $\delta \leftarrow \infty$

else $\left\{ \begin{array}{l} m \leftarrow \lfloor (\ell + r) / 2 \rfloor \\ \delta_L \leftarrow \textit{ClosestPair2}(\ell, m) \\ \text{comment: } Q[\ell], \dots, Q[m] \text{ is sorted WRT } y\text{-coordinates} \\ \delta_R \leftarrow \textit{ClosestPair2}(m + 1, r) \\ \text{comment: } Q[m + 1], \dots, Q[r] \text{ is sorted WRT } y\text{-coordinates} \\ \delta \leftarrow \min\{\delta_L, \delta_R\} \\ \textit{Merge}(\ell, m, r) \\ R \leftarrow \textit{SelectCandidates}(\ell, r, \delta, Q[m].x) \\ \delta \leftarrow \textit{CheckStrip}(R, \delta) \end{array} \right.$

return (δ)

Multiprecision Multiplication

Problem

Multiprecision Multiplication

Instance: Two k -bit positive integers, X and Y , having binary representations

$$X = [X[k - 1], \dots, X[0]]$$

and

$$Y = [Y[k - 1], \dots, Y[0]].$$

Question: Compute the $2k$ -bit positive integer $Z = XY$, where

$$Z = (Z[2k - 1], \dots, Z[0]).$$

We are interested in the **bit complexity** of algorithms that solve **Multiprecision Multiplication**, which means that the complexity is expressed as a function of k (the size of the problem instance is $2k$ bits).

Not-So-Fast D&C Multiprecision Multiplication

Algorithm: *NotSoFastMultiply*(X, Y, k)

if $k = 1$

then $Z \leftarrow X[0] \times Y[0]$

else
$$\begin{cases} Z_1 \leftarrow \textit{NotSoFastMultiply}(X_L, Y_L, k/2) \\ Z_2 \leftarrow \textit{NotSoFastMultiply}(X_R, Y_R, k/2) \\ Z_3 \leftarrow \textit{NotSoFastMultiply}(X_L, Y_R, k/2) \\ Z_4 \leftarrow \textit{NotSoFastMultiply}(X_R, Y_L, k/2) \\ Z \leftarrow \textit{LeftShift}(Z_1, k) + Z_2 + \textit{LeftShift}(Z_3 + Z_4, k/2) \end{cases}$$

return (Z)

Fast D&C Multiprecision Multiplication

Algorithm: *FastMultiply*(X, Y, k)

if $k = 1$

then $Z \leftarrow X[0] \times Y[0]$

else
$$\left\{ \begin{array}{l} X_T \leftarrow X_L + X_R \\ Y_T \leftarrow Y_L + Y_R \\ Z_1 \leftarrow \text{FastMultiply}(X_L, Y_L, k/2) \\ Z_2 \leftarrow \text{FastMultiply}(X_R, Y_R, k/2) \\ Z_3 \leftarrow \text{FastMultiply}(X_T, Y_T, k/2), \\ Z \leftarrow \text{LeftShift}(Z_1, k) + Z_2 + \text{LeftShift}(Z_3 - Z_1 - Z_2, k/2) \end{array} \right.$$

return (Z)

Matrix Multiplication

Problem

Matrix Multiplication

Instance: *Two n by n matrices, A and B .*

Question: *Compute the n by n matrix product $C = AB$.*

The naive algorithm for **Matrix Multiplication** has complexity $\Theta(n^3)$.

Matrix Multiplication: Problem Decomposition

$$A = \begin{pmatrix} a & b \\ c & d \end{pmatrix}, \quad B = \begin{pmatrix} e & f \\ g & h \end{pmatrix}, \quad C = AB = \begin{pmatrix} r & s \\ t & u \end{pmatrix}$$

If A, B are n by n matrices, then $a, b, \dots, h, r, s, t, u$ are $\frac{n}{2}$ by $\frac{n}{2}$ matrices, where

$$r = a e + b g$$

$$s = a f + b h$$

$$t = c e + d g$$

$$u = c f + d h$$

We require 8 multiplications of $\frac{n}{2}$ by $\frac{n}{2}$ matrices in order to compute $C = AB$.

Efficient D&C Matrix Multiplication

Define

$$P_1 = a(f - h)$$

$$P_3 = (c + d)e$$

$$P_5 = (a + d)(e + h)$$

$$P_7 = (a - c)(e + f).$$

$$P_2 = (a + b)h$$

$$P_4 = d(g - e)$$

$$P_6 = (b - d)(g + h)$$

Then, compute

$$r = P_5 + P_4 - P_2 + P_6$$

$$t = P_3 + P_4$$

$$s = P_1 + P_2$$

$$u = P_5 + P_1 - P_3 - P_7.$$

We now require only 7 multiplications of $\frac{n}{2}$ by $\frac{n}{2}$ matrices in order to compute $C = AB$.

Selection

Problem

Selection

Instance: *An array $A[1], \dots, A[n]$ of distinct integer values, and an integer k , where $1 \leq k \leq n$.*

Find: *The k th smallest integer in the array A .*

The problem **Median** is the special case of **Selection** where $k = \lceil \frac{n}{2} \rceil$.

QuickSelect

Suppose we choose a **pivot** element y in the array A , and we **restructure** A so that all elements less than y precede y in A , and all elements greater than y occur after y in A . (This is exactly what is done in *Quicksort*, and it takes **linear time**.)

Suppose that $A[\text{posn}] = y$ after restructuring. Let A_L be the subarray $A[1], \dots, A[\text{posn} - 1]$ and let A_R be the subarray (of size $n - \text{posn}$) $A[\text{posn} + 1], \dots, A[n]$.

Then the k th smallest element of A is

$$\begin{cases} y & \text{if } k = \text{posn} \\ \text{the } k\text{th smallest element of } A_L & \text{if } k < \text{posn} \\ \text{the } (k - \text{posn})\text{th smallest element of } A_R & \text{if } k > \text{posn}. \end{cases}$$

We make (at most) one recursive call at each level of the recursion.

Average-case Analysis of QuickSelect

We say that a pivot is **good** if $posn$ is in the middle half of A .

The probability that a pivot is good is $1/2$.

On average, after **two iterations**, we will encounter a good pivot.

If a pivot is good, then $|A_L| \leq 3n/4$ and $|A_R| \leq 3n/4$.

With an **expected** linear amount of work, the size of the subproblem is reduced by at least 25%.

It follows that the average-case complexity of the **QuickSelect** is linear.

Achieving $O(n)$ Worst-Case Complexity: A Strategy for Choosing the Pivot

We choose the pivot to be a certain **median-of-medians**:

- step 1** Given $n \geq 15$, write $n = 10r + 5 + \theta$, where $r \geq 1$ and $0 \leq \theta \leq 9$.
- step 2** Divide A into $2r + 1$ disjoint subarrays of 5 elements. Denote these subarrays by B_1, \dots, B_{2r+1} .
- step 3** For $1 \leq i \leq 2r + 1$, find the median of B_i (nonrecursively), and denote it by m_i .
- step 4** Define M to be the array consisting of elements m_1, \dots, m_{2r+1} .
- step 5** Find the median y of the array M (recursively).
- step 6** Use the element y as the pivot for A .

Median-of-medians-QuickSelect

Algorithm: *Mom-QuickSelect*(k, n, A)

1. **if** $n \leq 14$ **then** sort A and **return** ($A[k]$)
2. write $n = 10r + 5 + \theta$, where $0 \leq \theta \leq 9$
3. construct B_1, \dots, B_{2r+1} (subarrays of A , each of size 5)
4. find medians m_1, \dots, m_{2r+1} (non-recursively)
5. $M \leftarrow [m_1, \dots, m_{2r+1}]$
6. $y \leftarrow \textit{Mom-QuickSelect}(r + 1, 2r + 1, M)$
7. $(A_L, A_R, \textit{posn}) \leftarrow \textit{Restructure}(A, y)$
8. **if** $k = \textit{posn}$ **then return** (y)
9. **else if** $k < \textit{posn}$ **then return** ($\textit{Mom-QuickSelect}(k, \textit{posn} - 1, A_L)$)
10. **else return** ($\textit{Mom-QuickSelect}(k - \textit{posn}, n - \textit{posn}, A_R)$)

Worst-case Analysis of *Mom-QuickSelect*

When the pivot is the median-of-medians, we have that $|A_L| \leq \lfloor \frac{7n+12}{10} \rfloor$ and $|A_R| \leq \lfloor \frac{7n+12}{10} \rfloor$.

The *Mom-QuickSelect* algorithm requires **two recursive calls**.

The worst-case complexity $T(n)$ of this algorithm satisfies the following recurrence:

$$T(n) \leq \begin{cases} T(\lfloor \frac{n}{5} \rfloor) + T(\lfloor \frac{7n+12}{10} \rfloor) + \Theta(n) & \text{if } n \geq 15 \\ \Theta(1) & \text{if } n \leq 14. \end{cases}$$

It can be shown that $T(n)$ is $O(n)$.

Table of Contents

4 Greedy Algorithms

- Optimization Problems
- Design Strategy
- Interval Selection
- Interval Colouring
- Knapsack
- Coin Changing
- Stable Marriage Problem

Optimization Problems

Problem: Given a problem instance, find a feasible solution that maximizes (or minimizes) a certain objective function.

Problem Instance: **Input** for the specified problem.

Problem Constraints: **Requirements** that must be satisfied by any feasible solution.

Feasible Solution: For any problem instance I , $\text{feasible}(I)$ is the set of all outputs (i.e., solutions) for the instance I that satisfy the given constraints.

Objective Function: A function $f : \text{feasible}(I) \rightarrow \mathbb{R}^+ \cup \{0\}$. We often think of f as being a **profit** or a **cost** function.

Optimal Solution: A feasible solution $X \in \text{feasible}(I)$ such that the profit $f(X)$ is maximized (or the cost $f(X)$ is minimized).

The Greedy Method

partial solutions

Given a problem instance I , it should be possible to write a feasible solution X as a tuple $[x_1, x_2, \dots, x_n]$ for some integer n , where $x_i \in \mathcal{X}$ for all i . A tuple $[x_1, \dots, x_i]$ where $i < n$ is a **partial solution** if no constraints are violated. Note: it may be the case that a partial solution cannot be extended to a feasible solution.

choice set

For a partial solution $X = [x_1, \dots, x_i]$ where $i < n$, we define the **choice set**

$$\text{choice}(X) = \{y \in \mathcal{X} : [x_1, \dots, x_i, y] \text{ is a partial solution}\}.$$

The Greedy Method (cont.)

local evaluation criterion

For any $y \in \mathcal{X}$, $g(y)$ is a **local evaluation criterion** that measures the cost or profit of including y in a (partial) solution.

extension

Given a partial solution $X = [x_1, \dots, x_i]$ where $i < n$, choose $y \in \text{choice}(X)$ so that $g(y)$ is as small (or large) as possible. Update X to be the $(i + 1)$ -tuple $[x_1, \dots, x_i, y]$.

greedy algorithm

Starting with the “empty” partial solution, repeatedly extend it until a feasible solution X is constructed. This feasible solution may or may not be optimal.

Features of the Greedy Method

Greedy algorithms do no **looking ahead** and no **backtracking**.

Greedy algorithms can usually be implemented efficiently. Often they consist of a **preprocessing step** based on the function g , followed by a **single pass** through the data.

In a greedy algorithm, only **one feasible solution** is constructed.

The execution of a greedy algorithm is based on **local criteria** (i.e., the values of the function g).

Correctness: For certain greedy algorithms, it is possible to prove that they always yield optimal solutions. However, these proofs can be tricky and complicated!

Interval Selection

Problem

Interval Selection

Instance: A set $\mathcal{A} = \{A_1, \dots, A_n\}$ of **intervals**.

For $1 \leq i \leq n$, $A_i = [s_i, f_i)$, where s_i is the **start time** of interval A_i and f_i is the **finish time** of A_i .

Feasible solution: A subset $\mathcal{B} \subseteq \mathcal{A}$ of **pairwise disjoint intervals**.

Find: A feasible solution of maximum size (i.e., one that maximizes $|\mathcal{B}|$).

Possible Greedy Strategies for Interval Selection

- 1 Choose the **earliest starting** interval that is disjoint from all previously chosen intervals (i.e., the local evaluation criterion is s_i).
- 2 Choose the interval of **minimum duration** that is disjoint from all previously chosen intervals (i.e., the local evaluation criterion is $f_i - s_i$).
- 3 Choose the **earliest finishing** interval that is disjoint from all previously chosen intervals (i.e., the local evaluation criterion is f_i).

Does one of these strategies yield a **correct** greedy algorithm?

A Greedy Algorithm for Interval Selection

Algorithm: *GreedyIntervalSelection*(\mathcal{A})

rename the intervals, by sorting if necessary, so that $f_1 \leq \dots \leq f_n$

$\mathcal{B} \leftarrow \{A_1\}$

$prev \leftarrow 1$

comment: $prev$ is the index of the last selected interval

for $i \leftarrow 2$ **to** n

do $\left\{ \begin{array}{l} \text{if } s_i \geq f_{prev} \\ \text{then } \left\{ \begin{array}{l} \mathcal{B} \leftarrow \mathcal{B} \cup \{A_i\} \\ prev \leftarrow i \end{array} \right. \end{array} \right.$

return (\mathcal{B})

Interval Colouring

Problem

Interval Colouring

Instance: A set $\mathcal{A} = \{A_1, \dots, A_n\}$ of **intervals**.

For $1 \leq i \leq n$, $A_i = [s_i, f_i)$, where s_i is the **start time** of interval A_i and f_i is the **finish time** of A_i .

Feasible solution: A **c-colouring** is a mapping $col : \mathcal{A} \rightarrow \{1, \dots, c\}$ that assigns each interval a **colour** such that two intervals receiving the same colour are always disjoint.

Find: A c -colouring of \mathcal{A} with the minimum number of colours.

Greedy Strategies for Interval Colouring

As usual, we consider the intervals one at a time.

At a given point in time, suppose we have coloured the first $i < n$ intervals using d colours.

We will colour the $(i + 1)$ st interval with the **any permissible colour**. If it cannot be coloured using any of the existing d colours, then we introduce a **new colour** and d is increased by 1.

Question: In **what order** should we consider the intervals?

A Greedy Algorithm for Interval Colouring

Algorithm: *GreedyIntervalColouring*(\mathcal{A})

rename the intervals, by sorting if necessary, so that $s_1 \leq \dots \leq s_n$

$d \leftarrow 1$

$colour[1] \leftarrow 1$

$finish[1] \leftarrow f_1$

for $i \leftarrow 2$ **to** n

do $\left\{ \begin{array}{l} flag \leftarrow \text{false} \\ c \leftarrow 1 \\ \text{while } c \leq d \text{ and (not } flag) \\ \quad \text{do } \left\{ \begin{array}{l} \text{if } finish[c] \leq s_i \text{ then } \left\{ \begin{array}{l} colour[i] \leftarrow c \\ finish[c] \leftarrow f_i \\ flag \leftarrow \text{true} \end{array} \right. \\ \text{else } c \leftarrow c + 1 \end{array} \right. \\ \text{if not } flag \text{ then } \left\{ \begin{array}{l} d \leftarrow d + 1 \\ colour[i] \leftarrow d \\ finish[d] \leftarrow f_i \end{array} \right. \end{array} \right.$

return $(d, colour)$

Comments and Questions

In the algorithm on the previous slide, at any point in time, $finish[c]$ denotes the finishing time of the **last interval** that has received colour c . Therefore, a new interval A_i can be assigned colour c if $s_i \geq finish[c]$.

The complexity of the algorithm is $O(n \times D)$, where D is the value of d returned by the algorithm.

If it turns out that $D \in \Omega(n)$, then the best we can say is that the complexity is $O(n^2)$.

What **inefficiencies** exist in this algorithm?

What **data structure** would allow a more efficient algorithm to be designed?

What would be the complexity of an algorithm making use of an appropriate data structure?

Knapsack Problems

Problem

Knapsack

Instance: **Profits** $P = [p_1, \dots, p_n]$; **weights** $W = [w_1, \dots, w_n]$; and a **capacity**, M . These are all positive integers.

Feasible solution: An n -tuple $X = [x_1, \dots, x_n]$ where $\sum_{i=1}^n w_i x_i \leq M$. In the **0-1 Knapsack** problem (often denoted just as **Knapsack**), we require that $x_i \in \{0, 1\}$, $1 \leq i \leq n$.

In the **Rational Knapsack** problem, we require that $x_i \in \mathbb{Q}$ and $0 \leq x_i \leq 1$, $1 \leq i \leq n$.

Find: A feasible solution X that maximizes $\sum_{i=1}^n p_i x_i$.

Possible Greedy Strategies for Knapsack Problems

- 1 Consider the items in decreasing order of **profit** (i.e., the local evaluation criterion is p_i).
- 2 Consider the items in increasing order of **weight** (i.e., the local evaluation criterion is w_i).
- 3 Consider the items in decreasing order of **profit divided by weight** (i.e., the local evaluation criterion is p_i/w_i).

Does one of these strategies yield a **correct** greedy algorithm for the **0-1 Knapsack** or **Rational Knapsack** problem?

A Greedy Algorithm for Rational Knapsack

Algorithm: *GreedyRationalKnapsack*($P, W : \text{array}; M : \text{integer}$)

rename the items, sorting if necessary, so that $p_1/w_1 \geq \dots \geq p_n/w_n$

$X \leftarrow [0, \dots, 0]$

$i \leftarrow 1$

$CurW \leftarrow 0$

while ($CurW < M$) **and** ($i \leq n$)

do $\left\{ \begin{array}{ll} \text{if } CurW + w_i \leq M & \\ \text{then } \left\{ \begin{array}{l} x_i \leftarrow 1 \\ CurW \leftarrow CurW + w_i \\ i \leftarrow i + 1 \end{array} \right. & \\ \text{else } \left\{ \begin{array}{l} x_i \leftarrow (M - CurW)/w_i \\ CurW := M \end{array} \right. & \end{array} \right.$

return (X)

Coin Changing

Problem

Coin Changing

Instance: A list of **coin denominations**, d_1, d_2, \dots, d_n , and a positive integer T , which is called the **target sum**.

Find: An n -tuple of non-negative integers, say $A = [a_1, \dots, a_n]$, such that $T = \sum_{i=1}^n a_i d_i$ and such that $N = \sum_{i=1}^n a_i$ is minimized.

In the **Coin Changing** problem, a_i denotes the number of coins of denomination d_i that are used, for $i = 1, \dots, n$.

The total value of all the chosen coins must be exactly equal to T . We want to **minimize** the number of coins used, which is denoted by N .

A Greedy Algorithm for Coin Changing

Algorithm: *GreedyCoinChanging*($D : \text{array}; T : \text{integer}$)

comment: $D = [d_1, \dots, d_n]$

rename the coins, by sorting if necessary, so that $d_1 > \dots > d_n$

$N \leftarrow 0$

for $i \leftarrow 1$ **to** n

do
$$\begin{cases} a_i \leftarrow \left\lfloor \frac{T}{d_i} \right\rfloor \\ T \leftarrow T - a_i d_i \\ N \leftarrow N + a_i \end{cases}$$

if $T > 0$

then return (*fail*)

else return ($[a_1, \dots, a_n], N$)

The Stable Marriage Problem

Problem

Stable Marriage

Instance: A set of n **men**, say $M = [m_1, \dots, m_n]$, and a set of n **women**, $W = [w_1, \dots, w_n]$.

Each man m_i has a **preference ranking** of the n women, and each woman w_i has a preference ranking of the n men: $\text{pref}(m_i, j) = w_k$ if w_k is the j -th favourite woman of man m_i ; and $\text{pref}(w_i, j) = m_k$ if m_k is the j -th favourite man of woman w_i .

Find: A **matching** of the n men with the n women such that there **does not exist** a couple (m_i, w_j) who are **not** engaged to each other, but prefer each other to their existing matches. A matching with this property is called a **stable matching**.

Overview of the Gale-Shapley Algorithm

Men propose to women.

If a woman accepts a proposal, then the couple is **engaged**.

An unmatched woman **must accept** a proposal.

If an engaged woman receives a proposal from a man whom she prefers to her current match, then she **cancels** her existing engagement and she becomes engaged to the new proposer; her previous match is no longer engaged.

If an engaged woman receives a proposal from a man, but she prefers her current match, then the proposal is **rejected**.

Engaged women never become unengaged.

A man might make a number of proposals (up to n); the order of the proposals is determined by the man's preference list.

Gale-Shapley Algorithm

Algorithm: *Gale-Shapley*(M, W, pref)

$\text{Match} \leftarrow \emptyset$

while there exists an unengaged man m_i

do $\left\{ \begin{array}{l} \text{let } w_j \text{ be the next woman in } m_i \text{'s preference list} \\ \text{if } w_j \text{ is not engaged} \\ \quad \text{then } \text{Match} \leftarrow \text{Match} \cup \{m_i, w_j\} \\ \text{else} \left\{ \begin{array}{l} \text{suppose } \{m_k, w_j\} \in \text{Match} \\ \text{if } w_j \text{ prefers } m_i \text{ to } m_k \\ \quad \text{then } \left\{ \begin{array}{l} \text{Match} \leftarrow \text{Match} \setminus \{m_k, w_j\} \cup \{m_i, w_j\} \\ \text{comment: } m_k \text{ is now unengaged} \end{array} \right. \end{array} \right. \end{array} \right.$

return (Match)

Questions

How do we prove that the *Gale-Shapley* algorithm always **terminates**?

How many **iterations** does this algorithm require in the worst case?

How do we prove that this algorithm is **correct**, i.e., that it finds a stable matching?

Is there an efficient way to **identify** an unengaged man at any point in the algorithm? What **data structure** would be helpful in doing this?

What can we say about the **complexity** of the algorithm?

Table of Contents

5 Dynamic Programming Algorithms

- Fibonacci Numbers
- Design Strategy
- 0-1 Knapsack
- Coin Changing
- Longest Common Subsequence
- Minimum Length Triangulation

Computing Fibonacci Numbers Inefficiently

Algorithm: *BadFib*(n)

if $n = 0$ then $f \leftarrow 0$

else if $n = 1$ then $f \leftarrow 1$

else $\begin{cases} f_1 \leftarrow \text{BadFib}(n-1) \\ f_2 \leftarrow \text{BadFib}(n-2) \\ f \leftarrow f_1 + f_2 \end{cases}$

return (f);

Computing Fibonacci Numbers More Efficiently

Algorithm: *BetterFib*(n)

$f[0] \leftarrow 0$

$f[1] \leftarrow 1$

for $i \leftarrow 2$ **to** n

do $f[i] \leftarrow f[i - 1] + f[i - 2]$

return ($f[n]$)

Designing Dynamic Programming Algorithms for Optimization Problems

Optimal Structure

Examine the structure of an optimal solution to a problem instance I , and determine if an optimal solution for I can be expressed in terms of optimal solutions to certain **subproblems** of I .

Define Subproblems

Define a set of subproblems $\mathcal{S}(I)$ of the instance I , the solution of which enables the optimal solution of I to be computed. I will be the last or largest instance in the set $\mathcal{S}(I)$.

Designing Dynamic Programming Algorithms (cont.)

Recurrence Relation

Derive a **recurrence relation** on the optimal solutions to the instances in $\mathcal{S}(I)$. This recurrence relation should be completely specified in terms of optimal solutions to (smaller) instances in $\mathcal{S}(I)$ and/or base cases.

Compute Optimal Solutions

Compute the optimal solutions to all the instances in $\mathcal{S}(I)$. Compute these solutions using the recurrence relation in a **bottom-up** fashion, filling in a table of values containing these optimal solutions. Whenever a particular table entry is filled in using the recurrence relation, the optimal solutions of relevant subproblems can be looked up in the table (they have been computed already). The final table entry is the solution to I .

0-1 Knapsack

Problem

0-1 Knapsack

Instance: **Profits** $P = [p_1, \dots, p_n]$; **weights** $W = [w_1, \dots, w_n]$; and a **capacity**, M . These are all positive integers.

Feasible solution: An n -tuple $X = [x_1, \dots, x_n]$, where $x_i \in \{0, 1\}$ for $1 \leq i \leq n$, and $\sum_{i=1}^n w_i x_i \leq M$.

Find: A feasible solution X that maximizes $\sum_{i=1}^n p_i x_i$.

Let $P[i, m]$ denote the optimal solution to the subproblem consisting of the first i objects (having profits p_1, \dots, p_i and weights w_1, \dots, w_i , respectively) and capacity m .

A Dynamic Programming Algorithm for 0-1 Knapsack

Algorithm: *0-1Knapsack*($p_1, \dots, p_n, w_1, \dots, w_n, M$)

for $m \leftarrow 0$ **to** M

do $\left\{ \begin{array}{l} \text{if } m \geq w_1 \\ \quad \text{then } P[1, m] \leftarrow p_1 \\ \quad \text{else } P[1, m] \leftarrow 0 \end{array} \right.$

for $i \leftarrow 2$ **to** n

do $\left\{ \begin{array}{l} \text{for } m \leftarrow 0 \text{ to } M \\ \quad \text{do } \left\{ \begin{array}{l} \text{if } m < w_i \\ \quad \text{then } P[i, m] \leftarrow P[i - 1, m] \\ \quad \text{else } P[i, m] \leftarrow \max\{P[i - 1, m - w_i] + p_i, P[i - 1, m]\} \end{array} \right. \end{array} \right.$

return ($P[n, M]$);

Computing the Optimal Knapsack X

Algorithm: *0-1Knapsack*($p_1, \dots, p_n, w_1, \dots, w_n, M, P$)

$m \leftarrow M$

$p \leftarrow P[n, M]$

for $i \leftarrow n$ **downto** 2

do $\left\{ \begin{array}{l} \text{if } p = P[i - 1, m] \\ \text{then } x_i \leftarrow 0 \\ \text{else } \left\{ \begin{array}{l} x_i \leftarrow 1 \\ p \leftarrow p - p_i \\ m \leftarrow m - w_i \end{array} \right. \end{array} \right.$

if $p = 0$

then $x_1 \leftarrow 0$

else $x_1 \leftarrow 1$

return (X);

Coin Changing

Problem

Coin Changing

Instance: A list of **coin denominations**, $1 = d_1, d_2, \dots, d_n$, and a positive integer T , which is called the **target sum**.

Find: An n -tuple of non-negative integers, say $A = [a_1, \dots, a_n]$, such that $T = \sum_{i=1}^n a_i d_i$ and such that $N = \sum_{i=1}^n a_i$ is minimized.

Let $N[i, t]$ denote the optimal solution to the subproblem consisting of the first i coin denominations p_1, \dots, p_i and target sum t . Let $A[i, t]$ denote the number of coins of denomination d_i used in the optimal solution to this subproblem.

A Dynamic Programming Algorithm for Coin Changing

Algorithm: *Coin Changing*(d_1, \dots, d_n, T)

comment: $d_1 = 1$

for $t \leftarrow 0$ **to** T

do $\begin{cases} N[1, t] \leftarrow t \\ A[1, t] \leftarrow t \end{cases}$

for $i \leftarrow 2$ **to** n

do $\begin{cases} \text{for } t \leftarrow 0 \text{ to } T \\ \text{do } \begin{cases} N[i, t] \leftarrow N[i - 1, t] \\ A[i, t] \leftarrow 0 \\ \text{for } j \leftarrow 1 \text{ to } \lfloor (t/d_i) \rfloor \\ \text{do } \begin{cases} \text{if } j + N[i - 1, t - jd_i] < N[i, t] \\ \text{then } \begin{cases} N[i, t] \leftarrow j + N[i - 1, t - jd_i] \\ A[i, t] \leftarrow j \end{cases} \end{cases} \end{cases} \end{cases}$

return ($N[n, T]$);

Longest Common Subsequence

Problem

Longest Common Subsequence

Instance: Two sequences $X = (x_1, \dots, x_m)$ and $Y = (y_1, \dots, y_n)$ over some finite alphabet Γ .

Find: A maximum length sequence Z that is a subsequence of both X and Y .

$Z = (z_1, \dots, z_\ell)$ is a **subsequence** of X if there exist indices $1 \leq i_1 < \dots < i_\ell \leq m$ such that $z_j = x_{i_j}$, $1 \leq j \leq \ell$.

Similarly, Z is a subsequence of Y if there exist (possibly different) indices $1 \leq h_1 < \dots < h_\ell \leq n$ such that $z_j = y_{h_j}$, $1 \leq j \leq \ell$.

Computing the Length of the LCS of X and Y

Algorithm: *LCS1*($X = (x_1, \dots, x_m), Y = (y_1, \dots, y_n)$)

for $i \leftarrow 0$ **to** m

do $c[i, 0] \leftarrow 0$

for $j \leftarrow 0$ **to** n

do $c[0, j] \leftarrow 0$

for $i \leftarrow 1$ **to** m

do $\left\{ \begin{array}{l} \text{for } j \leftarrow 1 \text{ to } n \end{array} \right.$

do $\left\{ \begin{array}{l} \text{if } x_i = y_j \end{array} \right.$

then $c[i, j] \leftarrow c[i - 1, j - 1] + 1$

else $c[i, j] \leftarrow \max\{c[i, j - 1], c[i - 1, j]\}$

return ($c[m, n]$);

Finding the LCS of X and Y

Algorithm: $LCS2(X = (x_1, \dots, x_m), Y = (y_1, \dots, y_n))$

for $i \leftarrow 0$ **to** m **do** $c[i, 0] \leftarrow 0$

for $j \leftarrow 0$ **to** n **do** $c[0, j] \leftarrow 0$

for $i \leftarrow 1$ **to** m

do { **for** $j \leftarrow 1$ **to** n

do { **if** $x_i = y_j$

then $\begin{cases} c[i, j] \leftarrow c[i - 1, j - 1] + 1 \\ \pi[i, j] \leftarrow \text{UL} \end{cases}$

else if $c[i, j - 1] > c[i - 1, j]$

then $\begin{cases} c[i, j] \leftarrow c[i, j - 1] \\ \pi[i, j] \leftarrow \text{L} \end{cases}$

else $\begin{cases} c[i, j] \leftarrow c[i - 1, j] \\ \pi[i, j] \leftarrow \text{U} \end{cases}$

return $(c, \pi);$

Finding the LCS

Algorithm: *FindLCS*(c, π, v)

$seq \leftarrow ()$

$i \leftarrow m$

$j \leftarrow n$

while $\min\{i, j\} > 0$

do $\left\{ \begin{array}{l} \text{if } \pi[i, j] = \text{UL} \\ \quad \text{then } \left\{ \begin{array}{l} seq \leftarrow x_i \parallel seq \\ i \leftarrow i - 1 \\ j \leftarrow j - 1 \end{array} \right. \\ \text{else if } \pi[i, j] = \text{L} \text{ then } j \leftarrow j - 1 \\ \text{else } i \leftarrow i - 1 \end{array} \right.$

return (seq)

Minimum Length Triangulation

Problem

Minimum Length Triangulation v1

Instance: n points q_1, \dots, q_n in the Euclidean plane that form a convex n -gon P .

Find: A triangulation of P such that the sum S_c of the lengths of the $n - 3$ chords is minimized.

Problem

Minimum Length Triangulation v2

Instance: n points q_1, \dots, q_n in the Euclidean plane that form a convex n -gon P .

Find: A triangulation of P such that the sum S_p of the perimeters of the $n - 2$ triangles is minimized.

Let L denote the perimeter of P . Then we have that $S_p = L + 2S_c$.

Hence the two versions have the **same optimal solutions**.

Problem Decomposition

We consider version 2 of the problem.

The edge $q_n q_1$ is in a triangle with a third vertex q_k , where $k \in \{2, \dots, n-1\}$.

For a given k , we have:

- ① the triangle $q_1 q_k q_n$,
- ② the polygon with vertices q_1, \dots, q_k ,
- ③ the polygon with vertices q_k, \dots, q_n .

The optimal solution will consist of optimal solutions to the **two subproblems** in 2 and 3, along with the **triangle** in 1.

Recurrence Relation

For $1 \leq i < j \leq n$, let $S[i, j]$ denote the optimal solution to the subproblem consisting of the polygon having vertices q_i, \dots, q_j .

Let $\Delta(q_i, q_k, q_j)$ denote the perimeter of the triangle having vertices q_i, q_k, q_j .

Then we have the recurrence relation

$$S[i, j] = \min\{\Delta(q_i, q_k, q_j) + S[i, k] + S[k, j] : i < k < j\}.$$

The base cases are given by

$$S[i, i + 1] = 0$$

for all i .

We compute all $S[i, j]$ with $j - i = c$, for $c = 2, 3, \dots, n - 1$.

Table of Contents

6 Graph Algorithms

- Graphs
- Breadth-first Search
- Depth-first Search
- Topological Ordering
- Strongly Connected Components
- Minimum Spanning Trees
- Single Source Shortest Paths
- All-Pairs Shortest Paths

Graphs and Digraphs

A **graph** is a pair $G = (V, E)$. V is a set whose elements are called **vertices** and E is a set whose elements are called **edges**. Each edge joins two distinct vertices. An edge can be represented as a set of two vertices, e.g., $\{u, v\}$, where $u \neq v$. We may also write this edge as uv or vu .

We often denote the number of vertices by n and the number of edges by m . Clearly $m \leq \binom{n}{2}$.

A **directed graph** or **digraph** is also a pair $G = (V, E)$. The elements of E are called **directed edges** or **arcs** in a digraph. Each arc joins two vertices, and an arc can be represented as a ordered pair, e.g., (u, v) . The arc (u, v) is directed from u (the **tail**) to v (the **head**), and we allow $u = v$.

If we denote the number of vertices by n and the number of arcs by m , then $m \leq n^2$.

Data Structures for Graphs: Adjacency Matrices

There are two main data structures to represent graphs: an **adjacency matrix** and a set of **adjacency lists**.

Let $G = (V, E)$ be a graph with $|V| = n$ and $|E| = m$. The **adjacency matrix** of G is an n by n matrix $A = (a_{u,v})$, which is indexed by V , such that

$$a_{u,v} = \begin{cases} 1 & \text{if } \{u, v\} \in E \\ 0 & \text{otherwise.} \end{cases}$$

There are exactly $2m$ entries of A equal to 1.

If G is a digraph, then

$$a_{u,v} = \begin{cases} 1 & \text{if } (u, v) \in E \\ 0 & \text{otherwise.} \end{cases}$$

For a digraph, there are exactly m entries of A equal to 1.

Data Structures for Graphs: Adjacency Lists

Let $G = (V, E)$ be a graph with $|V| = n$ and $|E| = m$.

An **adjacency list representation** of G consists of n linked lists.

For every $u \in V$, there is a linked list (called an **adjacency list**) which is named $Adj[u]$.

For every $v \in V$ such that $uv \in E$, there is a node in $Adj[u]$ labelled v .
(This definition is used for both directed and undirected graphs.)

In an undirected graph, every edge uv corresponds to nodes in **two** adjacency lists: there is a node v in $Adj[u]$ and a node u in $Adj[v]$.

In a directed graph, every edge corresponds to a node in only **one** adjacency list.

Breadth-first Search of an Undirected Graph

A **breadth-first search** of an undirected graph begins at a specified vertex s .

The search “spreads out” from s , proceeding in **layers**.

First, all the neighbours of s are **explored**.

Next, the neighbours of those neighbours are explored.

This process continues until all vertices have been explored.

A **queue** is used to keep track of the vertices to be explored.

Breadth-first Search

Algorithm: *BFS*(G, s)

for each $v \in V(G)$

do $\begin{cases} colour[v] \leftarrow \text{white} \\ \pi[v] \leftarrow \emptyset \end{cases}$

$colour[s] \leftarrow \text{gray}$

InitializeQueue(Q)

Enqueue(Q, s)

while $Q \neq \emptyset$

do $\begin{cases} u \leftarrow \text{Dequeue}(Q) \\ \text{for each } v \in Adj[u] \\ \text{do } \begin{cases} \text{if } colour[v] = \text{white} \\ \text{do } \begin{cases} \text{if } colour[v] = \text{gray} \\ \text{then } \begin{cases} \pi[v] \leftarrow u \\ \text{Enqueue}(Q, v) \end{cases} \end{cases} \end{cases} \\ colour[u] \leftarrow \text{black} \end{cases}$

Properties of Breadth-first Search

A vertex is **white** if it is **undiscovered**.

A vertex is **gray** if it has been **discovered**, but we are still processing its adjacent vertices.

A vertex becomes **black** when all the adjacent vertices have been processed.

If G is **connected**, then every vertex eventually is coloured black.

When we explore an edge $\{u, v\}$ starting from u :

- if v is **white**, then uv is a **tree edge** and $\pi[v] = u$ is the **predecessor** of v in the **BFS tree**
- otherwise, uv is a **cross edge**.

The BFS tree consists of all the tree edges.

Every vertex $v \neq s$ has a unique predecessor $\pi[v]$ in the BFS tree.

Shortest Paths via Breadth-first Search

Algorithm: *BFS*(G, s)

for each $v \in V(G)$ do $\begin{cases} colour[v] \leftarrow \text{white} \\ \pi[v] \leftarrow \emptyset \end{cases}$

$colour[s] \leftarrow \text{gray}$

$dist[s] \leftarrow 0$

InitializeQueue(Q)

Enqueue(Q, s)

while $Q \neq \emptyset$

do $\begin{cases} u \leftarrow \text{Dequeue}(Q) \\ \text{for each } v \in Adj[u] \\ \text{do } \begin{cases} \text{if } colour[v] = \text{white} \text{ then } \begin{cases} colour[v] = \text{gray} \\ \pi[v] \leftarrow u \\ \text{Enqueue}(Q, v) \\ dist[v] \leftarrow dist[u] + 1 \end{cases} \\ colour[u] \leftarrow \text{black} \end{cases} \end{cases}$

Distances in Breadth-first Search

If $\{u, v\}$ is **any edge**, then $|dist[u] - dist[v]| \leq 1$.

If uv is a **tree edge**, then $dist[v] = dist[u] + 1$.

$dist[u]$ is the length of the **shortest path** from s to u .

This is also called the **distance** from s to u .

Bipartite Graphs and Breadth-first Search

A graph is **bipartite** if the vertex set can be partitioned as $V = X \cup Y$, in such a way that all edges have one endpoint in X and one endpoint in Y .

A graph is bipartite if and only if it does not contain an **odd cycle**.

BFS can be used to test if a graph is bipartite:

- if we encounter an edge $\{u, v\}$ with $dist[u] = dist[v]$, then G is not bipartite, whereas
- if no such edge is found, then define $X = \{u : dist[u] \text{ is even}\}$ and $Y = \{u : dist[u] \text{ is odd}\}$; then X, Y forms a bipartition.

Depth-first Search of a Directed Graph

A **depth-first search** uses a **stack** (or **recursion**) instead of a queue.

We define predecessors and colour vertices as in BFS.

It is also useful to specify a **discovery time** $d[v]$ and a **finishing time** $f[v]$ for every vertex v .

We increment a **time counter** every time a value $d[v]$ or $f[v]$ is assigned.

We eventually visit all the vertices, and the algorithm constructs a **depth-first forest**.

Depth-first Search

Algorithm: *DFS*(G)
 for each $v \in V(G)$
 do $\begin{cases} \text{colour}[v] \leftarrow \text{white} \\ \pi[v] \leftarrow \emptyset \end{cases}$
 $time \leftarrow 0$
 for each $v \in V(G)$
 do $\begin{cases} \text{if } \text{colour}[v] = \text{white} \\ \quad \text{then } \text{DFSvisit}(v) \end{cases}$

Depth-first Search (cont.)

Algorithm: *DFSvisit*(v)

$colour[v] \leftarrow \text{gray}$

$time \leftarrow time + 1$

$d[v] \leftarrow time$

comment: $d[v]$ is the discovery time for vertex v

for each $w \in Adj[v]$

do $\left\{ \begin{array}{l} \text{if } colour[w] = \text{white} \\ \text{then } \left\{ \begin{array}{l} \pi[w] \leftarrow v \\ \text{DFSvisit}(w) \end{array} \right. \end{array} \right.$

$colour[v] \leftarrow \text{black}$

$time \leftarrow time + 1$

$f[v] \leftarrow time$

comment: $f[v]$ is the finishing time for vertex v

Classification of Edges in Depth-first Search

- uv is a **tree edge** if $u = \pi[v]$
- uv is a **forward edge** if it is not a tree edge, and v is a descendant of u in a tree in the depth-first forest
- uv is a **back edge** if u is a descendant of v in a tree in the depth-first forest
- any other edge is a **cross edge**.

Properties of Edges in Depth-first Search

In the following table, we indicate the colour of a vertex v when an edge uv is discovered, and the relation between the start and finishing times of u and v , for each possible type of edge uv .

edge type	colour of v	discovery/finish times
tree	white	$d[u] < d[v] < f[v] < f[u]$
forward	black	$d[u] < d[v] < f[v] < f[u]$
back	gray	$d[v] < d[u] < f[u] < f[v]$
cross	black	$d[v] < f[v] < d[u] < f[u]$

Observe that two intervals $(d[u], f[u])$ and $(d[v], f[v])$ never **overlap**. Two intervals are either **disjoint** or **nested**. This is sometimes called the **parenthesis theorem**.

Topological Orderings and DAGs

A directed graph G is a **directed acyclic graph**, or **DAG**, if G contains no directed cycle.

A directed graph $G = (V, E)$ has a **topological ordering**, or **topological sort**, if there is a linear ordering $<$ of all the vertices in V such that $u < v$ whenever $uv \in E$.

Some interesting/useful facts:

- A DAG contains a vertex of indegree 0.
- A directed graph G has a topological ordering if and only if it is a DAG.
- A directed graph G is a DAG if and only if a DFS of G has no back edges.
- If uv is an edge in a DAG, then a DFS of G has $f[v] < f[u]$.

Topological Ordering via Depth-first Search

Algorithm: *DFS*(G)

InitializeStack(S)

$DAG \leftarrow true$

for each $v \in V(G)$

do $\begin{cases} colour[v] \leftarrow \text{white} \\ \pi[v] \leftarrow \emptyset \end{cases}$

$time \leftarrow 0$

for each $v \in V(G)$

do $\begin{cases} \text{if } colour[v] = \text{white} \\ \text{then } DFSvisit(v) \end{cases}$

if DAG **then return** (S) **else return** (DAG)

Topological Ordering via Depth-first Search (cont.)

Algorithm: *DFSvisit*(v)

$colour[v] \leftarrow \text{gray}$

$time \leftarrow time + 1$

$d[v] \leftarrow time$

comment: $d[v]$ is the discovery time for vertex v

for each $w \in Adj[v]$

do $\left\{ \begin{array}{l} \text{if } colour[w] = \text{white} \\ \quad \text{then } \left\{ \begin{array}{l} \pi[w] \leftarrow v \\ \text{DFSvisit}(w) \end{array} \right. \\ \text{if } colour[w] = \text{gray} \text{ then } DAG \leftarrow false \end{array} \right.$

$colour[v] \leftarrow \text{black}$

$\text{Push}(S, v)$

$time \leftarrow time + 1$

$f[v] \leftarrow time$

comment: $f[v]$ is the finishing time for vertex v

Strongly Connected Components of a Digraph G

For two vertices x and y of G , define $x \sim y$ if $x = y$; or if $x \neq y$ and there exist directed paths from x to y **and** from y to x .

The relation \sim is an **equivalence relation**.

The **strongly connected components** of G are the equivalence classes of vertices defined by the relation \sim .

The **component graph** of G is a directed graph whose vertices are the strongly connected components of G . There is an arc from C_i to C_j if and only if there is an arc in G from some vertex of C_i to some vertex of C_j .

For a strongly connected component C , define $f[C] = \max\{f[v] : v \in C\}$ and $d[C] = \min\{d[v] : v \in C\}$.

Some interesting/useful facts:

- The component graph of G is a DAG.
- If C_i, C_j are strongly connected components, and there is an arc from C_i to C_j in the component graph, then $f[C_i] > f[C_j]$.

An Algorithm to Find the Strongly Connected Components

- step 1** Perform a depth-first search of G , recording the finishing times $f[v]$ for all vertices v .
- step 2** Construct a directed graph H from G by **reversing** the direction of all edges in G .
- step 3** Perform a depth-first search of H , considering the vertices in **decreasing** order of the values $f[v]$ computed in step 1.
- step 4** The strongly connected components of G are the trees in the depth-first forest constructed in step 3.

Depth-first Search of H

Assume that $f[v_{i_1}] > f[v_{i_2}] > \dots > f[v_{i_n}]$.

Algorithm: $DFS(H)$

for $j \leftarrow 1$ **to** n

do $colour[v_{i_j}] \leftarrow$ **white**

$scc \leftarrow 0$

for $j \leftarrow 1$ **to** n

do $\left\{ \begin{array}{l} \text{if } colour[v_{i_j}] = \text{white} \\ \text{then } \left\{ \begin{array}{l} scc \leftarrow scc + 1 \\ DFSvisit(H, v_{i_j}, scc) \end{array} \right. \end{array} \right.$

return $(comp)$

comment: $comp[v]$ is the strongly connected component containing v

DFSvisit for H

Algorithm: *DFSvisit*(H, v, scc)

$colour[v] \leftarrow \text{gray}$

$comp[v] \leftarrow scc$

for each $w \in Adj[v]$

do $\left\{ \begin{array}{l} \text{if } colour[w] = \text{white} \\ \quad \text{then } DFSvisit(H, w, scc) \end{array} \right.$

$colour[v] \leftarrow \text{black}$

Minimum Spanning Trees

A **spanning tree** in a connected, undirected graph $G = (V, E)$ is a subgraph T that is a tree which contains every vertex of V .

T is a spanning tree of G if and only if T is an acyclic subgraph of G that has $n - 1$ edges (where $n = |V|$).

Problem

Minimum Spanning Tree

Instance: A connected, undirected graph $G = (V, E)$ and a **weight function** $w : E \rightarrow \mathbb{R}$.

Find: A spanning tree T of G such that

$$\sum_{e \in T} w(e)$$

is minimized (this is called a **minimum spanning tree**, or **MST**).

Kruskal's Algorithm

Assume that $w(e_1) \leq w(e_2) \leq \dots \leq w(e_m)$, where $m = |E|$.

Algorithm: *Kruskal*(G, w)

$A \leftarrow \emptyset$

for $j \leftarrow 1$ **to** m

do $\left\{ \begin{array}{l} \text{if } A \cup \{e_j\} \text{ does not contain a cycle} \\ \text{then } A \leftarrow A \cup \{e_j\} \end{array} \right.$

return (A)

Prim's Algorithm (idea)

We initially choose an arbitrary vertex u_0 and define $A = \{e\}$, where e is the **minimum weight** edge incident with u_0 .

A is always a **single tree**, and at each step we select the minimum weight edge that joins a vertex in V_A to a vertex not in V_A .

Remark: V_A denotes the set of vertices in the tree A .

For a vertex $v \notin V_A$, define

$N[v] = u$, where $\{u, v\}$ is a minimum weight edge such that $u \in V_A$

$W[v] = w(N[v], v)$.

Assume $w(u, v) = \infty$ if $\{u, v\} \notin E$.

Prim's Algorithm

Algorithm: *Prim*(G, w)

$A \leftarrow \emptyset$

$V_A \leftarrow \{u_0\}$, where u_0 is arbitrary

for all $v \in V \setminus \{u_0\}$

do $\begin{cases} W[v] \leftarrow w(u_0, v) \\ N[v] \leftarrow u_0 \end{cases}$

while $|A| < n - 1$

do $\begin{cases} \text{choose } v \in V \setminus V_A \text{ such that } W[v] \text{ is minimized} \\ V_A \leftarrow V_A \cup \{v\} \\ u \leftarrow N[v] \\ A \leftarrow A \cup \{uv\} \\ \text{for all } v' \in V \setminus V_A \\ \quad \text{do } \begin{cases} \text{if } w(v, v') < W[v'] \\ \quad \text{then } \begin{cases} W[v'] \leftarrow w(v, v') \\ N[v'] \leftarrow v \end{cases} \end{cases} \end{cases}$

return (A)

Definitions

Let $G = (V, E)$ be a graph. A **cut** is a partition of V into two non-empty (disjoint) sets, i.e., a pair $(S, V \setminus S)$, where $S \subseteq V$ and $1 \leq |S| \leq n - 1$.

Let $(S, V \setminus S)$ be a cut in a graph $G = (V, E)$. An edge $e \in E$ is a **crossing edge** with respect to the cut $(S, V \setminus S)$ if e has one endpoint in S and one endpoint in $V \setminus S$.

Let $A \subseteq E$. A cut $(S, V \setminus S)$ **respects** the set of edges A provided that no edge in A is a crossing edge.

A General Greedy Algorithm to Find an MST

Algorithm: *GreedyMST*(G, w)

$A \leftarrow \emptyset$

while $|A| < n - 1$

do $\begin{cases} \text{let } (S, V \setminus S) \text{ be a cut that respects } A \\ \text{let } e \text{ be a minimum weight crossing edge} \\ A \leftarrow A \cup \{e\} \end{cases}$

return (A)

Single Source Shortest Paths

Problem

Single Source Shortest Paths

Instance: A directed graph $G = (V, E)$, a non-negative **weight function** $w : E \rightarrow \mathbb{R}^+ \cup \{0\}$, and a **source vertex** $u_0 \in V$.

Find: For every vertex $v \in V$, a directed path P from u_0 to v such that

$$w(P) = \sum_{e \in P} w(e)$$

is minimized.

The term **shortest path** really means **minimum weight path**.

We are asked to find n different shortest paths, one for each vertex $v \in V$.

If all edges have weight 1, we can just use **BFS** to solve this problem.

Dijkstra's Algorithm (Main Ideas)

S is a subset of vertices such that the shortest paths from u_0 to all vertices in S are known; initially, $S = \{u_0\}$.

For all vertices $v \in S$, $D[v]$ is the weight of the shortest path P_v from u_0 to v , and all vertices on P_v are in the set S .

For all vertices $v \notin S$, $D[v]$ is the weight of the shortest path P_v from u_0 to v in which all interior vertices are in S .

For $v \neq u_0$, $\pi[v]$ is the **predecessor** of v on the path P_v .

At each stage of the algorithm, we choose $v \in V \setminus S$ so that $D[v]$ is minimized, and then we add v to S .

Then the arrays D and π are updated appropriately.

Dijkstra's Algorithm

Algorithm: *Dijkstra*(G, w, u_0)

$S \leftarrow \{u_0\}$

$D[u_0] \leftarrow 0$

for all $v \in V \setminus \{u_0\}$

do $\begin{cases} D[v] \leftarrow w(u_0, v) \\ \pi[v] \leftarrow u_0 \end{cases}$

while $|S| < n$

do $\begin{cases} \text{choose } v \in V \setminus S \text{ such that } D[v] \text{ is minimized} \\ S \leftarrow S \cup \{v\} \\ \text{for all } v' \in V \setminus S \\ \quad \text{do } \begin{cases} \text{if } D[v] + w(v, v') < D[v'] \\ \quad \text{then } \begin{cases} D[v'] \leftarrow D[v] + w(v, v') \\ \pi[v'] \leftarrow v \end{cases} \end{cases} \end{cases}$

return (D, π)

Finding the Shortest Paths

Algorithm: *FindPath*(u_0, π, v)

$path \leftarrow v$

$u \leftarrow v$

while $u \neq u_0$

do $\begin{cases} u \leftarrow \pi[u] \\ path \leftarrow u \parallel path \end{cases}$

return ($path$)

Shortest Paths in a DAG

If G is a DAG, we perform a topological ordering of the vertices. Suppose the resulting ordering is v_1, \dots, v_n . Then we find all the shortest paths in G with source v_1 .

Note: This algorithm is correct even if there are **negative-weight edges**.

Algorithm: *DAG Shortest paths*(G, w, v_1)

for $j \leftarrow 1$ **to** n

do $\begin{cases} D[v_1] \leftarrow \infty \\ \pi[v_j] \leftarrow \text{undefined} \end{cases}$

$D[v_1] \leftarrow 0$

for $j \leftarrow 1$ **to** $n - 1$

do $\begin{cases} \text{for all } v' \in \text{Adj}[v_j] \\ \text{do } \begin{cases} \text{if } D[v_j] + w(v_j, v') < D[v'] \\ \text{then } \begin{cases} D[v'] \leftarrow D[v_j] + w(v_j, v') \\ \pi[v'] \leftarrow v_j \end{cases} \end{cases} \end{cases}$

return (D, π)

All-Pairs Shortest Paths

Problem

All-Pairs Shortest Paths

Instance: A directed graph $G = (V, E)$, and a **weight matrix** W , where $W[i, j]$ denotes the weight of edge ij , for all $i, j \in V$, $i \neq j$.

Find: For all pairs of vertices $u, v \in V$, $u \neq v$, a directed path P from u to v such that

$$w(P) = \sum_{ij \in P} W[i, j]$$

is minimized.

We allow edges to have negative weights, but we assume there are no negative-weight directed cycles in G .

First Solution

Algorithm: *SlowAllPairsShortestPath*(W)

$L_1 \leftarrow W$

for $m \leftarrow 2$ **to** $n - 1$

do $\left\{ \begin{array}{l} \textbf{for } i \leftarrow 1 \textbf{ to } n \\ \textbf{do } \left\{ \begin{array}{l} \textbf{for } j \leftarrow 1 \textbf{ to } n \\ \textbf{do } \left\{ \begin{array}{l} \ell \leftarrow \infty \\ \textbf{for } k \leftarrow 1 \textbf{ to } n \\ \textbf{do } \ell \leftarrow \min\{\ell, L_{m-1}[i, k] + W[k, j]\} \\ L_m[i, j] \leftarrow \ell \end{array} \right. \end{array} \right. \end{array} \right.$

return (L_{n-1})

Second Solution

Algorithm: *FasterAllPairsShortestPath*(W)

$L_1 \leftarrow W$

$m \leftarrow 1$

while $m < n - 1$

do {
 for $i \leftarrow 1$ **to** n
do {
 for $j \leftarrow 1$ **to** n
do {
 $\ell \leftarrow \infty$
for $k \leftarrow 1$ **to** n
do $\ell \leftarrow \min\{\ell, L_m[i, k] + L_m[k, j]\}$
 $L_{2m}[i, j] \leftarrow \ell$
 }
 }
 }
 $m \leftarrow 2m$

return (L_m)

Third Solution

Algorithm: *FloydWarshall*(W)

$D_0 \leftarrow W$

for $m \leftarrow 1$ **to** n

do $\left\{ \begin{array}{l} \text{for } i \leftarrow 1 \text{ to } n \\ \quad \text{do } \left\{ \begin{array}{l} \text{for } j \leftarrow 1 \text{ to } n \text{ do} \\ \quad D_m[i, j] \leftarrow \min\{D_{m-1}[i, j], D_{m-1}[i, m] + D_{m-1}[m, j]\} \end{array} \right. \end{array} \right.$

return (D_n)

Table of Contents

7 Intractability and Undecidability

- Decision Problems
- The Complexity Class P
- Decision, Optimal Value and Optimization Problems
- The Complexity Class NP
- Reductions
- NP-completeness and NP-complete Problems
- Undecidability

Decision Problems

Decision Problem: Given a problem instance I , answer a certain question “yes” or “no”.

Problem Instance: Input for the specified problem.

Problem Solution: Correct answer (“yes” or “no”) for the specified problem instance. I is a **yes-instance** if the correct answer for the instance I is “yes”. I is a **no-instance** if the correct answer for the instance I is “no”.

Size of a problem instance: $Size(I)$ is the number of bits required to specify (or encode) the instance I .

The Complexity Class P

Algorithm Solving a Decision Problem: An algorithm A is said to **solve** a decision problem Π provided that A finds the correct answer (“yes” or “no”) for every instance I of Π in finite time.

Polynomial-time Algorithm: An algorithm A for a decision problem Π is said to be a **polynomial-time algorithm** provided that the complexity of A is $O(n^k)$, where k is a positive integer and $n = \text{Size}(I)$.

The Complexity Class P denotes the set of all decision problems that have polynomial-time algorithms solving them. We write $\Pi \in P$ if the decision problem Π is in the complexity class P .

Cycles in Graphs

Problem

Cycle

Instance: *An undirected graph $G = (V, E)$.*

Question: *Does G contain a cycle?*

Problem

Hamiltonian Cycle

Instance: *An undirected graph $G = (V, E)$.*

Question: *Does G contain a hamiltonian cycle?*

A **hamiltonian cycle** is a cycle that passes through every vertex in V exactly once.

Knapsack Problems

Problem

0-1 Knapsack-Dec

Instance: a list of **profits**, $P = [p_1, \dots, p_n]$; a list of **weights**, $W = [w_1, \dots, w_n]$; a **capacity**, M ; and a **target profit**, T .

Question: Is there an n -tuple $[x_1, x_2, \dots, x_n] \in \{0, 1\}^n$ such that $\sum w_i x_i \leq M$ and $\sum p_i x_i \geq T$?

Problem

Rational Knapsack-Dec

Instance: a list of **profits**, $P = [p_1, \dots, p_n]$; a list of **weights**, $W = [w_1, \dots, w_n]$; a **capacity**, M ; and a **target profit**, T .

Question: Is there an n -tuple $[x_1, x_2, \dots, x_n] \in [0, 1]^n$ such that $\sum w_i x_i \leq M$ and $\sum p_i x_i \geq T$?

Polynomial-time Turing Reductions

Suppose Π_1 and Π_2 are problems (not necessarily decision problems). A (hypothetical) algorithm A_2 to solve Π_2 is called an **oracle** for Π_2 .

Suppose that A is an algorithm that solves Π_1 , assuming the existence of an oracle A_2 for Π_2 . (A_2 is used as a subroutine within the algorithm A .)

Then we say that A is a **Turing reduction** from Π_1 to Π_2 , denoted $\Pi_1 \leq^T \Pi_2$.

A Turing reduction A is a **polynomial-time Turing reduction** if the running time of A is polynomial, under the assumption that the oracle A_2 has **unit cost** running time.

If there is a polynomial-time Turing reduction from Π_1 to Π_2 , we write $\Pi_1 \leq_P^T \Pi_2$.

Informally: **Existence of a polynomial-time Turing reduction means that if we can solve Π_2 in polynomial time, then we can solve Π_1 in polynomial time.**

Travelling Salesperson Problems

Problem

TSP-Optimization

Instance: A graph G and edge weights $w : E \rightarrow \mathbb{Z}^+$.

Find: A hamiltonian cycle H in G such that $w(H) = \sum_{e \in H} w(e)$ is minimized.

Problem

TSP-Optimal Value

Instance: A graph G and edge weights $w : E \rightarrow \mathbb{Z}^+$.

Find: The minimum T such that there exists a hamiltonian cycle H in G with $w(H) = T$.

Problem

TSP-Decision

Instance: A graph G , edge weights $w : E \rightarrow \mathbb{Z}^+$, and a target T .

Question: Does there exist a hamiltonian cycle H in G with $w(H) \leq T$?

TSP-Optimal Value \leq_P^T TSP-Dec

Algorithm: *TSP-OptimalValue-Solver*(G, w)

external *TSP-Dec-Solver*

$hi \leftarrow \sum_{e \in E} w(e)$

$lo \leftarrow 0$

if not *TSP-Dec-Solver*(G, w, hi) **then return** (∞)

while $hi > lo$

do $\left\{ \begin{array}{l} mid \leftarrow \lfloor \frac{hi+lo}{2} \rfloor \\ \text{if } \textit{TSP-Dec-Solver}(G, w, mid) \\ \quad \text{then } hi \leftarrow mid \\ \quad \text{else } lo \leftarrow mid + 1 \end{array} \right.$

return (hi)

TSP-Optimization \leq_P^T TSP-Dec

Algorithm: *TSP-Optimization-Solver*($G = (V, E), w$)
external *TSP-OptimalValue-Solver*, *TSP-Dec-Solver*
 $T^* \leftarrow \textit{TSP-OptimalValue-Solver}(G, w)$
if $T^* = \infty$ **then return** (“no hamiltonian cycle exists”)
 $w_0 \leftarrow w$
 $H \leftarrow \emptyset$
for all $e \in E$
 do $\begin{cases} w_0[e] \leftarrow \infty \\ \text{if not } \textit{TSP-Dec-Solver}(G, w_0, T^*) \\ \quad \text{then } \begin{cases} w_0[e] \leftarrow w[e] \\ H \leftarrow H \cup \{e\} \end{cases} \end{cases}$
return (H)

Certificates

Certificate: Informally, a certificate for a yes-instance I is some “extra information” C which makes it easy to **verify** that I is a yes-instance.

Certificate Verification Algorithm: Suppose that Ver is an algorithm that verifies certificates for yes-instances. Then $Ver(I, C)$ outputs “yes” if I is a yes-instance and C is a valid certificate for I . If $Ver(I, C)$ outputs “no”, then either I is a no-instance, or I is a yes-instance and C is an invalid certificate.

Polynomial-time Certificate Verification Algorithm: A certificate verification algorithm Ver is a polynomial-time certificate verification algorithm if the complexity of Ver is $O(n^k)$, where k is a positive integer and $n = Size(I)$.

The Complexity Class NP

Certificate Verification Algorithm for a Decision Problem: A

certificate verification algorithm Ver is said to **solve** a decision problem Π provided that

- **for every** yes-instance I , **there exists** a certificate C such that $Ver(I, C)$ outputs “yes”.
- **for every** no-instance I and **for every** certificate C , $Ver(I, C)$ outputs “no”.

The Complexity Class NP denotes the set of all decision problems that have polynomial-time certificate verification algorithms solving them. We write $\Pi \in NP$ if the decision problem Π is in the complexity class NP .

Finding Certificates vs Verifying Certificates: It is **not required** to be able to **find** a certificate C for a yes-instance in polynomial time in order to say that a decision problem $\Pi \in NP$.

Certificate Verification Algorithm for Hamiltonian Cycle

A certificate consists of an n -tuple, $X = [x_1, \dots, x_n]$, that might be a hamiltonian cycle for a given graph $G = (V, E)$ (where $n = |V|$).

Algorithm: *Hamiltonian Cycle Certificate Verification*(G, X)

$flag \leftarrow \mathbf{true}$

$Used \leftarrow \{x_1\}$

$j \leftarrow 2$

while ($j \leq n$) **and** $flag$

do $\begin{cases} flag \leftarrow (x_j \notin Used) \mathbf{and} (\{x_{j-1}, x_j\} \in E) \\ \mathbf{if} (j = n) \mathbf{then} flag \leftarrow flag \mathbf{and} (\{x_n, x_1\} \in E) \\ Used \leftarrow Used \cup \{x_j\} \end{cases}$

return ($flag$)

Polynomial-time Reductions

For a decision problem Π , let $\mathcal{I}(\Pi)$ denote the set of all instances of Π . Let $\mathcal{I}_{\text{yes}}(\Pi)$ and $\mathcal{I}_{\text{no}}(\Pi)$ denote the set of all yes-instances and no-instances (respectively) of Π .

Suppose that Π_1 and Π_2 are decision problems. We say that there is a **polynomial-time reduction** (AKA **polynomial transformation**) from Π_1 to Π_2 (denoted $\Pi_1 \leq_P \Pi_2$) if there exists a function $f : \mathcal{I}(\Pi_1) \rightarrow \mathcal{I}(\Pi_2)$ such that the following properties are satisfied:

- $f(I)$ is computable in polynomial time (as a function of $\text{size}(I)$, where $I \in \mathcal{I}(\Pi_1)$)
- if $I \in \mathcal{I}_{\text{yes}}(\Pi_1)$, then $f(I) \in \mathcal{I}_{\text{yes}}(\Pi_2)$
- if $I \in \mathcal{I}_{\text{no}}(\Pi_1)$, then $f(I) \in \mathcal{I}_{\text{no}}(\Pi_2)$

Two Graph Theory Decision Problems

Problem

Clique

Instance: An undirected graph $G = (V, E)$ and an integer k , where $1 \leq k \leq |V|$.

Question: Does G contain a clique of size $\geq k$? (A **clique** is a subset of vertices $W \subseteq V$ such that $uv \in E$ for all $u, v \in W$, $u \neq v$.)

Problem

Vertex Cover

Instance: An undirected graph $G = (V, E)$ and an integer k , where $1 \leq k \leq |V|$.

Question: Does G contain a vertex cover of size $\leq k$? (A **vertex cover** is a subset of vertices $W \subseteq V$ such that $\{u, v\} \cap W \neq \emptyset$ for all edges $uv \in E$.)

Clique \leq_P Vertex-Cover

Suppose that $I = (G, k)$ is an instance of **Clique**, where $G = (V, E)$, $V = \{v_1, \dots, v_n\}$ and $1 \leq k \leq n$.

Construct an instance $f(I) = (H, \ell)$ of **Vertex Cover**, where $H = (V, F)$, $\ell = n - k$ and

$$v_i v_j \in F \Leftrightarrow v_i v_j \notin E.$$

H is called the **complement** of G , because every edge of G is a non-edge of H and every non-edge of G is an edge of H .

Properties of Polynomial-time Reductions

Suppose that Π_1, Π_2, \dots are decision problems.

Theorem

If $\Pi_1 \leq_P \Pi_2$ and $\Pi_2 \in P$, then $\Pi_1 \in P$.

Theorem

$\Pi_1 \leq_P \Pi_2$ and $\Pi_2 \leq_P \Pi_3$, then $\Pi_1 \leq_P \Pi_3$.

The Complexity Class \mathbf{NPC}

The complexity class \mathbf{NPC} denotes the set of all decision problems Π that satisfy the following two properties:

- $\Pi \in \mathbf{NP}$
- For all $\Pi' \in \mathbf{NP}$, $\Pi' \leq_P \Pi$.

\mathbf{NPC} is an abbreviation for **NP-complete**.

Theorem

If $P \cap \mathbf{NPC} \neq \emptyset$, then $P = \mathbf{NP}$.

Satisfiability and the Cook-Levin Theorem

Problem

CNF-Satisfiability

Instance: A boolean formula F in n boolean variables x_1, \dots, x_n , such that F is the **conjunction** (logical “and”) of m **clauses**, where each clause is the **disjunction** (logical “or”) of literals. (A **literal** is a boolean variable or its negation.)

Question: Is there a truth assignment such that F evaluates to **true**?

Theorem

CNF-Satisfiability \in **NPC**.

Proving Problems NP-complete

Now, given any NP-complete problem, say Π_1 , other problems in NP can be proven to be NP-complete via polynomial reductions **from** Π_1 , as stated in the following theorem:

Theorem

Suppose that the following conditions are satisfied:

- $\Pi_1 \in NPC$,
- $\Pi_1 \leq_P \Pi_2$, and
- $\Pi_2 \in NP$.

Then $\Pi_2 \in NPC$.

More Satisfiability Problems

Problem

3-CNF-Satisfiability

Instance: A boolean formula F in n boolean variables, such that F is the conjunction of m clauses, where each clause is the disjunction of exactly **three** literals.

Question: Is there a truth assignment such that F evaluates to **true**?

Problem

2-CNF-Satisfiability

Instance: A boolean formula F in n boolean variables, such that F is the conjunction of m clauses, where each clause is the disjunction of exactly **two** literals.

Question: Is there a truth assignment such that F evaluates to **true**?

3-CNF-Satisfiability \in **NPC**, while **2-CNF-Satisfiability** \in **P**.

CNF-Satisfiability \leq_P 3-CNF-Satisfiability

Suppose that (X, \mathcal{C}) is an instance of **CNF-SAT**, where $X = \{x_1, \dots, x_n\}$ and $\mathcal{C} = \{C_1, \dots, C_m\}$. For each C_j , do the following:

case 1 If $|C_j| = 1$, say $C_j = \{z\}$, construct four clauses

$$\{z, a, b\}, \{z, a, \bar{b}\}, \{z, \bar{a}, b\}, \{z, \bar{a}, \bar{b}\}.$$

case 2 If $|C_j| = 2$, say $C_j = \{z_1, z_2\}$, construct two clauses

$$\{z_1, z_2, c\}, \{z_1, z_2, \bar{c}\}.$$

case 3 If $|C_j| = 3$, then leave C_j unchanged.

case 4 If $|C_j| \geq 4$, say $C_j = \{z_1, z_2, \dots, z_k\}$, then construct $k - 2$ new clauses

$$\{z_1, z_2, d_1\}, \{\bar{d}_1, z_3, d_2\}, \{\bar{d}_2, z_4, d_3\}, \dots, \\ \{\bar{d}_{k-4}, z_{k-2}, d_{k-3}\}, \{\bar{d}_{k-3}, z_{k-1}, z_k\}.$$

3-CNF-Satisfiability \leq_P Clique

Let I be the instance of **3-CNF-SAT** consisting of n variables, x_1, \dots, x_n , and m clauses, C_1, \dots, C_m . Let $C_i = \{z_1^i, z_2^i, z_3^i\}$, $1 \leq i \leq m$.

Define $f(I) = (G, k)$, where $G = (V, E)$ according to the following rules:

- $V = \{v_j^i : 1 \leq i \leq m, 1 \leq j \leq 3\}$,
- $v_j^i v_{j'}^{i'} \in E$ if and only if $i \neq i'$ and $z_j^i \neq \overline{z_{j'}^{i'}}$, and
- $k = m$.

Subset Sum

Problem

Subset Sum

Instance: A list of **sizes** $S = [s_1, \dots, s_n]$; and a **target sum**, T . These are all positive integers.

Question: Does there exist a subset $J \subseteq \{1, \dots, n\}$ such that $\sum_{i \in J} s_i = T$?

Vertex Cover \leq_P Subset Sum

Suppose $I = (G, k)$, where $G = (V, E)$, $|V| = n$, $|E| = m$ and $1 \leq k \leq n$.

Suppose $V = \{v_1, \dots, v_n\}$ and $E = \{e_0, \dots, e_{m-1}\}$. For $1 \leq i \leq n$, $0 \leq j \leq m-1$, let

$$c_{ij} = \begin{cases} 1 & \text{if } e_j \text{ is incident with } v_i \\ 0 & \text{otherwise.} \end{cases}$$

Define $n + m$ sizes and a target sum W as follows:

$$a_i = 10^m + \sum_{j=0}^{m-1} c_{ij} 10^j \quad (1 \leq i \leq n)$$

$$b_j = 10^j \quad (0 \leq j \leq m-1)$$

$$W = k \cdot 10^m + \sum_{j=0}^{m-1} 2 \cdot 10^j$$

Then define $f(I) = (a_1, \dots, a_n, b_0, \dots, b_{m-1}, W)$.

Reductions among NP-complete Problems (summary)



In the above diagram, arrows denote polynomial reductions.

Undecidable Problems

A decision problem Π is **undecidable** if there does not exist an algorithm that solves Π .

If Π is undecidable, then for every algorithm A , there exists at least one instance $I \in \mathcal{I}(\Pi)$ such that $A(I)$ does not find the correct answer (“yes” or “no”) in finite time.

Problem

Halting

Instance: *A computer program A and input x for the program A .*

Question: *When program A is executed with input x , will it halt in finite time?*

Undecidability of the Halting Problem

Suppose that *Halt* is a program that solves the **Halting Problem**. Consider the following algorithm *Strange*.

Algorithm: *Strange*(*A*)
external *Halt*
if not *Halt*(*A*, *A*)
 then return (!)
 else $\begin{cases} i \leftarrow 1 \\ \textbf{while } i \neq 0 \textbf{ do } i \leftarrow i + 1 \end{cases}$

What happens when we run *Strange*(*Strange*)?