CS 350 Operating Systems Course Notes

Fall 2015

David R. Cheriton
School of Computer Science
University of Waterloo

Intro 1

What is an Operating System?

• Three views of an operating system



Application View: what services does it provide?

System View: what problems does it solve?

Implementation View: how is it built?

An operating system is part cop, part facilitator.

CS350 Operating Systems Fall 2015

Intro 2

Application View of an Operating System

- The OS provides an execution environment for running programs.
 - The execution environment provides a program with the processor time and memory space that it needs to run.
 - The execution environment provides interfaces through which a program can use networks, storage, I/O devices, and other system hardware components.
 - * Interfaces provide a simplified, abstract view of hardware to application programs.
 - The execution environment isolates running programs from one another and prevents undesirable interactions among them.

Intro 3

Other Views of an Operating System

System View: The OS manages the hardware resources of a computer system.

- Resources include processors, memory, disks and other storage devices, network interfaces, I/O devices such as keyboards, mice and monitors, and so on.
- The operating system allocates resources among running programs. It controls the sharing of resources among programs.
- The OS itself also uses resources, which it must share with application programs.

Implementation View: The OS is a concurrent, real-time program.

- Concurrency arises naturally in an OS when it supports concurrent applications, and because it must interact directly with the hardware.
- Hardware interactions also impose timing constraints.

CS350 Operating Systems Fall 2015

Intro 4

The Operating System and the Kernel

• Some terminology:



manages hardware

to be fastest (from a

resources so that the things you want

user perspective) are running faster

than background

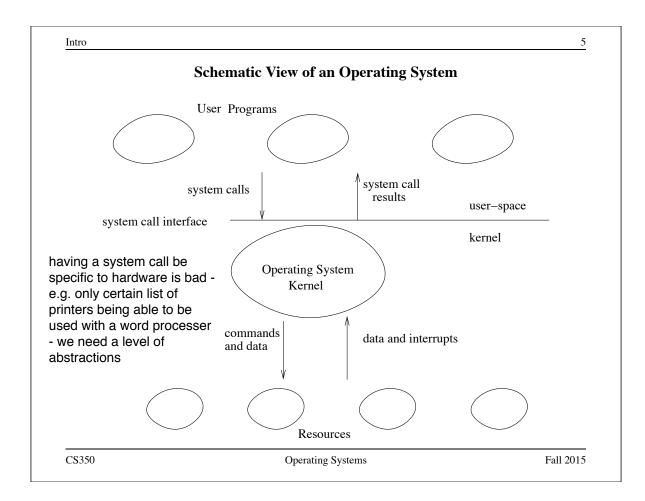
stuff

kernel: The operating system kernel is the part of the operating system that responds to system calls, interrupts and exceptions.

operating system: The operating system as a whole includes the kernel, and may include other related programs that provide services for applications. This may include things like:

- utility programs
- command interpreters
- programming libraries

we're focusing on the kernel



Intro

6

Operating System Abstractions

• The execution environment provided by the OS includes a variety of abstract entities that can be manipulated by a running program. Examples:

files and file systems: abstract view of secondary storage

address spaces: abstract view of primary memory

processes, threads: abstract view of program execution

sockets, pipes: abstract view of network or other message channels

- This course will cover
 - why these abstractions are designed the way they are
 - how these abstractions are manipulated by application programs
 - how these abstractions are implemented by the OS

Intro		7

Course Outline

- Introduction
- Threads and Concurrency
- Synchronization
- Processes and the Kernel
- Virtual Memory
- Scheduling
- Devices and Device Management
- File Systems
- Interprocess Communication and Networking (time permitting)

Review: Program Execution

• Registers



- program counter, stack pointer, . . .
- Memory
 - program code
 - program data
 - program stack containing procedure activation records
- CPU
 - fetches and executes instructions

CS350 Operating Systems Fall 2015

Threads and Concurrency

2

What is a Thread?

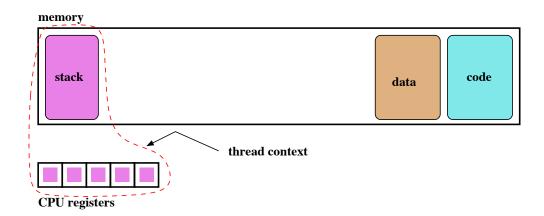
- A thread represents the control state of an executing program.
- A thread has an associated *context* (or state), which consists of
 - the processor's CPU state, including the values of the program counter (PC),
 the stack pointer, other registers, and the execution mode
 (privileged/non-privileged)
 privileged: in kernel space
 unprivileged: in user space (we haven't covered this yet)
 - a stack, which is located in the address space of the thread's process

suspending a thread = pausing a thread

Imagine that you would like to suspend the program execution, and resume it again later. Think of the thread context as the information you would need in order to restart program execution from where it left off when it was suspended.

when you switch threads, the registers are shared by both, we might not want values in registers to be changed while other thread is running





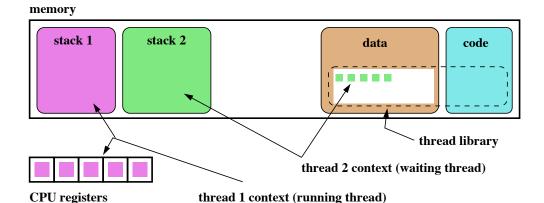
CS350 Operating Systems Fall 2015

Threads and Concurrency 4

Concurrent Threads

- more than one thread may exist simultaneously (why might this be a good idea?)
- each thread has its own context, though they share access to program code and data
- on a uniprocessor (one CPU), at most one thread is actually executing at any time. The others are paused, waiting to resume execution.
- on a multiprocessor, multiple threads may execute at the same time, but if there are more threads than processors then some threads will be paused and waiting

Two Threads, One Running



we know thread 1 is running because its data is in the registers, thread 2 is waiting because its saved stuff in data (though often these are stored to its stack)

CS350 Operating Systems Fall 2015

Threads and Concurrency 6

Thread Interface (Partial), With OS/161 Examples

- a thread library implements threads
- thread library provides a thread interface, used by program code to manipulate threads
- common thread interface functions include
 - create new thread

- end (and destroy) the current thread



void thread exit(void);

- cause current thread to yield (to be discussed later)
 void thread_yield(void);
- see kern/include/thread.h

Example: Creating Threads Using thread_fork()

What kern/synchprobs/catmouse.c actually does is slightly more elaborate than this.

CS350 Operating Systems Fall 2015

Threads and Concurrency

8

Example: Concurrent Mouse Simulation Threads (simplified)

Context Switch, Scheduling, and Dispatching

- the act of pausing the execution of one thread and resuming the execution of another is called a *(thread) context switch*
- what happens during a context switch?
 - 1. decide which thread will run next
 - 2. save the context of the currently running thread
 - 3. restore the context of the thread that is to run next
- the act of saving the context of the current thread and installing the context of the next thread to run is called *dispatching* (the next thread)
- sounds simple, but . . .
 - architecture-specific implementation
 different implementations of stack/heap,
 different register set up
 - thread must save/restore its context carefully, since thread execution
 continuously changes the context ordering of loads and stores needs to be in a specific
 order or you'll overwrite registers you want to save
 - can be tricky to understand (at what point does a thread actually stop? what is it executing when it resumes?)

CS350 Operating Systems Fall 2015

Threads and Concurrency

10

Scheduling

- scheduling means deciding which thread should run next
- scheduling is implemented by a *scheduler*, which is part of the thread library there are lots of complicated ways to decide this, but we're focusing on
- simple *round robin* scheduling:
 - scheduler maintains a queue of threads, often called the *ready queue*
 - the first thread in the ready queue is the running thread
 - on a context switch the running thread is moved to the end of the ready queue, and new first thread is allowed to run
 - newly created threads are placed at the end of the ready queue
- more on scheduling later . . .

Causes of Context Switches

- a call to **thread_yield** by a running thread
 - running thread *voluntarily* allows other threads to run
 - yielding thread remains runnable, and on the ready queue
- a call to **thread_exit** by a running thread
 - running thread is terminated don't need to save context
- running thread *blocks*, via a call to wchan_sleep
 - thread is no longer runnable, moves off of the ready queue and into a wait threads that are waiting on things, put in a wait channel, when they

are ready to be run again

the wait queue

something else moves it back into

- more on this later . . .
- running thread is *preempted*
 - running thread *involuntarily* stops running
 - remains runnable, and on the ready queue
 - e.g. you give it a time limit to run
 - you can stop it with a hardware interrupt

CS350 Operating Systems Fall 2015

Threads and Concurrency

12

Preemption

- without preemption, a running thread could potentially run forever, without yielding, blocking, or exiting
- to ensure *fair* access to the CPU for all threads, the thread library may preempt a running thread
- to implement preemption, the thread library must have a means of "getting control" (causing thread library code to be executed) even though the running thread has not called a thread library function
- this is normally accomplished using *interrupts*

Review: Interrupts

- an interrupt is an event that occurs during the execution of a program
- interrupts are caused by system devices (hardware), e.g., a timer, a disk controller, a network interface
- when an interrupt occurs, the hardware automatically transfers control to a fixed location in memory
- at that memory location, the thread library places a procedure called an *interrupt handler*
- the interrupt handler normally:
 - 1. saves the current thread context (in OS/161, this is saved in a *trap frame* on the current thread's stack)
 - 2. determines which device caused the interrupt and performs device-specific processing
 - 3. restores the saved thread context and resumes execution in that context where it left off at the time of the interrupt.

CS350 Operating Systems Fall 2015

Threads and Concurrency 14

Preemptive Round-Robin Scheduling

- In preemptive round-robin scheduling, the thread library imposes a limit on the amount of time that a thread can run before being preempted
- the amount of time that a thread is allocated is called the scheduling *quantum* (or time slice) the amount of time it gets before context switch
- when the running thread's quantum expires, it is preempted and moved to the back of the ready queue. The thread at the front of the ready queue is dispatched and allowed to run.
- the quantum is an *upper bound* on the amount of time that a thread can run once it has been dispatched
- the dispatched thread may run for less than the scheduling quantum if it yields, exits, or blocks before its quantum expires

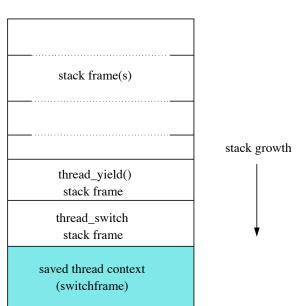
Implementing Preemptive Scheduling

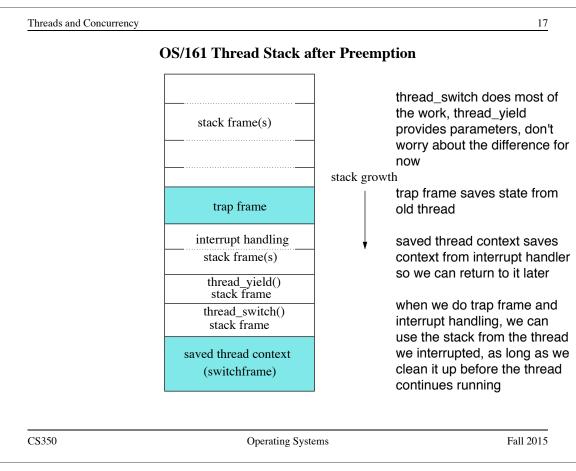
- suppose that the system timer generates an interrupt every t time units, e.g.,
 once every millisecond
- suppose that the thread library wants to use a scheduling quantum q=500t, i.e., it will preempt a thread after half a second of execution
- to implement this, the thread library can maintain a variable called running-time to track how long the current thread has been running:
 - when a thread is intially dispatched, running_time is set to zero
 - when an interrupt occurs, the timer-specific part of the interrupt handler can increment running_time and then test its value
 - * if running_time is less than q, the interrupt handler simply returns and the running thread resumes its execution
 - * if running_time is equal to q, then the interrupt handler invokes thread_yield to cause a context switch

CS350 Operating Systems Fall 2015

Threads and Concurrency 16

OS/161 Thread Stack after Voluntary Context Switch (thread_yield())







piazza example

SEP 22



Threads and Concurrency 18

Implementing Threads

- the thread library is responsibile for implementing threads
- the thread library stores threads' contexts (or pointers to the threads' contexts) when they are not running
- the data structure used by the thread library to store a thread context is sometimes called a *thread control block*

In the OS/161 kernel's thread implementation, thread contexts are stored in thread structures.

The OS/161 thread Structure

```
/* see kern/include/thread.h */
struct thread {
char *t name;
                           /* Name of this thread */
const char *t_wchan_name; /* Wait channel name, if sleeping */
threadstate_t t_state;
                           /* State this thread is in */
                                (running, paused, blocked)
 /* Thread subsystem internal fields. */
struct thread_machdep t_machdep; /* Any machine-dependent goo */
struct threadlistnode t listnode; /* run/sleep/zombie lists */
                                /* Kernel-level stack */
void *t stack;
struct switchframe *t context; /* Register context (on stack) */
                                /* CPU thread runs on */
struct cpu *t_cpu;
 struct proc *t_proc;
                                /* Process thread belongs to */
```

CS350 Operating Systems Fall 2015

Threads and Concurrency

Review: MIPS Register Usage

20

```
zero = ## zero (always returns 0)
  R0,
       at = ## reserved for use by assembler
  R1,
  R2,
             = ## return value / system call number
  R3,
       v1
             = ## return value
             = ## 1st argument (to subroutine)
  R4,
       a0
  R5,
             = ## 2nd argument
       a1
  R6,
       a2
             = ## 3rd argument
             = ## 4th argument
       a3
Review from 241: MIPS registers
- reserved
- arguments
```

- temps (not saved)
- registers that are saved
- reserved for interrupt
- global ptr
- stack ptr
- frame pointer (or just another subroutine register that's saved)
- return address
- note thread yield doesn't take or return parameters, so we don't need to save registers

Review: MIPS Register Usage

```
R08-R15,
          t0-t7 = ## temps (not preserved by subroutines)
          t8-t9 = ## temps (not preserved by subroutines)
R24-R25,
                  ##
                       can be used without saving
          s0-s7 = ## preserved by subroutines
R16-R23,
                  ##
                       save before using,
                  ##
                       restore before return
R26-27,
          k0-k1 = ## reserved for interrupt handler
                = ## global pointer
R28,
          qр
                  ## (for easy access to some variables)
R29,
                = ## stack pointer
          sp
R30,
          s8/fp = ## 9th subroutine reg / frame pointer
R31,
          ra
                = ## return addr (used by jal)
```

CS350 Operating Systems Fall 2015

Threads and Concurrency 22

```
Dispatching on the MIPS (1 of 2)
/* See kern/arch/mips/thread/switch.S */
switchframe switch:
  /* a0: address of switchframe pointer of old thread. */
  /* al: address of switchframe pointer of new thread. */
   /* Allocate stack space for saving 10 registers. 10*4 = 40 */
   addi sp, sp, -40
                    /* Save the registers */
        ra, 36(sp)
   SW
   sw
        gp, 32(sp)
   sw
        s8, 28(sp)
        s6, 24(sp)
   sw
        s5, 20(sp)
   SW
        s4, 16(sp)
   sw
        s3, 12(sp)
   SW
   sw
        s2, 8(sp)
   SW
        s1, 4(sp)
        s0, 0(sp)
   /* Store the old stack pointer in the old thread */
        sp, 0(a0)
   SW
```

Fall 2015 CS350 Operating Systems

Dispatching on the MIPS (2 of 2)

```
/* Get the new stack pointer from the new thread */
lw
     sp, 0(a1)
               /* delay slot for load */
nop
/* Now, restore the registers */
     s0, 0(sp)
lw
     s1, 4(sp)
lw
     s2, 8(sp)
lw
     s3, 12(sp)
lw
     s4, 16(sp)
     s5, 20(sp)
lw
lw
     s6, 24(sp)
lw
     s8, 28(sp)
lw
     gp, 32(sp)
lw
     ra, 36(sp)
                      /* delay slot for load */
nop
/* and return. */
j ra
addi sp, sp, 40
                      /* in delay slot */
.end switchframe_switch
```

CS350 Operating Systems Fall 2015

Threads and Concurrency

24

Dispatching on the MIPS (Notes)

- Not all of the registers are saved during a context switch
- This is because the context switch code is reached via a call to thread_switch and by convention on the MIPS not all of the registers are required to be preserved across subroutine calls
- thus, after a call to switchframe_switch returns, the caller
 (thread_switch) does not expect all registers to have the same values as
 they had before the call to save time, those registers are not preserved by the
 switch
- if the caller wants to reuse those registers it must save and restore them

For trap frame:

- we can't only save the subroutine registers when we get to the interrupt handler
- the compiler can't predict when interrupts will happen and is using temp registers so we need to save allll of them

Concurrency

 On multiprocessors, several threads can execute simultaneously, one on each processor.

• On uniprocessors, only one thread executes at a time. However, because of preemption and timesharing, threads appear to run concurrently.

Concurrency and synchronization are important even on uniprocessors.

CS350 Operating Systems Fall 2015

Synchronization 2

Thread Synchronization

- Concurrent threads can interact with each other in a variety of ways:
 - Threads share access, through the operating system, to system devices (more on this later . . .)
 - Threads may share access to program data, e.g., global variables.
- A common synchronization problem is to enforce *mutual exclusion*, which means making sure that only one thread at a time uses a shared object, e.g., a variable or a device.
- The part of a program in which the shared object is accessed is called a *critical* section.

important term: critical section: part of a program in which the shared object is accessed important for mutual exclusion: only one thread is inside the critical section at a time

Critical Section Example (Part 0)

```
/* Note the use of volatile */
int volatile total = 0;

void add() {
    int i;
    for (i=0; i<N; i++) {
        total++;
    }
}</pre>

    void sub() {
    int i;
    for (i=0; i<N; i++) {
        total---;
    }
}</pre>
```

If one thread executes add and another executes sub what is the value of total when they have finished?

CS350 Operating Systems Fall 2015

- global variable that uses keyword volatile

CS350

- see in these slides how doing both operations and saving back to the global at the same time can create problems

Fall 2015

```
Synchronization
                  Critical Section Example (Part 0)
/* Note the use of volatile */
int volatile total = 0;
void add() {
                                 void sub() {
   loadaddr R8 total
                                    loadaddr R10 total
   for (i=0; i<N; i++) {
                                    for (i=0; i<N; i++) {
      lw R9 0(R8)
                                        lw R11 0(R10)
      add R9 1
                                        sub R11 1
                                        sw R11 0(R10)
      sw R9 0(R8)
}
                                 }
```

Operating Systems

Critical Section Example (Part 0)

Thread 1 Thread 2

loadaddr R8 total lw R9 0(R8) R9=0 add R9 1 R9=1

<INTERRUPT>

loadaddr R10 total lw R11 0(R10) R11=0 sub R11 1 R11=-1 sw R11 0(R10) total=-1

<INTERRUPT>

sw R9 0(R8) total=1

One possible order of execution.

CS350 Operating Systems Fall 2015

Synchronization 6

Critical Section Example (Part 0)

Thread 1 Thread 2

loadaddr R8 total
lw R9 0(R8) R9=0

<INTERRUPT>

loadaddr R10 total

lw R11 0(R10) R11=0

<INTERRUPT>

add R9 1 R9=1 sw R9 0(R8) total=1

<INTERRUPT>

sub R11 1 R11=-1 sw R11 0(R10) total=-1

Another possible order of execution. Many interleavings of instructions are possible. Synchronization is required to ensure a correct ordering.

The use of volatile

Without volatile the compiler could optimize the code. If one thread executes add and another executes sub, what is the value of total when they have finished?

CS350 Operating Systems Fall 2015

Synchronization 8

The use of volatile

The compiler could aggressively optimize the code. Volatile tells the compiler that the object may change for reasons which cannot be determined from the local code (e.g., due to interaction with a device or because of another thread).

- volatile tells the compiler that the variable can change value in other threads
- without volatile keyword, will optimize code in ways that can cause problems

The use of volatile

```
/* Note the use of volatile */
int volatile total = 0;
void add() {
                               void sub() {
   loadaddr R8 total
                                  loadaddr R10 total
   for (i=0; i<N; i++) {
                                  for (i=0; i<N; i++) {
      lw R9 0(R8)
                                      lw R11 0(R10)
      add R9 1
                                      sub R11 1
                                      sw R11 0(R10)
      sw R9 0(R8)
}
                               }
```

The volatile declaration forces the compiler to load and store the value on every use. Using volatile is necessary but not sufficient for correct behaviour. Mutual exclusion is also required to ensure a correct ordering of instructions.

CS350 Operating Systems Fall 2015

Synchronization 10

Ensuring Correctness with Multiple Threads

Threads must enforce mutual exclusion.



Another Critical Section Example (Part 1)

```
int list_remove_front(list *lp) {
   int num;
   list_element *element;
   assert(!is_empty(lp));
   element = lp->first;
   num = lp->first->item;
   if (lp->first == lp->last) {
     lp->first = lp->last = NULL;
   } else {
     lp->first = element->next;
   }
   lp->num_in_list--;
   free(element);
   return num;
}
```

The list_remove_front function is a critical section. It may not work properly if two threads call it at the same time on the same list. (Why?)

CS350 Operating Systems Fall 2015

Synchronization 12

Another Critical Section Example (Part 2)

```
void list_append(list *lp, int new_item) {
   list_element *element = malloc(sizeof(list_element));
   element->item = new_item
   assert(!is_in_list(lp, new_item));
   if (is_empty(lp)) {
      lp->first = element; lp->last = element;
   } else {
      lp->last->next = element; lp->last = element;
   }
   lp->num_in_list++;
}
```

The list_append function is part of the same critical section as list_remove_front. It may not work properly if two threads call it at the same time, or if a thread calls it while another has called list_remove_front

Enforcing Mutual Exclusion

- mutual exclusion algorithms ensure that only one thread at a time executes the code in a critical section
- several techniques for enforcing mutual exclusion
 - exploit special hardware-specific machine instructions, e.g.,
 - * test-and-set,
 - * compare-and-swap, or
 - * load-link / store-conditional,

that are intended for this purpose

 control interrupts to ensure that threads are not preempted while they are executing a critical section

atomic instructions do multiple things but are indivisible, you cannot context switch in the middle

CS350 Operating Systems Fall 2015

Synchronization 14

Disabling Interrupts

- On a uniprocessor, only one thread at a time is actually running.
- If the running thread is executing a critical section, mutual exclusion may be violated if
 - 1. the running thread is preempted (or voluntarily yields) while it is in the critical section, and
 - 2. the scheduler chooses a different thread to run, and this new thread enters the same critical section that the preempted thread was in
- Since preemption is caused by timer interrupts, mutual exclusion can be enforced by disabling timer interrupts before a thread enters the critical section, and re-enabling them when the thread leaves the critical section.

Interrupts in OS/161

This is one way that the OS/161 kernel enforces mutual exclusion on a single processor. There is a simple interface

- spl0() sets IPL to 0, enabling all interrupts.
- splhigh() sets IPL to the highest value, disabling all interrupts.
- splx(s) sets IPL to S, enabling whatever state S represents. (can only really be 0 or 1 in this course)



These are used by splx() and by the spinlock code.

- splraise(int oldipl, int newipl)
- spllower(int oldipl, int newipl)
- For splraise, NEWIPL > OLDIPL, and for spllower, NEWIPL < OLDIPL.

See kern/include/spl.h and kern/thread/spl.c

CS350 Operating Systems Fall 2015

Synchronization 16

Pros and Cons of Disabling Interrupts

- advantages:
 - does not require any hardware-specific synchronization instructions
 - works for any number of concurrent threads
- disadvantages:
 - indiscriminate: prevents all preemption, not just preemption that would threaten the critical section (even safe things can't be run)
 - ignoring timer interrupts has side effects, e.g., kernel unaware of passage of time. (Worse, OS/161's splhigh() disables *all* interrupts, not just timer interrupts.) Keep critical sections *short* to minimize these problems.
 - will not enforce mutual exclusion on multiprocessors (why??)
 because multiple threads can run without context switches

Hardware-Specific Synchronization Instructions

- a test-and-set instruction *atomically* sets the value of a specified memory location and either places that memory location's *old* value into a register
- abstractly, a test-and-set instruction works like the following function:

```
TestAndSet(addr,value)
  old = *addr;    // get old value at addr
  *addr = value;    // write new value to addr
  return old;
```

these steps happen atomically

• example: x86 xchg instruction:

```
xchg src,dest
```

where src is typically a register, and dest is a memory address. Value in register src is written to memory at address dest, and the old value at dest is placed into src.

CS350 Operating Systems Fall 2015



bathroom analogy

Synchronization 18

Alternatives to Test-And-Set

Compare-And-Swap

• example: SPARC cas instruction

```
cas addr,R1,R2
```

if value at addr matches value in R1 then swap contents of addr and R2

- load-linked and store-conditional (the one we use, because it works with MIPS)
 - Load-linked returns the current value of a memory location, while a subsequent store-conditional to the same memory location will store a new value only if no updates have occurred to that location since the load-linked.
 - more on this later . . .

A Spin Lock Using Test-And-Set

- a test-and-set instruction can be used to enforce mutual exclusion
- for each critical section, define a lock variable, in memory
 boolean volatile lock; /* shared, initially false */
 We will use the lock variable to keep track of whether there is a thread in the critical section, in which case the value of lock will be true
- before a thread can enter the critical section, it does the following:
 while (TestAndSet(&lock,true)) { } /* busy-wait */
- when the thread leaves the critical section, it doeslock = false;
- this enforces mutual exclusion (why?), but starvation is a possibility because only one thread has access at a time

This construct is sometimes known as a *spin lock*, since a thread "spins" in the while loop until the critical section is free.

CS350 Operating Systems Fall 2015

disadvantages:

Synchronization 20

Spinlocks in OS/161

```
struct spinlock {
  volatile spinlock_data_t lk_lock; /* word for spin */
  struct cpu *lk_holder; /* CPU holding this lock */
};

void spinlock_init(struct spinlock *lk);
void spinlock_cleanup(struct spinlock *lk);
void spinlock_acquire(struct spinlock *lk);
void spinlock_release(struct spinlock *lk);
bool spinlock_do_i_hold(struct spinlock *lk);
```

Spinning happens in spinlock_acquire

Spinlocks in OS/161

```
spinlock_init(struct spinlock *lk)
{
   spinlock_data_set(&lk->lk_lock, 0);
   lk->lk_holder = NULL;
}

void spinlock_cleanup(struct spinlock *lk)
{
   KASSERT(lk->lk_holder == NULL);
   KASSERT(spinlock_data_get(&lk->lk_lock) == 0);
}

void spinlock_data_set(volatile spinlock_data_t *sd, unsigned val)
{
   *sd = val;
}
```

CS350 Operating Systems Fall 2015

Synchronization 22

Acquiring a Spinlock in OS/161

```
void spinlock_acquire(struct spinlock *lk)
{
    /* note: code that sets lk->holder has been removed! */
    splraise(IPL_NONE, IPL_HIGH);
    while (1) {
        /* Do test-and-test-and-set to reduce bus contention */
        if (spinlock_data_get(&lk->lk_lock) != 0) {
            continue;
        }
        if (spinlock_data_testandset(&lk->lk_lock) != 0) {
            continue;
        }
        break;
    }
}
```

Using Load-Linked / Store-Conditional

```
spinlock data testandset(volatile spinlock data t *sd)
     spinlock data t x,y;
     /* Test-and-set using LL/SC.
      * Load the existing value into X, and use Y to store 1.
      * After the SC, Y contains 1 if the store succeeded,
      * 0 if it failed. On failure, return 1 to pretend
      * that the spinlock was already held.
      */
     y = 1;
- load link is like a normal load operation, but it puts a watch on the variable that you just loaded
- if any other threads modifies the global memory, store conditional will fail
```

- test and set: reads old value, overwrites with new value, returns old value
- x could be 1 and it could be 0

CS350 Operating Systems Fall 2015

Synchronization 24

Using Load-Linked / Store-Conditional (Part 2)

```
asm volatile(
    ".set push;"
                     /* save assembler mode */
    ".set mips32;"
                     /* allow MIPS32 instructions */
                     /* avoid unwanted optimization */
    ".set volatile;"
    "11 %0, 0(%2);"
                     /*
                          x = *sd */
    "sc %1, 0(%2);"
                     /* *sd = y; y = success? */
                     /* restore assembler mode */
    ".set pop"
    : "=r" (x), "+r" (y) : "r" (sd));
  if (y == 0) {
    return 1;
  }
 return x;
}
```

Fall 2015 CS350 Operating Systems

Releasing a Spinlock in OS/161

```
void spinlock_release(struct spinlock *lk)
{
   /* Note: code that sets lk->holder has been removed! */
   spinlock_data_set(&lk->lk_lock, 0);
   spllower(IPL_HIGH, IPL_NONE);
}
```

CS350 Operating Systems Fall 2015

Synchronization 26

Pros and Cons of Spinlocks

- Pros:
 - can be efficient for short critical sections
 - works on multiprocessors
- Cons:
 - CPU is busy (nothing else runs) while waiting for lock
 - starvation is possible

Main disadvantage

- wasting energy waiting
- metaphor: have to get up at 7, checking the clock every 10 hearbeats to see if it's 7 yet and then going back to sleep for 10 hearbeats (this is ridiculous!)

Thread Blocking

- Sometimes a thread will need to wait for an event. For example, if a thread needs to access a critical section that is busy, it must wait for the critical section to become free before it can enter
- other examples that we will see later on:
 - wait for data from a (relatively) slow device
 - wait for input from a keyboard
 - wait for busy device to become idle
- With spinlocks, threads *busy wait* when they cannot enter a critical section. This means that a processor is busy doing useless work. If a thread may need to wait for a long time, it would be better to avoid busy waiting.
- To handle this, the thread scheduler can *block* threads.
- A blocked thread stops running until it is signaled to wake up, allowing the processor to run some other thread.

CS350 Operating Systems Fall 2015

Synchronization 28

Thread Blocking in OS/161

- OS/161 thread library functions for blocking and unblocking threads:
 - void wchan_lock(struct wchan *wc);
 - ; <mark>|</mark>
 - void wchan_unlock(struct wchan *wc);
 - * locks/unlocks the wait channel wc
 - void wchan_sleep(struct wchan *wc);
 - * blocks calling thread on wait channel wc
 - * channel must be locked, will be unlocked upon return
 - void wchan_wakeall(struct wchan *wc);
 - * unblock all threads sleeping on wait channel wc
 - void wchan_wakeone(struct wchan *wc);
 - * unblock one thread sleeping on wait channel wc

both wakeall and wakeone automatically and internally acquire and release lock

Note: current implementation is FIFO but not promised by the interface

Thread Blocking in OS/161

- wchan_sleep() is much like thread_yield(). The calling thread is voluntarily giving up the CPU, so the scheduler chooses a new thread to run, the state of the running thread is saved and the new thread is dispatched. However:
 - after a thread_yield(), the calling thread is ready to run again as soon as it is chosen by the scheduler
 - after a wchan_sleep(), the calling thread is *blocked*, and must not be scheduled to run again until after it has been explicitly unblocked by a call to wchan_wakeone() or wchan_wakeall().

wchan_sleep:

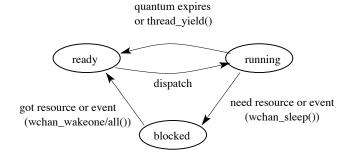
- can only do when wc is locked, makes channel go to sleep
- blocks calling thread on wait channel wc
- channel must be locked, will be unlocked upon return
- right before it goes to sleep, we release the lock so the thread that wakes it up can aquire the lock

CS350 Operating Systems Fall 2015

Synchronization 30

Thread States

• a very simple thread state transition diagram



• the states:

running: currently executing

ready: ready to execute

blocked: waiting for something, so not ready to execute.

Semaphøres

- A semaphore is a synchronization primitive that can be used to enforce mutual exclusion requirements. It can also be used to solve other kinds of synchronization problems.
- A semaphore is an object that has an integer value, and that supports two operations:
 - **P:** if the semaphore value is greater than 0, decrement the value. Otherwise, wait until the value is greater than 0 and then decrement it.
 - V: increment the value of the semaphore
- Two kinds of semaphores:

counting semaphores: can take on any non-negative value **binary semaphores:** take on only the values 0 and 1. (V on a binary semaphore with value 1 has no effect.)

By definition, the P and V operations of a semaphore are *atomic*.

CS350 Operating Systems Fall 2015

Synchronization 32

Mutual Exclusion Using a Semaphore

A Simple Example using Semaphores

```
volatile int total = 0;

void add() {
    int i;
    for (i=0; i<N; i++) {
        P(sem);
        total++;
        V(sem);
    }
}</pre>
void sub() {
    int i;
    for (i=0; i<N; i++) {
        P(sem);
        total---;
        V(sem);
    }
}
```

What type of semaphore can be used for sem?



CS350 Operating Systems Fall 2015

Synchronization

34

OS/161 Semaphores

see kern/include/synch.h and kern/thread/synch.c

```
OS/161 Semaphores: P() from kern/thread/synch.c
P(struct semaphore *sem)
  KASSERT(sem != NULL);
  KASSERT(curthread->t_in_interrupt == false);
  spinlock_acquire(&sem->sem_lock);
    while (sem->sem_count == 0) {
      /* Note: we don't maintain strict FIFO ordering */
      wchan_lock(sem->sem_wchan);
      spinlock_release(&sem->sem_lock);
      wchan_sleep(sem->sem_wchan);
      spinlock_acquire(&sem->sem_lock);
    KASSERT(sem->sem_count > 0);
    sem->sem_count--;
  spinlock_release(&sem->sem_lock);
CS350
                         Operating Systems
                                                        Fall 2015
```

Synchronization 36

OS/161 Semaphores: V() from kern/thread/synch.c

```
V(struct semaphore *sem)
{
   KASSERT(sem != NULL);

   spinlock_acquire(&sem->sem_lock);

   sem->sem_count++;
   KASSERT(sem->sem_count > 0);
   wchan_wakeone(sem->sem_wchan);

   spinlock_release(&sem->sem_lock);
}
```

Producer/Consumer Synchronization

- suppose we have threads that add items to a list (producers) and threads that remove items from the list (consumers)
- suppose we want to ensure that consumers do not consume if the list is empty instead they must wait until the list has something in it
- this requires synchronization between consumers and producers
- semaphores can provide the necessary synchronization, as shown on the next slide

CS350 Operating Systems Fall 2015

Synchronization 38

Producer/Consumer Synchronization using Semaphores

```
struct semaphore *s;
s = sem_create("Items", 0); /* initial value is 0 */
Producer's Pseudo-code:
   add item to the list (call list_append())
   V(s);

Consumer's Pseudo-code:
   P(s);
   remove item from the list (call list_remove_front())
```

The Items semaphore does not enforce mutual exclusion on the list. If we want mutual exclusion, we can also use semaphores to enforce it. (How?)



Bounded Buffer Producer/Consumer Synchronization

- suppose we add one more requirement: the number of items in the list should not exceed N
- producers that try to add items when the list is full should be made to wait until the list is no longer full
- We can use an additional semaphore to enforce this new constraint:
 - semaphore Occupied is used to count the number of occupied entries in the list (to ensure nothing is consumed if there are no occupied entries)
 - semaphore Unoccupied is used to count the number of unoccupied entries in the list (to ensure nothing is produced if there are no unoccupied entries)

CS350 Operating Systems Fall 2015

Synchronization 40

Bounded Buffer Producer/Consumer Synchronization with Semaphores

```
Producer's Pseudo-code:
   P(unoccupied);
   add item to the list (call list_append())
   V(occupied);

Consumer's Pseudo-code:
   P(occupied);
   remove item from the list (call list_remove_front())
   V(unoccupied);
```

OS/161 Locks

• OS/161 also uses a synchronization primitive called a *lock*. Locks are intended to be used to enforce mutual exclusion.

```
struct lock *mylock = lock_create("LockName");
lock_aquire(mylock);
   critical section /* e.g., call to list_remove_front */
lock_release(mylock);
```

- A lock is similar to a binary semaphore with an initial value of 1. However, locks also enforce an additional constraint: the thread that releases a lock must be the same thread that most recently acquired it.
- The system enforces this additional constraint to help ensure that locks are used as intended.

we will implement this on our assignment

CS350 Operating Systems Fall 2015

Synchronization 42

Condition Variables

- OS/161 supports another common synchronization primitive: condition variables
- each condition variable is intended to work together with a lock: condition variables are only used *from within the critical section that is protected by the lock*
- three operations are possible on a condition variable:
 - wait: This causes the calling thread to block, and it releases the lock associated with the condition variable. Once the thread is unblocked it reacquires the lock.
 - **signal:** If threads are blocked on the signaled condition variable, then one of those threads is unblocked.
 - **broadcast:** Like signal, but unblocks all threads that are blocked on the condition variable.



Using Condition Variables

- Condition variables get their name because they allow threads to wait for arbitrary conditions to become true inside of a critical section.
- Normally, each condition variable corresponds to a particular condition that is
 of interest to an application. For example, in the bounded buffer
 producer/consumer example on the following slides, the two conditions are:
 - count > 0 (condition variable notempty)
 - count < N (condition variable notfull)
- when a condition is not true, a thread can wait on the corresponding condition variable until it becomes true
- when a thread detects that a condition is true, it uses signal or broadcast to notify any threads that may be waiting

Note that signalling (or broadcasting to) a condition variable that has no waiters has *no effect*. Signals do not accumulate.

CS350 Operating Systems Fall 2015

warning: don't mix semaphores with locks and condition variables

Synchronization 44

Waiting on Condition Variables

- when a blocked thread is unblocked (by signal or broadcast), it reacquires the lock before returning from the wait call
- a thread is in the critical section when it calls wait, and it will be in the critical section when wait returns. However, in between the call and the return, while the caller is blocked, the caller is out of the critical section, and other threads may enter.
- In particular, the thread that calls signal (or broadcast) to wake up the waiting thread will itself be in the critical section when it signals. The waiting thread will have to wait (at least) until the signaller releases the lock before it can unblock and return from the wait call.

This describes Mesa-style condition variables, which are used in OS/161. There are alternative condition variable semantics (Hoare semantics), which differ from the semantics described here.

Bounded Buffer Producer Using Locks and Condition Variables

```
int volatile count = 0;  /* must initially be 0 */
struct lock *mutex;  /* for mutual exclusion */
struct cv *notfull, *notempty; /* condition variables */

/* Initialization Note: the lock and cv's must be created
 * using lock_create() and cv_create() before Produce()
 * and Consume() are called */

Produce(itemType item) {
    lock_acquire(mutex);
    while (count == N) {
        cv_wait(notfull, mutex);
    }
    add item to buffer (call list_append())
    count = count + 1;
    cv_signal(notempty, mutex);
    lock_release(mutex);
}
```

CS350 Operating Systems Fall 2015

Synchronization

Bounded Buffer Consumer Using Locks and Condition Variables

46

```
itemType Consume() {
   lock_acquire(mutex);
   while (count == 0) {
      cv_wait(notempty, mutex);
   }
   remove item from buffer (call list_remove_front())
   count = count - 1;
   cv_signal(notfull, mutex);
   lock_release(mutex);
   return(item);
}
```

Both Produce() and Consume() call cv_wait() inside of a while loop. Why?

Deadlocks

- Suppose there are two threads and two locks, lockA and lockB, both initially unlocked.
- Suppose the following sequence of events occurs
 - 1. Thread 1 does lock_acquire(lockA).
 - 2. Thread 2 does lock_acquire(lockB).
 - 3. Thread 1 does lock_acquire(lockB) and blocks, because lockB is held by thread 2.
 - 4. Thread 2 does lock_acquire(lockA) and blocks, because lockA is held by thread 1.

These two threads are *deadlocked* - neither thread can make progress. Waiting will not resolve the deadlock. The threads are permanently stuck.

CS350 Operating Systems Fall 2015

Synchronization 48

Deadlocks (Another Simple Example)

- Suppose a machine has 64 MB of memory. The following sequence of events occurs.
 - 1. Thread A starts, requests 30 MB of memory.
 - 2. Thread B starts, also requests 30 MB of memory.
 - 3. Thread A requests an additional 8 MB of memory. The kernel blocks thread A since there is only 4 MB of available memory.
 - 4. Thread *B* requests an additional 5 MB of memory. The kernel blocks thread *B* since there is not enough memory available.

These two threads are deadlocked.

Deadlock Prevention





No Hold and Wait: prevent a thread from requesting resources if it currently has resources allocated to it. A thread may hold several resources, but to do so it must make a single request for all of them.



Preemption: take resources away from a thread and give them to another (usually not possible). Thread is restarted when it can acquire all the resources it needs.



Resource Ordering: Order (e.g., number) the resource types, and require that each thread acquire resources in increasing resource type order. That is, a thread may make no requests for resources of type less than or equal to i if it is holding resources of type i.

An example of a *resource allocation graph* ->



discussion on the board at the beginning of class about deadlock and intersection — the intersection is divided into 4 quadrants, to cross the street each car has to acquire two locks to get through 2 quadrants

CS350 Operating Systems Fall 2015

Synchronization 50

Deadlock Detection and Recovery

- main idea: the system maintains the resource allocation graph and tests it to determine whether there is a deadlock. If there is, the system must recover from the deadlock situation.
- deadlock recovery is usually accomplished by terminating one or more of the threads involved in the deadlock
- when to test for deadlocks? Can test on every blocked resource request, or can simply test periodically. Deadlocks persist, so periodic detection will not "miss" them.

Deadlock detection and deadlock recovery are both costly. This approach makes sense only if deadlocks are expected to be infrequent.

What is a Process?

Answer 1: a process is an abstraction of a program in execution



Answer 2: a process consists of



an address space, which represents the memory that holds the program's note that processes *cannot* share information with each other - this creates isolation



- a *thread* of execution (possibly several threads)
- other resources associated with the running program. For example:
 - open files
 - sockets
 - attributes, such as a name (process identifier)
 - **–** ...

A process with one thread is a *sequential* process. A process with more than one thread is a *concurrent* process.

CS350 Operating Systems Fall 2015

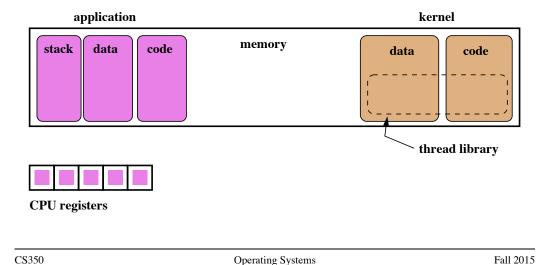
Processes and the Kernel 2

Multiprogramming

- multiprogramming means having multiple processes existing at the same time
- most modern, general purpose operating systems support multiprogramming
- all processes share the available hardware resources, with the sharing coordinated by the operating system:
 - Each process uses some of the available memory to hold its address space.
 The OS decides which memory and how much memory each process gets
 - OS can coordinate shared access to devices (keyboards, disks), since processes use these devices indirectly, by making system calls.
 - Processes timeshare the processor(s). Again, timesharing is controlled by the operating system.
- OS ensures that processes are isolated from one another. Interprocess communication should be possible, but only at the explicit request of the processes involved.

The OS Kernel

- The kernel is a program. It has code and data like any other program.



Processes and the Kernel 4

Kernel Privilege, Kernel Protection

- What does it mean to run in privileged mode?
- Kernel uses privilege to
 - control hardware
 - protect and isolate itself from processes



- privileges vary from platform to platform, but may include:
 - ability to execute special instructions (like halt)



- ability to manipulate processor state (like execution mode)
- ability to access memory addresses that can't be accessed otherwise
- kernel ensures that it is *isolated* from processes. No process can execute or change kernel code, or read or write kernel data, except through controlled mechanisms like system calls.

System Calls

- System calls are an interface between processes and the kernel.
- A process uses system calls to request operating system services.
- Some examples:

Service	OS/161 Examples		
create,destroy,manage processes	fork, execv, waitpid, getpid		
create,destroy,read,write files open,close,remove,read,write			
manage file system and directories	mkdir,rmdir,link,sync		
interprocess communication	pipe,read,write		
manage virtual memory	sbrk		
query,manage system	reboot,time		

CS350 Operating Systems Fall 2015

Processes and the Kernel 6

How System Calls Work

- The hardware provides a mechanism that a running program can use to cause a system call. Often, it is a special instruction, e.g., the MIPS syscall instruction.
- What happens on a system call:
 - the processor is switched to system (privileged) execution mode

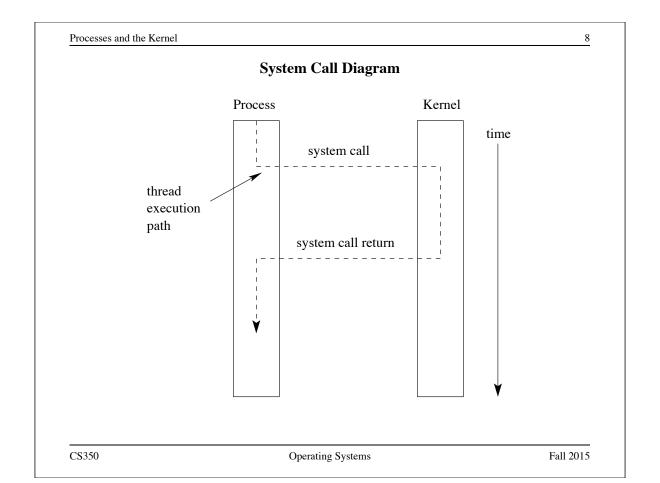


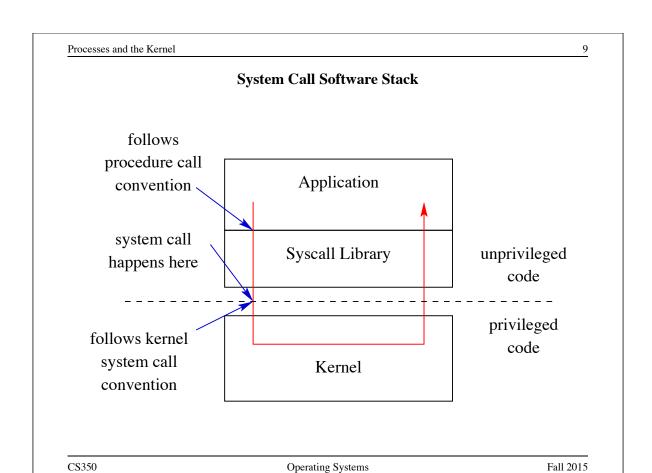
- key parts of the current thread context, such as the program counter, are saved
- the program counter is set to a fixed (specified by the hardware) memory address, which is within the kernel's address space

System Call Execution and Return

- Once a system call occurs, the calling thread will be executing a system call handler, which is part of the kernel, in privileged mode.
- The kernel's handler determines which service the calling process wanted, and performs that service.
- When the kernel is finished, it returns from the system call. This means:
 - restore the key parts of the thread context that were saved when the system call was made
 - switch the processor back to unprivileged (user) execution mode
- Now the thread is executing the calling process' program again, picking up where it left off when it made the system call.

A system call causes a thread to stop executing application code and to start executing kernel code in privileged mode. The system call return switches the thread back to executing application code in unprivileged mode.





OS/161 close System Call Description

Library: standard C library (libc)

Synopsis:

#include <unistd.h>

closing a file

int

close(int fd);

Description: The file handle fd is closed. . . .

Return Values: On success, close returns 0. On error, -1 is returned and errno

is set according to the error encountered.

Errors:

EBADF: fd is not a valid file handle

EIO: A hard I/O error occurred hard drive is not working

An Example System Call: A Tiny OS/161 Application that Uses close

```
/* Program: user/uw-testbin/syscall.c */
#include <unistd.h>
#include <errno.h>

int
main()
{
   int x;
   x = close(999);
   if (x < 0) {
     return errno;
   }
   return x;
}</pre>
```

CS350 Operating Systems Fall 2015

Processes and the Kernel 12

Disassembly listing of user/uw-testbin/syscall

```
00400050 <main>:
  400050: 27bdffe8
                      addiu sp, sp, -24
                                         set up stack, save return address
                           ra,16(sp)
  400054: afbf0010
  400058: 0c100077
                      jal 4001dc <close>
                                            call close function
  40005c: 240403e7
                      li
                           a0,999 in the delay slot (happens before jump)
                             400060: 04410003
                      bgez
                                                           parameter
  400064: 00000000
                      nop
  400068: 3c021000
                      lui v0,0x1000 put address of errno in v0
                                       put value of errno in v0
  40006c: 8c420000
                           v0,0(v0)
                      lw
  400070: 8fbf0010
                      lw
                           ra, 16(sp)
  400074: 00000000
                      nop
                                     return
  400078: 03e00008
                      jr
                           ra
  40007c: 27bd0018
                      addiu sp, sp, 24
```

MIPS procedure call convention: arguments in a0,a1,..., return value in v0.

The above can be obtained using cs350-objdump -d.

OS/161 MIPS System Call Conventions

- When the syscall instruction occurs:
 - An integer system call code should be located in register R2 (v0)
 - Any system call arguments should be located in registers R4 (a0), R5 (a1), R6 (a2), and R7 (a3), much like procedure call arguments.
- When the system call returns
 - register R7 (a3) will contain a 0 if the system call succeeded, or a 1 if the system call failed
 - register R2 (v0) will contain the system call return value if the system call succeeded, or an error number (errno) if the system call failed.

CS350 Operating Systems Fall 2015

Processes and the Kernel 14

OS/161 System Call Code Definitions

```
/* Contains a number for every more-or-less standard
/* Unix system call (you will implement some subset). */
#define SYS close
                            49
#define SYS read
                            50
#define SYS pread
                            51
//#define SYS readv
                            52
                                 /* won't be implementing */
//#define SYS preadv
                            53
                                 /* won't be implementing */
#define SYS getdirentry
                            54
                                  if you can pass a pointer to a function, then
#define SYS write
                            55
                                  you could pass any function you want! this
                                  helps us limit what can be called
```

This comes from kern/include/kern/syscall.h. The files in kern/include/kern define things (like system call codes) that must be known by both the kernel and applications.

System Call Wrapper Functions from the Standard Library

The above is disassembled code from the standard C library (libc), which is linked with user/uw-testbin/syscall.o.

CS350 Operating Systems Fall 2015

Processes and the Kernel 16

The OS/161 System Call and Return Processing

```
004000c0 <__syscall>:
  4000c0: 0000000c
                      syscall
  4000c4: 10e00005
                      beqz
                             a3,4000dc < syscall+0x1c>
                                           if no error, jump to returning
  4000c8: 00000000
                      nop
  4000cc: 3c011000
                      lui at,0x1000
  4000d0: ac220000
                           v0,0(at)
                      sw
                                      store errno in the global var
  4000d4: 2403ffff
                      li
                           v1,-1
                                      and set v0 and v1 to -1
  4000d8: 2402ffff
                       li
                           v0,-1
  4000dc: 03e00008
                       jr
                           ra
  4000e0: 00000000
                      nop
```

The system call and return processing, from the standard C library. Like the rest of the library, this is unprivileged, user-level code.

OS/161 MIPS Exception Handler

```
first few instructions see if we're coming from user space
common_exception:
                    or system call (this could be interrupts/exceptions/sys call)
  mfc0 k0, c0_status /* Get status register */
  andi k0, k0, CST_KUp /* Check the we-were-in-user-mode bit */
  beg k0, $0, 1f /* If clear, from kernel, already have stack */
                   /* 1f is branch forward to label 1: */
                   /* delay slot */
  nop
  /* Coming from user mode - find kernel stack */
  mfc0 k1, c0_context
                          /* we keep the CPU number here */
  srl k1, k1, CTX PTBASESHIFT /* shift to get the CPU number */
  sll k1, k1, 2
                          /* shift back to make array index */
  lui k0, %hi(cpustacks) /* get base address of cpustacks[] */
  addu k0, k0, k1
                          /* index it */
  move k1, sp
                          /* Save previous stack pointer */
                          /* Skip to common code */
  b 2f
  lw sp, %lo(cpustacks)(k0) /* Load kernel sp (in delay slot) */
```

CS350 Operating Systems Fall 2015

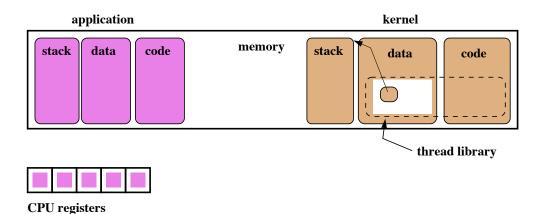
Processes and the Kernel

18

OS/161 MIPS Exception Handler

When the syscall instruction occurs, the MIPS transfers control to address 0x80000080. This kernel exception handler lives there. See kern/arch/mips/locore/exception-mips1.S

OS/161 User and Kernel Thread Stacks



Each OS/161 thread has two stacks, one that is used while the thread is executing unprivileged application code, and another that is used while the thread is executing privileged kernel code.

CS350 Operating Systems Fall 2015



Processes and the Kernel 20

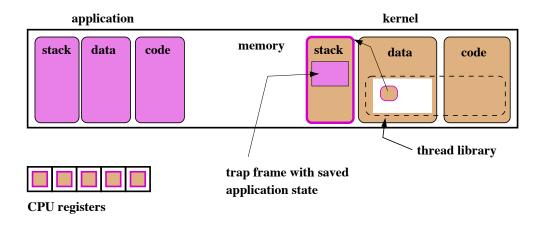
OS/161 MIPS Exception Handler (cont'd)

The common_exception code then does the following:

- 1. allocates a *trap frame* on the thread's kernel stack and saves the user-level application's complete processor state (all registers except k0 and k1) into the trap frame.

 an exception is when something illegal happens, but is also when we use the syscall instruction and also when we do an interrupt
- 2. calls the mips_trap function to continue processing the exception.
- 3. when mips_trap returns, restores the application processor state from the trap frame to the registers
- 4. issues MIPS jr and rfe (restore from exception) instructions to return control to the application code. The jr instruction takes control back to the location specified by the application program counter when the syscall occurred (i.e., exception PC) and the rfe (which happens in the delay slot of the jr) restores the processor to unprivileged mode

OS/161 Trap Frame



While the kernel handles the system call, the application's CPU state is saved in a trap frame on the thread's kernel stack, and the CPU registers are available to hold kernel execution state.

CS350 Operating Systems Fall 2015

Processes and the Kernel 22

mips_trap: Handling System Calls, Exceptions, and Interrupts

- On the MIPS, the same exception handler is invoked to handle system calls, exceptions and interrupts
- The hardware sets a code to indicate the reason (system call, exception, or interrupt) that the exception handler has been invoked
- OS/161 has a handler function corresponding to each of these reasons. The mips_trap function tests the reason code and calls the appropriate function: the system call handler (syscall) in the case of a system call.
- mips_trap can be found in kern/arch/mips/locore/trap.c.

Interrupts and exceptions will be presented shortly



}

OS/161 System Call Handler

```
we can't just look at v0 and find out
syscall(struct trapframe *tf)
                                          what the problem was
   callno = tf->tf v0; retval = 0;
   switch (callno) {
                                          the trapframe is a snapshot of the
     case SYS reboot:
                                          registers when the thread went into
        err = sys_reboot(tf->tf_a0);
                                          the kernel - the v0 there has the
       break;
                                          syscall number
     case SYS
                 time:
        err = sys___time((userptr_t)tf->tf_a0,
          (userptr_t)tf->tf_a1);
        break;
     /* Add stuff here */
     default:
       kprintf("Unknown syscall %d\n", callno);
       err = ENOSYS;
        break;
   }
```

system call checks system call code and invokes a handler for the indicated system call. See kern/arch/mips/syscall/syscall.c

CS350 Operating Systems Fall 2015

Processes and the Kernel

OS/161 MIPS System Call Return Handling

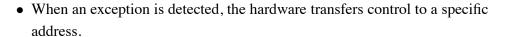
24

```
if (err) {
                                                   we change the
  tf->tf v0 = err;
                                                  trapframe and not the
  tf->tf_a3 = 1;
                        /* signal an error */
                                                   local registers, because
} else {
                                                  the trapframe will
  /* Success. */
                                                   restore later and
  tf->tf v0 = retval;
                                                   overwrite and local
  tf->tf a3 = 0;
                        /* signal no error */
                                                   registers
/* Advance the PC, to avoid the syscall again. */
tf->tf epc += 4;
/* Make sure the syscall code didn't forget to lower spl */
KASSERT(curthread->t curspl == 0);
/* ...or leak any spinlocks */
KASSERT(curthread->t iplhigh count == 0);
```

syscall must ensure that the kernel adheres to the system call return convention.

Exceptions

- Exceptions are another way that control is transferred from a process to the kernel.
- Exceptions are conditions that occur during the execution of an instruction by a process. For example, arithmetic overflows, illegal instructions, or page faults (to be discussed later).
- Exceptions are detected by the hardware.



• Normally, a kernel exception handler is located at that address.

Exception handling is similar to, but not identical to, system call handling. (What is different?)

CS350 Operating Systems Fall 2015

Processes and the Kernel 26

MIPS Exceptions

```
EX_IRQ
                /* Interrupt */
                /* TLB Modify (write to read-only page) */
EX MOD
EX TLBL
                /* TLB miss on load */
          2
                /* TLB miss on store */
EX TLBS
          3
EX ADEL
                /* Address error on load */
                /* Address error on store */
EX ADES
          5
EX IBE
                /* Bus error on instruction fetch */
           6
EX DBE
          7
                /* Bus error on data load *or* store */
EX SYS
                /* Syscall */
          8
                /* Breakpoint */
EX BP
          9
EX RI
          10
                /* Reserved (illegal) instruction */
EX CPU
           11
                /* Coprocessor unusable */
                                               11 — happens when using floating point
           12
                /* Arithmetic overflow */
EX OVF
                                               numbers, we don't have a floating point
```

In OS/161, mips_trap uses these codes to decide whether it has been invoked because of an interrupt, a system call, or an exception.

processor, so don't use floats or doubles!

Interrupts (Revisited)

- Interrupts are a third mechanism by which control may be transferred to the kernel
- Interrupts are similar to exceptions. However, they are caused by hardware devices, not by the execution of a program. For example:
 - a network interface may generate an interrupt when a network packet arrives
 - a disk controller may generate an interrupt to indicate that it has finished writing data to the disk
 - a timer may generate an interrupt to indicate that time has passed
- Interrupt handling is similar to exception handling current execution context is saved, and control is transferred to a kernel interrupt handler at a fixed address.

CS350 Operating Systems Fall 2015

PIAZZA EXAMPLE

Processes and the Kernel 28

Interrupts, Exceptions, and System Calls: Summary

- Interrupts, exceptions and system calls are three mechanisms by which control is transferred from an application program to the kernel
- When these events occur, the hardware switches the CPU into privileged mode and transfers control to a predefined location, at which a kernel *handler* should be located.
- The kernel handler creates a *trap frame* and uses it to saves the application thread context so that the handler code can be executed on the CPU.
- Just before the kernel handler finishes executing, it restores the application thread context from the trap frame, before returning control to the application.

In OS/161, trap frames are placed on the *kernel stack* of the thread performed the system call, or of the thread that was running when the interrupt or exception occurred

System Calls for Process Management

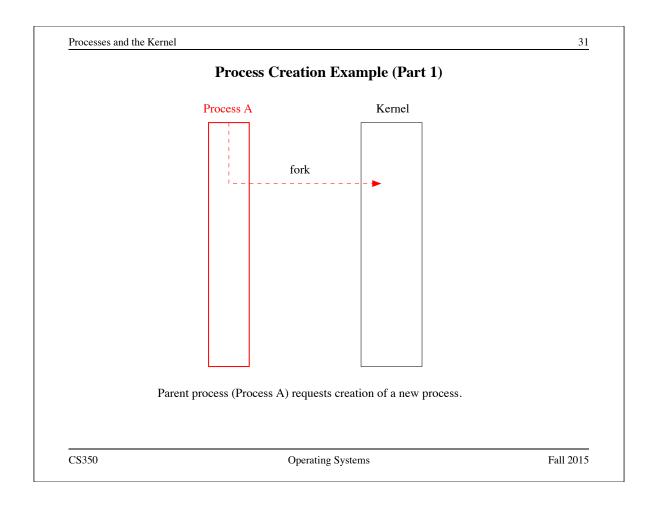
	Linux	OS/161
Creation	fork,execv	fork,execv
Destruction	_exit,kill	_exit
Synchronization	wait,waitpid,pause,	waitpid
Attribute Mgmt	getpid,getuid,nice,getrusage,	getpid

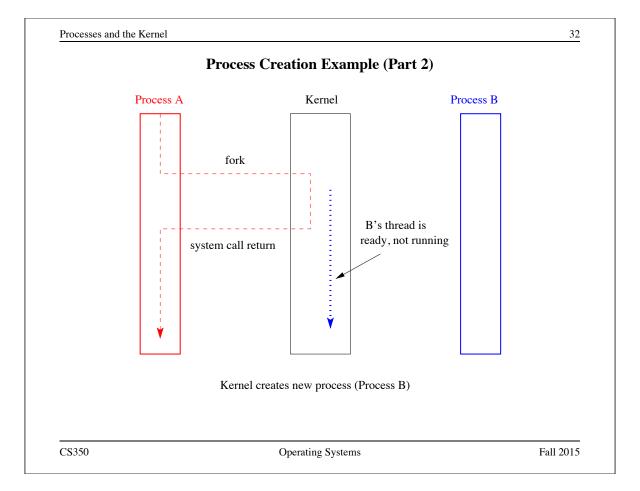
CS350 Operating Systems Fall 2015

Processes and the Kernel 30

The fork, _exit, getpid and waitpid system calls

```
when we fork, we copy the thread and they both
main()
                         continue in the exact same way - the only way to
                         know which process we are now is to check rc
{
    rc = fork(); /* returns 0 to child, pid to parent */
    if (rc == 0) {
        my pid = getpid();
        x = child code();
                             x is status code of the process when it exits (int value)
        _exit(x);
                             that value (think of it as a return value from a function is passed back to the
    } else {
                             parent - encoded with the ~reason for return~ as well as the value
       child_pid = rc;
       parent_code();
       child_exit = waitpid(child_pid); child_exit is now equal to the encoded return
value of the child process
       parent pid = getpid();
                                           wrong signature for waitpid
    }
                                           waitpid(pid_t child_id, int* status, int options)
}
                                           - don't worry about options
                                           - status is the status on exit of child thread
```





The execv system call

takes a binary file in ELF format

Fall 2015

```
int main()
                                     system image consists of code region, instructions,
                                     global data, stack — this is all replaced
  int rc = 0;
                                     pretty much replacing a program with another program
  char *args[4];
  args[0] = (char *) "/testbin/argtest";
  args[1] = (char *) "first";
                                      The first argument, by convention, should
                                      point to the filename associated with the file
  args[2] = (char *) "second";
                                      being executed. The array of pointers must
  args[3] = 0;
                                      be terminated by a NULL pointer.
  rc = execv("/testbin/argtest", args);
  printf("If you see this execv failed\n");
  printf("rc = %d errno = %d\n", rc, errno);
  exit(0);
}
```

Processes and the Kernel 34

Operating Systems

CS350

Combining fork and execv

```
main()
{
   char *args[4];
   /* set args here */
   rc = fork(); /* returns 0 to child, pid to parent */
   if (rc == 0) \{ <-child
     status = execv("/testbin/argtest",args);
     printf("If you see this execv failed\n");
     printf("status = %d errno = %d\n", status, errno);
     exit(0);
   } else {
                                  parent still waits for child
                                 pid doesn't change
     child pid = rc;
                                 same process, different code executing
     parent code();
     child exit = waitpid(child pid);
   }
```

Implementation of Processes

- The kernel maintains information about all of the processes in the system in a data structure often called the process table.
- Per-process information may include:
 - process identifier and owner
 - the address space for the process
 - threads belonging to the process
 - lists of resources allocated to the process, such as open files
 - accounting information

accounting —> misc things like how many open files you have what sockets you're using

CS350 Operating Systems Fall 2015

Processes and the Kernel 36

OS/161 Process

OS/161 Process

CS350 Operating Systems Fall 2015

Processes and the Kernel 38

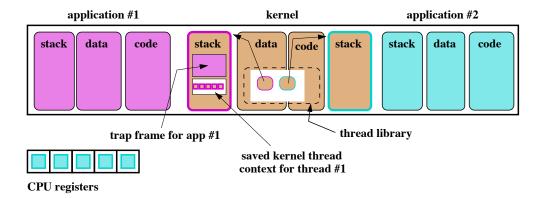
Implementing Timesharing

- whenever a system call, exception, or interrupt occurs, control is transferred from the running program to the kernel
- at these points, the kernel has the ability to cause a context switch from the running process' thread to another process' thread
- notice that these context switches always occur while a process' thread is executing kernel code

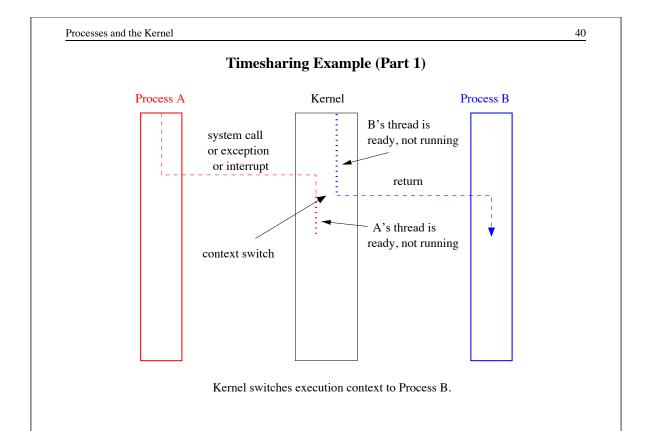
By switching from one process's thread to another process's thread, the kernel timeshares the processor among multiple processes.

processor = CPU

Two Processes in OS/161



CS350 Operating Systems Fall 2015

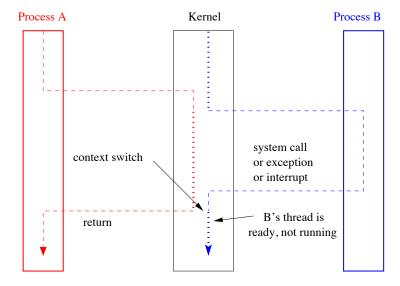


Operating Systems

CS350

Fall 2015

Timesharing Example (Part 2)



Kernel switches execution context back to process A.

CS350 Operating Systems Fall 2015

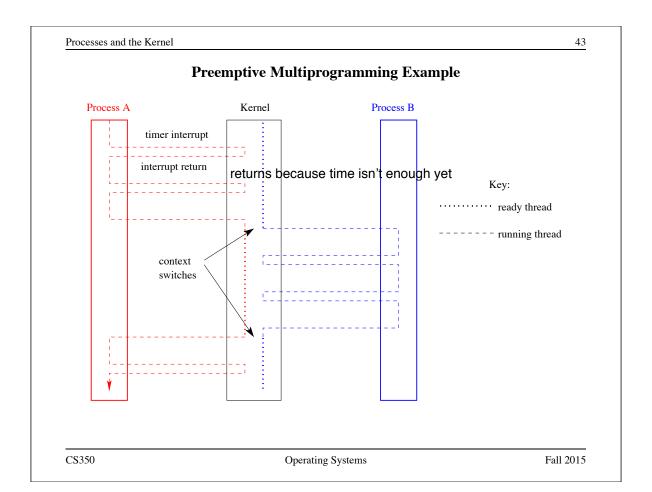
Processes and the Kernel 42

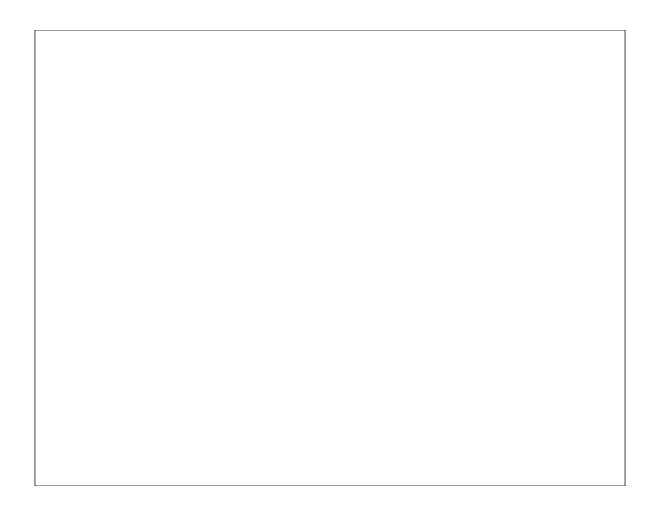
Implementing Preemption

- the kernel uses interrupts from the system timer to measure the passage of time and to determine whether the running process's quantum has expired.
- a timer interrupt (like any other interrupt) transfers control from the running program to the kernel.
- this gives the kernel the opportunity to preempt the running thread and dispatch a new one.

time sharing — the idea I can have lots of things going on at the same time (multitasking) - switching between

preemption — more specific term, can force a thread to stop running





Virtual and Physical Addresses

- Physical addresses are provided directly by the machine.
 - one physical address space per machine

talked about in CS251

- the size of a physical address determines the maximum amount of addressable physical memory
- Virtual addresses (or logical addresses) are addresses provided by the OS to processes.
 prevents accessing things it shouldn't
 - one virtual address space per process
- Programs use virtual addresses. As a program runs, the hardware (with help from the operating system) converts each virtual address to a physical address.
- The conversion of a virtual address to a physical address is called *address translation*.

On the MIPS, virtual addresses and physical addresses are 32 bits long. This limits the size of virtual and physical address spaces.

CS350 Operating Systems Fall 2015

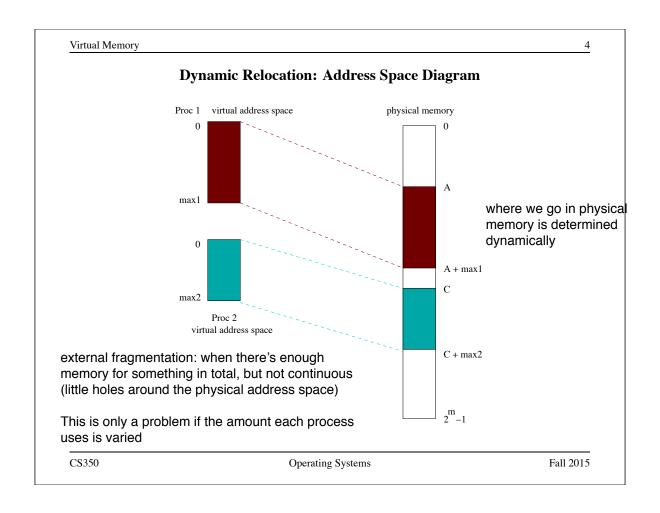
Virtual Memory 2

Simple Address Translation: Dynamic Relocation

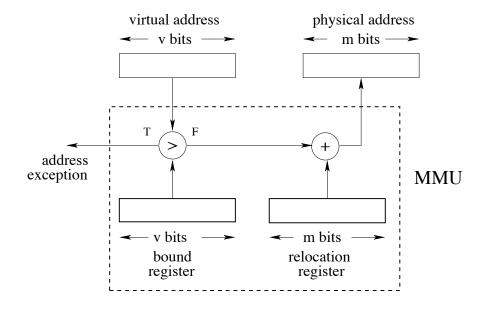
- hardware provides a memory management unit (MMU) which includes a
 relocation register and a limit register (or bound register).
 bound makes sure you're not accessing someone else's memory
- to translate a virtual address to a physical address, the MMU: by mistake
 - checks whether the virtual address is larger than the limit in the limit register
 - if it is, the MMU raises an exception
 - otherwise, the MMU adds the base address (stored in the relocation register) to the virtual address to produce the physical address
- The OS maintains a separate base address and limit for each process, and ensures that the relocation and limit registers in the MMU always contain the base address and limit of the currently-running process.
- To ensure this, the OS must normally change the values in the MMU's registers during each address space switch (i.e., when context switching to a thread of a different process).

Properties of Dynamic Relocation

- each virtual address space corresponds to a contiguous range of physical addresses
- the OS is responsible for deciding *where* each virtual address space should map to in physical memory
 - the OS must track which parts of physical memory are in use, and which parts are free
 - since different address spaces may have different sizes, the OS must allocate/deallocate variable-sized chunks of physical memory
 - this creates the potential for external fragmentation of physical memory: wasted, unallocated space
- the MMU is responsible for performing all address translations, using base and limit information provided to it by the the OS



Dynamic Relocation Mechanism



CS350 Fall 2015 Operating Systems

Virtual Memory 6

Address Translation: Dynamic Relocation Example

Bound register:	0x0011 8000		0x00010000
Relocation register:	0x0010 0000		+0x00A1024
Process 1: virtual addr	0x000A 1024		=
Process 1: physical addr =		?	0x001A1024
Process 1: virtual addr	0x0010 E048		
Process 1: physical addr =		?	0x0020E048
Bound register:	0x0001 0000		
Relocation register:	0x0030 0000		
Process 2: virtual addr	0x0001 8090		
Process 2: physical addr =		?	EXCEPTION
Bound register:	0x000A 1184		
Relocation register	0x0020 0000		
Process 3: virtual addr	0x000A 1024		
Process 3: physical addr =		?	0x02A1024

Bound register:

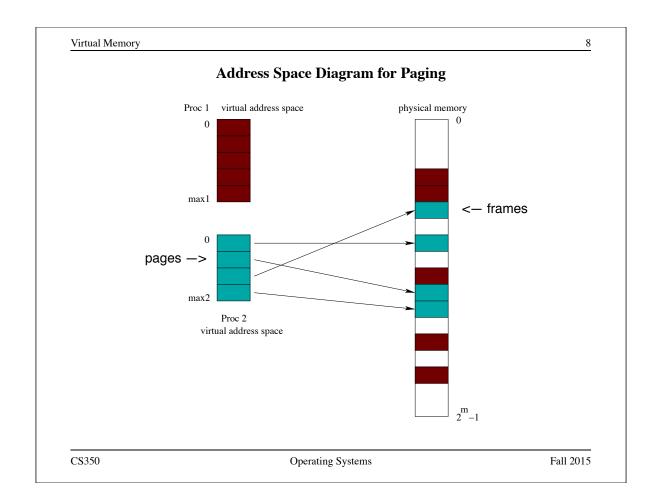
Fall 2015 CS350 Operating Systems

Address Translation: Paging

- Each virtual address space is divided into fixed-size chunks called pages
- Physical address space is divided into *frames*. Frame size matches page size.
- OS maintains a *page table* for each process. Page table specifies the frame in which each of the process's pages is located.
- At run time, MMU translates virtual addresses to physical using the page table of the running process.

from 251:D

this means we don't have fragmentation, because all chunks we want to allocate are the same size



Properties of Paging

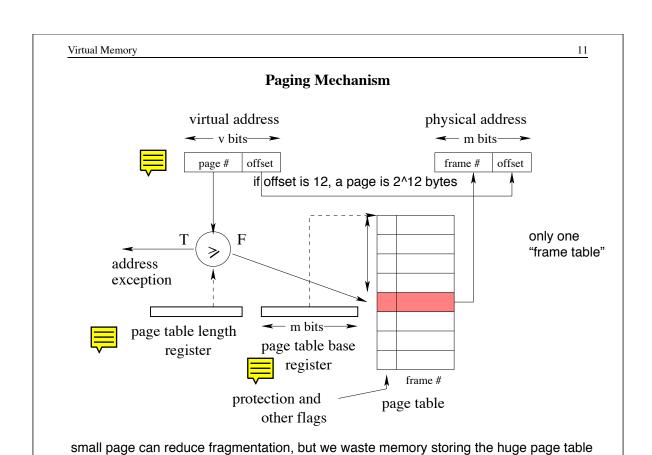
- OS is responsible for deciding which frame will hold each page
 - simple physical memory management
 we could have a frame with 1 byte used wasted space
 - potential for *internal fragmentation* of physical memory: wasted, allocated
 space external fragmentation (from earlier slide) is much worse
 - virtual address space need not be physically contiguous in physical space after translation.
- MMU is responsible for performing all address translations using the page table that is created and maintained by the OS.
- The OS must normally change the values in the MMU registers on each address space switch, so that they refer to the page table of the currently-running process.

CS350 Operating Systems Fall 2015

Virtual Memory 10

How the MMU Translates Virtual Addresses

- The MMU includes a page table base register and a page table length register.
 - the base register contains the (physical) address of the first page table entry for the currently-running process
 - the length register contains the number of entries in the page table of the currently running process.
- To translate a virtual address, the MMU:
 - determines the *page number* and *offset* of the virtual address
 - checks whether the page number is larger than the value in the page table length register
 - if it is, the MMU raises an exception
 - otherwise, the MMU uses the page table to determine the *frame number* of the frame that holds the virtual page, and combines the frame number and offset to determine the physical address



Operating Systems

Fall 2015

Address Translation: Paging Example

• Virtual addr = 32 bits, physical addr = 32 bits, offset = 12 bits 3 digits

CS350

Page Table Proc 1		Page Table Proc 2		
0x0010 0000		0x0050 0000		
Page	Frame	Page	Frame	
0	0x5 0002	0	0x4 0001	
1	0x5 0001	1	0x7 9005	
2	0x5 0000	2	0x7 FADD	
3	0x6 0002			

Process 1 Process 2 Page Table base register 0x0010 0000 0x0050 0000 Page Table Len register 0x0000 0004 0x0000 0003 Virtual addr 0x0000 2004 0x0000 2004 0x7FADD004 0x50000004 ? Physical addr = Virtual addr 0x0000 31A4 0x0000 31A4 0x60021A4 **EXCEPTION!** Physical addr =

Page Table Entries

- the primary payload of each page table entry (PTE) is a frame number
- PTEs typically contain other information as well, such as
 - information provided by the kernel to control address translation by the MMU, such as:
 - * valid bit: is the process permitted to use this part of the address space?
 - * present bit: is this page mapped into physical memory (useful with page replacement, to be discussed later)
 - * protection bits: to be discussed
 - information provided by the MMU to help the kernel manage address spaces, such as:
 - * reference (use) bit: has the process used this page recently?
 - * dirty bit: has the process changed the contents of this page?

CS350 Operating Systems Fall 2015

Virtual Memory 14

Validity and Protection

- during address translation, the MMU checks that the page being used by the process has a *valid* page table entry
 - typically, each PTE contains a *valid bit*
 - invalid PTEs indicate pages that the process is not permitted to use
- the MMU may also enforce other protection rules, for example
 - each PTE may contain a *read-only* bit that indicates whether the corresponding page is read-only, or can be modified by the process
- if a process attempts to access an invalid page, or violates a protection rule, the MMU raises an exception, which is handled by the kernel

The kernel controls which pages are valid and which are protected by setting the contents of PTEs and/or MMU registers.

Summary: Roles of the Kernel and the MMU

- Kernel:
 - manage MMU state on address space switches (context switch from thread in one process to thread in a different process)
 - create and manage page tables
 - manage (allocate/deallocate) physical memory
 - handle exceptions raised by the MMU
- MMU (hardware):
 - translate virtual addresses to physical addresses
 - check for and raise exceptions when necessary

CS350 Operating Systems Fall 2015

Virtual Memory 16

Speed of Address Translation

- Execution of each machine instruction may involve one, two or more memory operations
 - one to fetch instruction
 - one or more for instruction operands
- Address translation through a page table adds one extra memory operation (for page table entry lookup) for each memory operation performed during instruction execution
 - Simple address translation through a page table can cut instruction execution rate in half.
 - More complex translation schemes (e.g., multi-level paging) are even more expensive.
 251 stuff:D
- Solution: include a Translation Lookaside Buffer (TLB) in the MMU
 - TLB is a fast, fully associative address translation cache
 - TLB hit avoids page table lookup

TLB

- Each entry in the TLB contains a (page number, frame number) pair.
- If address translation can be accomplished using a TLB entry, access to the page table is avoided.
 - This is called a *TLB hit*.
- Otherwise, translate through the page table.
 - This is called a *TLB miss*.
- TLB lookup is much faster than a memory access. TLB is an associative memory - page numbers of all entries are checked simultaneously for a match. However, the TLB is typically small (typically hundreds, e.g. 128, or 256 entries).
- If the MMU cannot distinguish TLB entries from different address spaces, then
 the kernel must clear or invalidate the TLB on each address space switch.
 (Why?) we can cache a page#->frame# and switch address and then it won't be valid
 anymore we need to tell the diff between old invalid and ones that are relevant

CS350 Operating Systems Fall 2015

Virtual Memory 18

TLB Management

- A TLB may be hardware-controlled or software-controlled
- In a hardware-controlled TLB, when there is a TLB miss:
 - The MMU (hardware) finds the frame number by performing a page table lookup, translates the virtual address, and adds the translation (page number, frame number pair) to the TLB.
 - If the TLB is full, the MMU evicts an entry to make room for the new one.
- In a software-controlled TLB, when there is a TLB miss:
 - the MMU simply causes an exception, which triggers the kernel exception handler to run



- the kernel must determine the correct page-to-frame mapping and load the mapping into the TLB (evicting an entry if the TLB is full), before returning from the exception
- after the exception handler runs, the MMU retries the instruction that caused the exception.

The MIPS R3000 TLB

- The MIPS has a software-controlled TLB that can hold 64 entries.
- Each TLB entry includes a virtual page number, a physical frame number, an address space identifier (not used by OS/161), and several flags (valid, read-only).
- OS/161 provides low-level functions for managing the TLB:

TLB_Write: modify a specified TLB entry

TLB_Read: read a specified TLB entry

TLB_Probe: look for a page number in the TLB

• If the MMU cannot translate a virtual address using the TLB it raises an exception, which must be handled by OS/161.

See kern/arch/mips/include/tlb.h

CS350 Operating Systems Fall 2015

What is in a Virtual Address Space?

0x00400000 – 0x00401a0c
text (program code) and read—only data

growth

0x100000000 – 0x101200b0
data

0x000000000
0xffffffff

This diagram illustrates the layout of the virtual address space for the OS/161 test application user/testbin/sort

user space

- starts at 0

- goes until 0x8

- stack starts at

0x8, 12 pages in

length, moves

backwards into

- has some code

kernel addr space - from 0x8 to

2³2-1 (4 GB)

user space

data

Address Translation In OS/161: dumbvm

- OS/161 starts with a very simple virtual memory implementation
- virtual address spaces are described by addrspace objects, which record the mappings from virtual to physical addresses

```
struct addrspace {
#if OPT_DUMBVM    1 is for segment 1, 2 is for segment 2
    vaddr_t as_vbase1; /* base virtual address of code segment */
    paddr_t as_pbase1; /* base physical address of code segment */
    size_t as_npages1; /* size (in pages) of code segment */
    vaddr_t as_vbase2; /* base virtual address of data segment */
    paddr_t as_pbase2; /* base physical address of data segment */
    size_t as_npages2; /* size (in pages) of data segment */
    paddr_t as_stackpbase; /* base physical address of stack */
#else
    /* Put stuff here for your VM system */
#endif
};
```

• Notice that each segment must be mapped contiguously into physical memory.

CS350 Operating Systems Fall 2015

Virtual Memory

~~~ covered later on Oct 27 ~~~

22



#### Address Translation Under dumbvm

- the MIPS MMU tries to translate each virtual address using the entries in the TLB
- If there is no valid entry for the page the MMU is trying to translate, the MMU generates a TLB fault (called an *address exception*)
- The vm\_fault function (see kern/arch/mips/vm/dumbvm.c) handles this exception for the OS/161 kernel. It uses information from the current process' addrspace to construct and load a TLB entry for the page.
- On return from exception, the MIPS retries the instruction that caused the exception. This time, it may succeed.

vm\_fault is not very sophisticated. If the TLB fills up, OS/161 will crash!



## Address Translation: OS/161 dumbvm Example

• Note: in OS/161 the stack is 12 pages and the page size is 4 KB = 0 x 1000.

| Variable/Field | Process 1   | Process 2   |
|----------------|-------------|-------------|
| as_vbase1      | 0x0040 0000 | 0x0040 0000 |
| as_pbase1      | 0x0020 0000 | 0x0050 0000 |
| as_npages1     | 0x0000 0008 | 0x0000 0002 |
| as_vbase2      | 0x1000 0000 | 0x1000 0000 |
| as_pbase2      | 0x          | 0x00A0 0000 |
| as_npages2     | 0x600 0010  | 0x0000 0008 |
| as_stackpbase  | 0x0010 0000 | 0x00B0 0000 |

|                 | Process 1   |   | Process 2   |   |
|-----------------|-------------|---|-------------|---|
| Virtual addr    | 0x0040 0004 |   | 0x0040 0004 |   |
| Physical addr = |             | ? |             | ? |
| Virtual addr    | 0x1000 91A4 |   | 0x1000 91A4 |   |
| Physical addr = |             | ? |             | ? |
| Virtual addr    | 0x7FFF 41A4 |   | 0x7FFF 41A4 |   |
| Physical addr = |             | ? |             | ? |
| Virtual addr    | 0x7FFF 32B0 |   | 0x2000 41BC |   |
| Physical addr = |             | ? |             | ? |
|                 |             |   |             |   |

CS350 Operating Systems Fall 2015

Virtual Memory 24

#### **Initializing an Address Space**

 When the kernel creates a process to run a particular program, it must create an address space for the process, and load the program's code and data into that address space

OS/161 *pre-loads* the address space before the program runs. Many other OS load pages *on demand*. (Why?)

- A program's code and data is described in an *executable file*, which is created when the program is compiled and linked
- OS/161 (and some other operating systems) expect executable files to be in ELF (Executable and Linking Format) format very similar to 241
- The OS/161 execv system call re-initializes the address space of a process int execv(const char \*program, char \*\*args)
- The program parameter of the execv system call should be the name of the ELF executable file for the program that is to be loaded into the address space.

#### **ELF Files**

- ELF files contain address space segment descriptions, which are useful to the kernel when it is loading a new address space
- the ELF file identifies the (virtual) address of the program's first instruction
- the ELF file also contains lots of other information (e.g., section descriptors, symbol tables) that is useful to compilers, linkers, debuggers, loaders and other tools used to build programs

a section is variables/data that has same property we'll talk about properties later (e.g. executable, read only)

CS350 Operating Systems Fall 2015

Virtual Memory 26

#### **Address Space Segments in ELF Files**

- The ELF file contains a header describing the segments and segment *images*.
- Each ELF segment describes a contiguous region of the virtual address space.
- The header includes an entry for each segment which describes:
  - the virtual address of the start of the segment
  - the length of the segment in the virtual address space
  - the location of the start of the segment image in the ELF file (if present)
  - the length of the segment image in the ELF file (if present)
- the image is an exact copy of the binary data that should be loaded into the specified portion of the virtual address space
- the image may be smaller than the address space segment, in which case the rest
  of the address space segment is expected to be zero-filled

To initialize an address space, the OS/161 kernel copies segment images from the ELF file to the specifed portions of the virtual address space.

## ELF Files and OS/161

- OS/161's dumbvm implementation assumes that an ELF file contains two segments:
  - a text segment, containing the program code and any read-only data
  - a data segment, containing any other global program data

stack starts out empty

- the ELF file does not describe the stack (why not?) in os161 it's always pointing to the same place (12 pages)
- dumbvm creates a *stack segment* for each process. It is 12 pages long, ending at virtual address 0x7fffffff

Look at kern/syscall/loadelf.c to see how OS/161 loads segments from ELF files

CS350 Operating Systems Fall 2015

Virtual Memory 28

#### **ELF Sections and Segments**

• In the ELF file, a program's code and data are grouped together into *sections*, based on their properties. Some sections:

.text: program code

.rodata: read-only global data.data: initialized global data

.bss: uninitialized global data (Block Started by Symbol)

.sbss: small uninitialized global data

- not all of these sections are present in every ELF file
- normally
  - the .text and .rodata sections together form the text segment
  - the .data, .bss and .sbss sections together form the data segement
- space for *local* program variables is allocated on the stack when the program runs

## The user/uw-testbin/segments.c Example Program (1 of 2)

```
#include <unistd.h>
#define N
                (200)
                              it's not a variable, so it's not stored anywhere
int x = 0xdeadbeef;
                                global variable, not read only, initialized — stored in data
int t1;
int t2;
                                small, global, uninitialized - sbss
int t3;
int array[4096];
                               big, global, uninitialized - bss
char const *str = "Hello World\n"; the const refers to the thing it's pointing to -
                                               so "Hello World\n" is rodata, and str is data
const int z = 0xabcddcba;
                            global constant - rodata
struct example {
  int ypos;
  int xpos;
};
CS350
                                Operating Systems
                                                                       Fall 2015
```

Virtual Memory 30

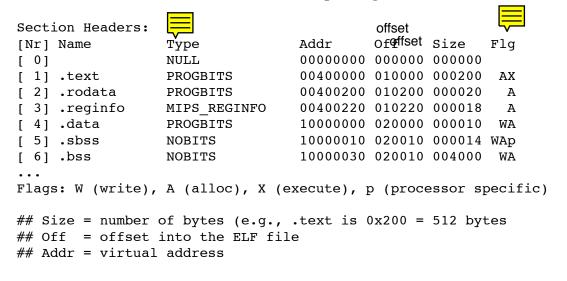
## The user/uw-testbin/segments.c Example Program (2 of 2)

```
int
main()
{
  int count = 0;
  const int value = 1;
  t1 = N;
  t2 = 2;
  count = x + t1;
  t2 = z + t2 + value;

reboot(RB_POWEROFF);
  return 0; /* avoid compiler warnings */
}
```



## **ELF Sections for the Example Program**



The cs350-readelf program can be used to inspect OS/161 MIPS ELF files: cs350-readelf -a segments

CS350 Operating Systems Fall 2015

Virtual Memory 32

## **ELF Segments for the Example Program**

```
Program Headers:
```

```
Type Offset VirtAddr PhysAddr FileSiz MemSiz Flg Align REGINFO 0x010220 0x00400220 0x00400220 0x00018 0x00018 R 0x4 LOAD 0x010000 0x00400000 0x00400000 0x00238 0x00238 R E 0x10000 LOAD 0x020000 0x10000000 0x00010 0x04030 RW 0x10000
```

- segment info, like section info, can be inspected using the cs350-readelf program
- the REGINFO section is not used
- the first LOAD segment includes the .text and .rodata sections
- the second LOAD segment includes .data, .sbss, and .bss

## Contents of the Example Program's .text Section

```
Contents of section .text:
 400000 3c1c1001 279c8000 2408fff8 03a8e824 <...'...$.....$
## Decoding 3c1c1001 to determine instruction
## 0x3c1c1001 = binary 11110000011100000100000000001
you can inspect the code
## instr
         rs
                 rt
                            immediate
                                              to see this information
                                              "if you really wanted to"
## 6 bits | 5 bits | 5 bits |
                            16 bits
## 001111 | 00000 | 11100 |
                           0001 0000 0000 0001
## LUI
                 reg 28
                            0x1001
## LUI
         unused reg 28
                            0x1001
## Load upper immediate into rt (register target)
## lui gp, 0x1001
```

The cs350-objdump program can be used to inspect OS/161 MIPS ELF file section contents: cs350-objdump -s segments

CS350 Operating Systems Fall 2015

Virtual Memory 34

## Contents of the Example Program's .rodata Section

The .rodata section contains the "Hello World" string literal and the constant integer variable z.

## Contents of the Example Program's .data Section

```
Contents of section .data:

10000000 deadbeef 00400210 00000000 00000000 ....@....

## Size = 0x10 bytes = 16 bytes (padding for alignment)

## int x = deadbeef (4 bytes)

## char const *str = "Hello World\n"; (4 bytes)

## address of str = 0x10000004

## value stored in str = 0x00400210.

## NOTE: this is the address of the start

## of the string literal in the .rodata section
```

The .data section contains the initialized global variables str and x.

CS350 Operating Systems Fall 2015

Virtual Memory 36

#### Contents of the Example Program's .bss and .sbss Sections

```
10000000 D x
10000004 D str
                      ## S indicates sbss section
10000010 S t3
                              symbol table tells you
10000014 S t2
10000018 S t1
                              which things aren't initialized yet
1000001c S errno
                              memory is allocated, but no information about what
10000020 S __argv
                              to set that memory to
10000030 B array
                      ## B indicates bss section
10004030 A end
10008000 A qp
```

The t1,t2, and t3 variables are in the .sbss section. The array variable is in the .bss section. There are no values for these variables in the ELF file, as they are uninitialized. The cs350-nm program can be used to inspect symbols defined in ELF files: cs350-nm -n <filename>, in this case cs350-nm -n segments.

## An Address Space for the Kernel

- Each process has its own address space. What about the kernel?
- Three possibilities:



**Kernel in physical space:** disable address translation in privileged system execution mode, enable it in unprivileged mode



**Kernel in separate virtual address space:** need a way to change address translation (e.g., switch page tables) when moving between privileged and unprivileged code



Kernel mapped into portion of address space of every process: OS/161,

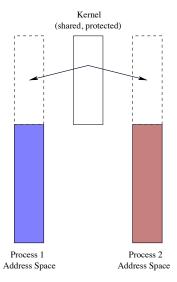
Linux, and other operating systems use this approach

- memory protection mechanism is used to isolate the kernel from applications
- one advantage of this approach: application virtual addresses (e.g., system call parameters) are easy for the kernel to use

CS350 Operating Systems Fall 2015

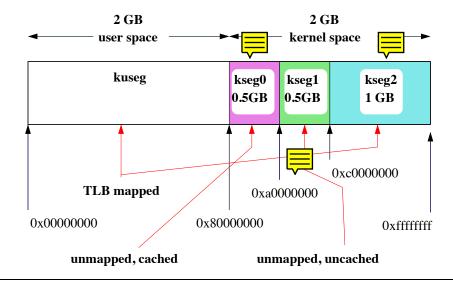
Virtual Memory 38

## The Kernel in Process' Address Spaces



Attempts to access kernel code/data in user mode result in memory protection exceptions, not invalid address exceptions.

### Address Translation on the MIPS R3000



In OS/161, user programs live in kuseg, kernel code and data structures live in kseg0, devices are accessed through kseg1, and kseg2 is not used.

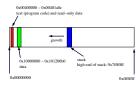
CS350 Operating Systems Fall 2015

Virtual Memory 40

segments: first code

# The Problem of Sparse Address Spaces

then data then stack, which starts at the back of user space and grows back the heap increases the size of the data segment forward



- Consider the page table for user/testbin/sort, assuming a 4 Kbyte page:
  - need 2<sup>19</sup> page table entries (PTEs) to cover the bottom half of the virtual address space (2GB).
  - the text segment occupies 2 pages, the data segment occupies 289 pages,
     and OS/161 sets the initial stack size to 12 pages, so there are only 303 valid
     pages (of 2<sup>19</sup>).
     2^19 is approx half a million
     page table is then bigger than program!
- If dynamic relocation is used, the kernel will need to map a 2GB address space contiguously into physical memory, even though only a tiny fraction of that address space is actually used by the program. so no other programs could use space!
- If paging is used, the kernel will need to create a page table with  $2^{19}$  PTEs, marked as invalid, but we still have to waste 4 bytes per page to mark it as invalid

## **Handling Sparse Address Spaces**

- Use dynamic relocation, but provide separate base and length for each valid segment of the address space. Do not map the rest of the address space.
  - OS/161 dumbvm uses a simple variant of this idea, which depends on having a software-managed TLB.
  - A more general approach is *segmentation*.
- A second approach is to use *multi-level paging* 
  - replace the single large linear page table with a hierarchy of smaller page tables
  - a sparse address space can be mapped by a sparse tree hierarchy
  - easier to manage several smaller page tables than one large one (remember: each page table must be continguous in physical memory!)

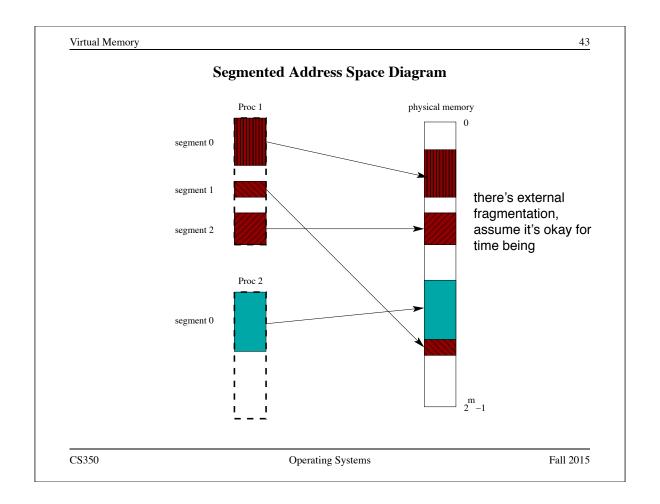
CS350 Operating Systems Fall 2015

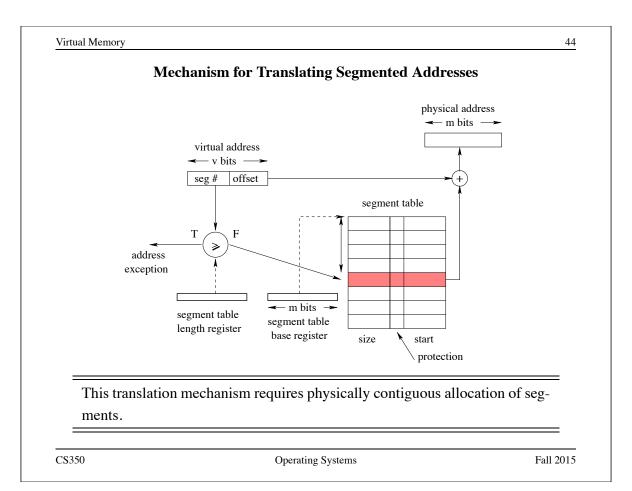
Virtual Memory 42

## Segmentation

- Often, programs (like sort) need several virtual address segments, e.g, for code, data, and stack.
- With segmentation, a virtual address can be thought of as having two parts: (segment ID, address within segment)
- Each segment also has a length.

hard code the segments address space refers to byte in memory, but we're using the first part to identify which segment





## **Address Translation: Segmented Addresses Example**

• Virtual addr = 32 bits, physical addr = 32 bits, offset = 28 bits

| Segn | nent Table 1 | @ 0x0 | 0010 0000 | Segment Table 2 @ 0x0050 0000 |        |      |           |  |  |  |
|------|--------------|-------|-----------|-------------------------------|--------|------|-----------|--|--|--|
| Seg  | Size         | Prot  | Start     | Seg                           | Size   | Prot | Start     |  |  |  |
| 0    | 0x6000       | Х-    | 0x7 0000  | 0                             | 0x9000 | Х-   | 0x90 0000 |  |  |  |
| 1    | 0x1000       | -W    | 0x6 0000  | 1                             | 0xD000 | -W   | 0xAF 0000 |  |  |  |
| 2    | 0xC000       | -W    | 0x5 0000  | 2                             | 0xA000 | -W   | 0x70 0000 |  |  |  |
| 3    | 0xB000       | -W    | 0000 8x0  |                               |        |      |           |  |  |  |

4 bits / 1 digit — first digit tells us which segment

Process 1 Process 2 Seg Table base reg 0x0010 0000 0x0050 0000 Seg Table Len 0x0000 0004 0x0000 0003 Virtual addr 0x1000 B481 0x0000 2004 Physical addr = 0x7 2004 0xAFB481 Virtual addr 0x2000 31A4 0x2000 B1A4 0x831A4 Physical addr = exception B1A4 > A000

CS350 Operating Systems Fall 2015

we did a page translation example (on piazza)

Oct 27

Virtual Memory go back to slides 21-23 46

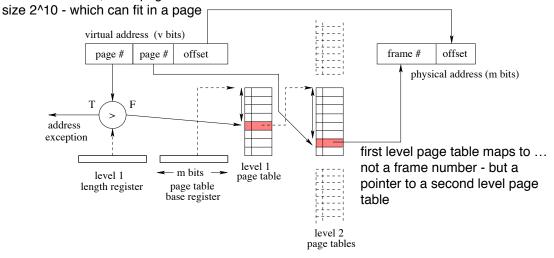
## **Handling Sparse Paged Virtual Address Spaces**

- Large paged virtual address spaces require large page tables.
- example:  $2^{48}$  byte virtual address space, 8 Kbyte ( $2^{13}$  byte) pages, 4 byte page table entries means

$$\frac{2^{48}}{2^{13}}2^2 = 2^{37}$$
 bytes per page table

- page tables for large address spaces may be very large, and
  - they must be in memory, and
  - they must be physically contiguous

# e.g. 12 byte offset, 20 bytes - let's do **Two-Level Paging** 10 for each level, each page table is



we could make level2 page tables only when needed - if we have 3 segments we only need 4 total page tables — if we only had one page table, it'd need to be really big, because addresses are sparse across user virtual address space (dumbvm)

disadvantage to 2level paging - extra lookup

CS350 Operating Systems Fall 2015

Virtual Memory 48

## **Address Translation: Two-Level Paging Example**

• Bits: Virtual addr = 32, physical addr = 32, level 1 = 4, level 2 = 12

|   | Level 1     |          |    | Level 2   |     | Level 2   |     | Level 2   |  |
|---|-------------|----------|----|-----------|-----|-----------|-----|-----------|--|
|   | 0x0010 0000 |          | 0x | 0005 0000 | 0 x | 0004 0000 | 0 x | 0007 0000 |  |
| S | Len         | Address  | P  | Frame     | Р   | Frame     | Р   | Frame     |  |
| 0 | 3           | 0x7 0000 | 0  | 0x8001    | 0   | 0x4001    | 0   | 0x200A    |  |
| 1 | 3           | 0x5 0000 | 1  | 0x9005    | 1   | 0x7005    | 1   | 0x200B    |  |
| 2 | 4           | 0x4 0000 | 2  | 0x6ADD    | 2   | 0x7A00    | 2   | 0x200C    |  |
|   |             |          |    |           | 3   | 0x1023    |     |           |  |
|   |             | • • •    |    | • • •     |     | •••       |     | • • •     |  |

?





Base register

Physical addr =

Level 1 Len 0x0000 0003Virtual addr 0x1000 2004Physical addr = 0x80012004Virtual addr 0x2003 31A4Physical addr = 0x102331A4Virtual addr 0x0002 A049

0x0000 0003 0x3004 020F <u>exception (out of</u> ange in level 1 table) 0x1003 20A8

0x0010 0000

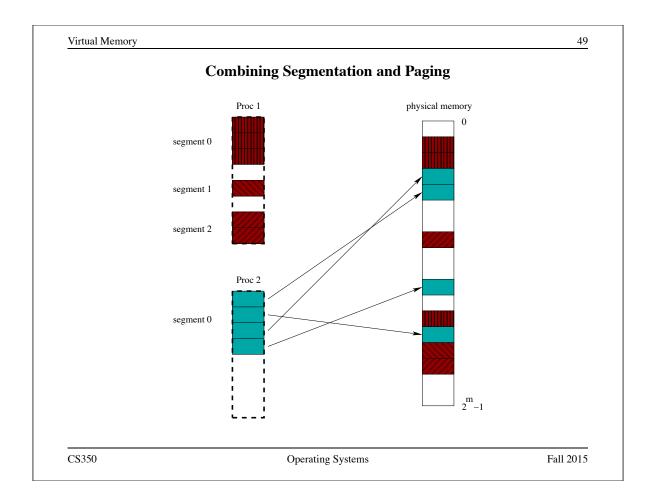
exception (out of range in level 2 table)

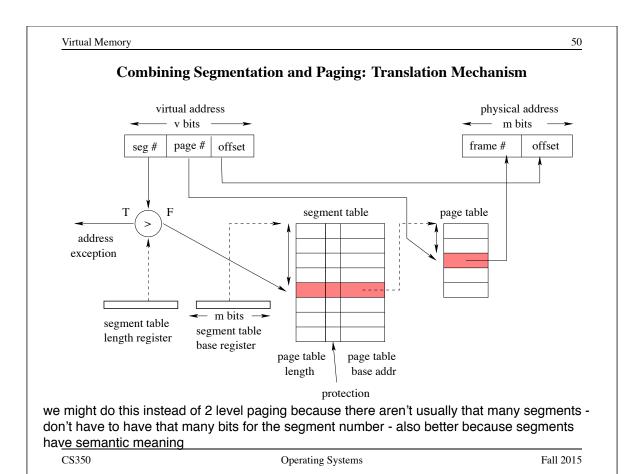
Fall 2015

CS350 Operating Systems

0x200CA049

0x0010 0000





## **Address Translation: Segmentation and Paging Example**

• Bits: Virtual addr = 32, physical addr = 32, Segment = 8, Page = 8

|             | Seg | gment 7 | Table 1  | P   | age Table | Pa  | age Table | Page Table  |        |  |
|-------------|-----|---------|----------|-----|-----------|-----|-----------|-------------|--------|--|
| 0x0020 0000 |     |         |          | 0 x | 0003 0000 | 0x0 | 0001 0000 | 0x0002 0000 |        |  |
| S           | Len | Prot    | Address  | P   | Frame     | P   | Frame     | Р           | Frame  |  |
| 0           | 3   | Х-      | 0x3 0000 | 0   | 0x8001    | 0   | 0x4001    | 0           | 0x200A |  |
| 1           | 4   | -W      | 0x1 0000 | 1   | 0x9005    | 1   | 0x7005    | 1           | 0x200B |  |
| 2           | 3   | -W      | 0x2 0000 | 2   | 0x6ADD    | 2   | 0x7A00    | 2           | 0x200C |  |
|             |     |         |          |     |           | 3   | 0x1023    |             |        |  |

| Base register   | 0x0020 | 0000 |   | 0x0020 | 0000 |   |
|-----------------|--------|------|---|--------|------|---|
| Seg Table Len   | 0x0000 | 0003 |   | 0x0000 | 0003 |   |
| Virtual addr    | 0x0100 | 2004 |   | 0x0102 | C104 |   |
| Physical addr = |        |      | ? |        |      | ? |
| Virtual addr    | 0x0203 | 31A4 |   | 0x0304 | 020F |   |
| Physical addr = |        |      | ? |        |      | ? |
| Virtual addr    | 0x0002 | A049 |   | 0x0104 | 20A8 |   |
| Physical addr = |        |      | ? |        |      | ? |

CS350 Operating Systems Fall 2015

Virtual Memory 52

## **Exploiting Secondary Storage**

#### **Goals:**

- Allow virtual address spaces that are larger than the physical address space.
- Allow greater multiprogramming levels by using less of the available (primary) memory for each process.

## **Method:**

- Allow pages (or segments) from the virtual address space to be stored in secondary storage, e.g., on disks, as well as primary memory.
- Move pages (or segments) between secondary storage and primary memory so that they are in primary memory when they are needed.

## **Paging Policies**

## When to Page?:

reactive/lazy approach: demand paging

Demand paging brings pages into memory when they are used. Alternatively,

the OS can attempt to guess which pages will be used, and prefetch them.

## What to Replace?:

prefetch is proactive



Unless there are unused frames, one page must be replaced for each page that is loaded into memory. A *replacement policy* specifies how to determine which page to replace.

when we swap in memory, it's probably full and we need to swap stuff out (it's usually gonna be full)

Similar issues arise if (pure) segmentation is used, only the unit of data transfer is segments rather than pages. Since segments may vary in size, segmentation also requires a *placement policy*, which specifies where, in memory, a newly-fetched segment should be placed.

just using segmentation while needing swapping is quite rare

CS350 Operating Systems Fall 2015

Virtual Memory

existing bits



• When paging is used, some valid pages may be loaded into memory, and some may not be.

**Page Faults** 

- To account for this, each PTE may contain a *present* bit, to indicate whether the page is or is not loaded into memory
  - V = 1, P = 1: page is valid and in memory (no exception occurs)
  - V = 1, P = 0: page is valid, but is not in memory (exception!)
  - -V=0, P=x: invalid page (exception!) in this case, we'd probably kill the process exception lets the OS move around and fix the memory so we can access
- If V = 1 and P = 0, the MMU will generate an exception if a process tries to access the page. This is called a *page fault*.
- To handle a page fault, the kernel must:
  - issue a request to the disk to copy the missing page into memory,
  - block the faulting process
  - once the copy has completed, set P = 1 in the PTE and unblock the process



• the processor will then retry the instrution that caused the page fault

+ modify frame #

54



## Page Faults in OS/161

- things are a bit different in systems with software-managed TLBs, such as OS/161 on the MIPS processor
- MMUs with software-managed TLBs never check page tables, and thus do not interpret P bits in page table entries
- In an MMU with a software-managed TLB, either there is a valid translation for a page in the TLB, or there is not.
  - If there is not, the MMU generates an exception. It is up to the kernel to determine the reason for the exception. Is this:
    - \* an access to a valid page that is not in memory (a page fault)?
    - \* an access to a valid page that is in memory?
    - \* an access to an invalid page?
  - The kernel should ensure that a page has a translation in the TLB *only* if the page is valid and in memory. (Why?)

CS350 Operating Systems Fall 2015

Virtual Memory 56

## A Simple Replacement Policy: FIFO

- the FIFO policy: replace the page that has been in memory the longest
- a three-frame example:

| Num     | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
|---------|---|---|---|---|---|---|---|---|---|----|----|----|
| Refs    | a | b | c | d | a | b | e | a | b | c  | d  | e  |
| Frame 1 | a | a | a | d | d | d | e | e | e | e  | e  | e  |
| Frame 2 |   | b | b | b | a | a | a | a | a | c  | c  | c  |
| Frame 3 |   |   | c | c | c | b | b | b | b | b  | d  | d  |
| Fault ? | X | X | X | X | X | X | X |   |   | X  | X  |    |

kick e and b are in a memory cause so we're accessed good

longest ago

as you can see, this isn't a super great strategy

## **Optimal Page Replacement**

"it's literally the best"

- There is an optimal page replacement policy for demand paging.
- The OPT policy: replace the page that will not be referenced for the longest time.

| Num     | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
|---------|---|---|---|---|---|---|---|---|---|----|----|----|
| Refs    | a | b | c | d | a | b | e | a | b | c  | d  | e  |
| Frame 1 | a | a | a | a | a | a | a | a | a | c  | c  | c  |
| Frame 2 |   | b | b | b | b | b | b | b | b | b  | d  | d  |
| Frame 3 |   |   | c | d | d | d | e | e | e | e  | e  | e  |
| Fault ? | X | Х | X | X |   |   | X |   |   | X  | X  |    |

a and b

are used soonest, so we replace c

• OPT requires knowledge of the future.

welp, this is pretty hard

CS350 Operating Systems Fall 2015

... but we can make some educated guesses

Virtual Memory 58

## **Other Replacement Policies**

• FIFO is simple, but it does not consider:

some things (root of tree, global variables)
might be accessed a lot more than say, a line of

Frequency of Use: how often a page has been used? code that executes once

**Recency of Use:** when was a page last used? what we were saying with dirty bit earlier

**Cleanliness:** has the page been changed while it is in memory?

- The *principle of locality* suggests that usage ought to be considered in a replacement decision.
- Cleanliness may be worth considering for performance reasons.

## Locality

- Locality is a property of the page reference string. In other words, it is a property of programs themselves.
- *Temporal locality* says that pages that have been used recently are likely to be used again.
- Spatial locality says that pages "close" to those that have been used are likely to be used next.

In practice, page reference strings exhibit strong locality. Why?

CS350 Operating Systems Fall 2015

Virtual Memory 60

## Least Recently Used (LRU) Page Replacement

- LRU is based on the principle of temporal locality: replace the page that has not been used for the longest time
- To implement LRU, it is necessary to track each page's recency of use. For example: maintain a list of in-memory pages, and move a page to the front of the list when it is used.
- Although LRU and variants have many applications, true LRU is difficult to implement in virtual memory systems. (Why?)

## **Least Recently Used: LRU**

• the same three-frame example:

| Num     | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
|---------|---|---|---|---|---|---|---|---|---|----|----|----|
| Refs    | a | b | c | d | a | b | e | a | b | c  | d  | e  |
| Frame 1 | a | a | a | d | d | d | e | e | e | c  | c  | c  |
| Frame 2 |   | b | b | b | a | a | a | a | a | a  | d  | d  |
| Frame 3 |   |   | c | c | c | b | b | b | b | b  | b  | e  |
| Fault ? | X | X | X | X | X | X | X |   |   | х  | X  | X  |

from least recently used to most recently added:

a b c d a b e a b etc

problem: have to keep track of entire memory access ordering have to scan through and update list - this is super expensive

CS350 Operating Systems Fall 2015

Virtual Memory 62

## The "Use" Bit

- A use bit (or reference bit) is a bit found in each page table entry that:
  - is set by the MMU each time the page is used, i.e., each time the MMU translates a virtual address on that page these will eventually overflow we need to reset things every so often
  - can be read and cleared by the operating system
- The use bit provides a small amount of efficiently-maintainable usage information that can be exploited by a page replacement algorithm.

## The Clock Replacement Algorithm

- The clock algorithm (also known as "second chance") is one of the simplest algorithms that exploits the use bit.
- Clock is identical to FIFO, except that a page is "skipped" if its use bit is set. we give it a second chance we skip it and set the bit to unused
- The clock algorithm can be visualized as a victim pointer that cycles through the page frames. The pointer moves whenever a replacement is necessary:

```
while use bit of victim is set
   clear use bit of victim
  victim = (victim + 1) % num_frames
choose victim for replacement
victim = (victim + 1) % num frames
```

CS350 Operating Systems Fall 2015

Virtual Memory 64

## Page Cleanliness: the "Modified" Bit

- A page is *modified* (sometimes called dirty) if it has been changed since it was loaded into memory.
- A modified page is more costly to replace than a clean page. (Why?)
- The MMU identifies modified pages by setting a *modified bit* in page table entry of a page when a process *writes* to a virtual address on that page, i.e., when the page is changed.
- The operating system can clear the modified bit when it cleans the page
- The modified bit potentially has two roles:
   dirty bit tells us if we need to swap back out to disk when changing frame content
   Indicates which pages need to be cleaned. clean —> written back to disk
  - Can be used to influence the replacement policy, because it's slower

## **How Much Physical Memory Does a Process Need?**

- Principle of locality suggests that some portions of the process's virtual address space are more likely to be referenced than others.
- A refinement of this principle is the *working set model* of process reference behaviour.
- According to the working set model, at any given time some portion of a
  program's address space will be heavily used and the remainder will not be.
  The heavily used portion of the address space is called the *working set* of the
  process.
- The working set of a process may change over time.
- The resident set of a process is the set of pages that are located in memory.

According to the working set model, if a process's resident set includes its working set, it will rarely page fault. resident set = pages in memory, not on disk

CS350 Operating Systems Fall 2015

Virtual Memory 66

## **Resident Set Sizes (Example)**

resident set size - pages in memory (always <=RSZ)

```
PID
      VSZ
           RSS COMMAND
805 13940 5956 /usr/bin/gnome-session
831
           848 /usr/bin/ssh-agent
834
     7936 5832 /usr/lib/gconf2/gconfd-2 11
838
     6964 2292 gnome-smproxy
840 14720 5008 gnome-settings-daemon
848
     8412 3888 sawfish
851 34980 7544 nautilus
853 19804 14208 gnome-panel
857
     9656 2672 gpilotd
867
     4608 1252 gnome-name-service
     virtual
```

address size

(been allocated, but might not be in memory)

we could dynamically allocate only the pages needed at a given time

in os161 we don't do this - why? because it's easier (when this course had group work, they did do it)

## **Thrashing and Load Control**

• What is a good multiprogramming level?

- If too low: resources are idle

- If too high: too few resources per process

- A system that is spending too much time paging is said to be *thrashing*.
   Thrashing occurs when there are too many processes competing for the available memory.
   we keep going to disk disk is slow, this is slowww bad
- Thrashing can be cured by load shedding, e.g.,
  - Killing processes (not nice) e.g. too many tabs open, slow, close some tabs
  - Suspending and swapping out processes (nicer)

suspended state - a new state swap out all it's memory (including page table) to disk later swap back in and continue the kernel doesn't have to kill stuff necessarily if the user will close stuff when the system is getting slow

CS350 Operating Systems Fall 2015

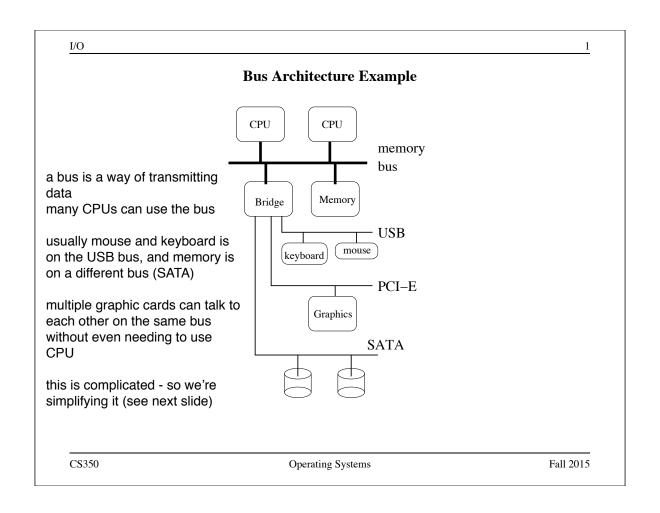
Virtual Memory 68

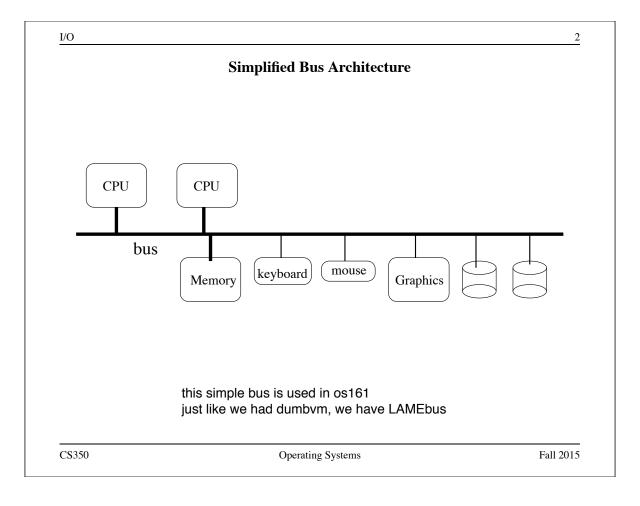
## **Swapping Out Processes**

- Swapping a process out means removing all of its pages from memory, or marking them so that they will be removed by the normal page replacement process. Suspending a process ensures that it is not runnable while it is swapped out.
- Which process(es) to suspend?
  - low priority processes



- blocked processes
- large processes (lots of space freed) or small processes (easier to reload)
- There must also be a policy for making suspended processes ready when system load has decreased.





## Sys/161 LAMEbus Device Examples

- timer/clock current time, timer, beep
- disk drive persistent storage
- serial console character input/output
- text screen character-oriented graphics
- network interface packet input/output

CS350 Operating Systems Fall 2015

I/O

## Device Register Example: Sys/161 timer/clock

## timer and clock have these registers available

| 0 4  | 4 | status             | current time (seconds)                 |                     |                              |
|------|---|--------------------|----------------------------------------|---------------------|------------------------------|
| 4 4  | 4 | status             | current time (nanoseconds)             |                     | onds to find<br>in nancoseco |
| 8 4  | 4 | command            | restart-on-expiry timer - when hits 0, | restarts            |                              |
| 12 4 | 4 | status and command | interrupt (reading clears) when you    | u read it, th       | at clears it                 |
| 16 4 | 4 | status and command | countdown time (microseconds)          | how much before ite |                              |
| 20 4 | 4 | command            | speaker (causes beeps)                 | 20.010 110          |                              |

YOU CAN PITCH THE BEEPS AND WRITE MUSIC WITH IT :0000

tal time in nancoseconds starts ead it, that clears it ow much time efore iterrupt

4

## Device Register Example: Sys/161 disk controller

| Offset | Size | Туре               | Description                    |                     |     |
|--------|------|--------------------|--------------------------------|---------------------|-----|
| 0      | 4    | status             | number of sectors              |                     |     |
| 4      | 4    | status and command | status find out if it's finish | ed reading or writi | ing |
| 8      | 4    | command            | sector number                  |                     |     |
| 12     | 4    | status             | rotational speed (RPM)         |                     |     |
| 32768  | 512  | data               | transfer buffer                |                     |     |

5

not 4 bytes

size is sector size

a disk is an array of sectors sector is smallest unit you can read or write to disk usually 512 bytes

CS350 Operating Systems Fall 2015

I/O 6

## **Device Drivers**

- a device driver is a part of the kernel that interacts with a device
- example: write character to serial output device

```
write character to device data register
write output command to device command register
repeat {
    read device status register
```

} until device status is ''completed''
clear the device status register

• this example illustrates *polling*: the driver repeatedly checks whether the device is finished, until it is finished.

## **Another Polling Example**

```
write target sector number into sector number register
write output data (512 bytes) into transfer buffer
write ''write'' command into status register
repeat {
   read status register
} until status is ''completed'' (or error)
clear the status register
```

usually we write 4 bytes at a time - that's a lot of writes Disk operations are slow. The device driver may have to poll for a long time.

CS350 Operating Systems Fall 2015

I/O 8

## **Using Interrupts to Avoid Polling**

- polling can be avoided if the device can use interrupts to indicate that it is finished
- example: disk write operation using interrupts:

write target sector number into sector number register write output data (512 bytes) into transfer buffer write ''write'' command into status register block until device generates completion interrupt read status register to check for errors clear status register

- while thread running the driver is blocked, the CPU is free to run other threads
- kernel synchronization primitives (e.g., semaphores) can be used to implement blocking

or like, a wait channel

Fall 2015 CS350 Operating Systems

#### **Device Data Transfer**

- Sometimes, a device operation will involve a large chunk of data much larger than can be moved with a single instruction.
  - example: disk read or write operation
- Devices may have data buffers for such data but how to get the data between the device and memory?
  - Option 1: program-controlled I/O PIO
     The device driver moves the data iteratively, one word at a time.
    - \* Simple, but the CPU is busy while the data is being transferred.
  - Option 2: direct memory access (DMA)
    - \* CPU is *not busy* during data transfer, and is free to do something else.

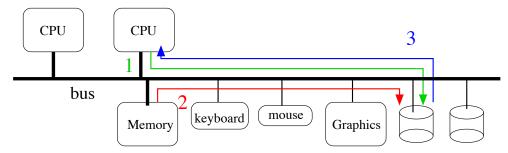
Sys/161 LAMEbus devices do program-controlled I/O.

CS350 Operating Systems Fall 2015

I/O 10

## **Direct Memory Access (DMA)**

- DMA is used for block data transfers between devices (e.g., a disk controller) and memory
- Under DMA, the CPU initiates the data transfer and is notified when the transfer is finished. However, the device (not the CPU) controls the transfer itself.



1. CPU issues DMA request to controller

'please copy this memory whenever you next get control of the bus'

- 2. controller directs data transfer
- 3. controller interrupts CPU

#### **Device Driver for Disk Write with DMA**

write target disk sector number into sector number register write source memory address into address register write ''write'' command into status register block (sleep) until device generates completion interrupt read status register to check for errors clear status register

Note: driver no longer copies data into device transfer buffer before the CPU wrote to a buffer

now the DMA is writing to a buffer that isn't accessible to the CPU so we don't call it a transfer buffer - it happens in a different place

CS350 Operating Systems Fall 2015

I/O 12

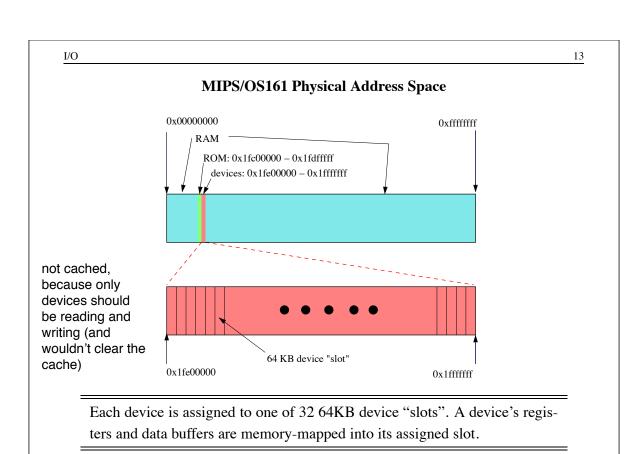
#### **Accessing Devices**

- how can a device driver access device registers?
- Option 1: special I/O instructions
  - such as in and out instructions on x86
  - device registers are assigned "port" numbers
  - instructions transfer data between a specified port and a CPU register

\_

- Option 2: memory-mapped I/O
  - each device register has a physical memory address
  - device drivers can read from or write to device registers using normal load and store instructions, as though accessing memory

like in 241



Operating Systems

Fall 2015

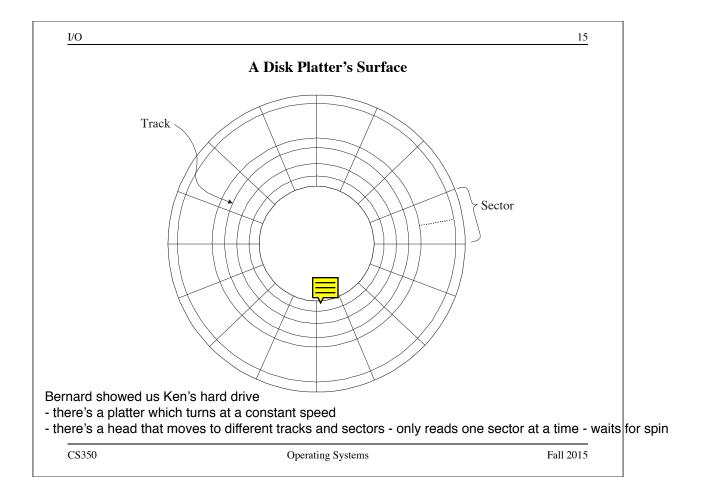
## Logical View of a Disk Drive

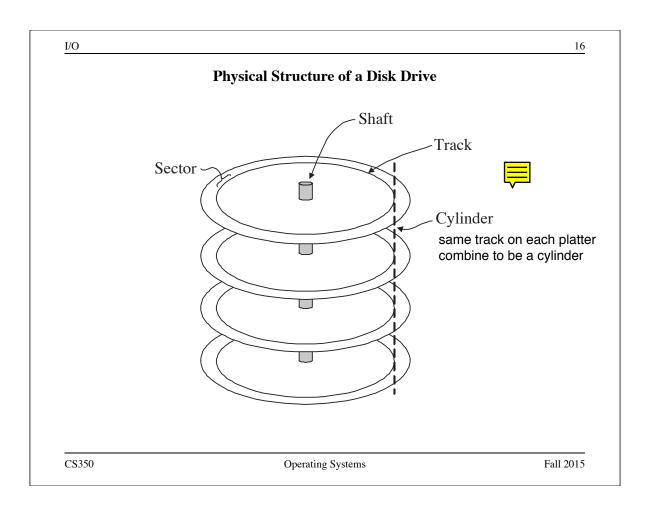
- disk is an array of numbered blocks (or sectors)
- each block is the same size (e.g., 512 bytes)

CS350

- blocks are the unit of transfer between the disk and memory
  - typically, one or more contiguous blocks can be transferred in a single operation
- storage is non-volatile, i.e., data persists even when the device is without power

RAM is volatile but if you turn off your computer, turn it back on - data is still there





## **Simplified Cost Model for Disk Block Transfer**

• moving data to/from a disk involves:

seek time: move the read/write heads to the appropriate cylinder

rotational latency: wait until the desired sectors spin to the read/write heads

**transfer time:** wait while the desired sectors spin past the read/write heads e.g. 9 sectors, want 3 in a row, 1s to do a full rotation —> 1/3 s transfer time

• request service time is the sum of seek time, rotational latency, and transfer time

$$t_{service} = t_{seek} + t_{rot} + t_{transfer}$$

• note that there are other overheads but they are typically small relative to these three

CS350 Operating Systems Fall 2015

I/O 18

## **Rotational Latency and Transfer Time**

- rotational latency depends on the rotational speed of the disk
- if the disk spins at  $\omega$  rotations per second: spin speed is measured in rpm (rot/min) 7200 higher end

$$0 \leq t_{rot} \leq \frac{1}{\omega}$$
 1/omega is worst case (full spin)

• expected rotational latency:

$$ar{t}_{rot} = rac{1}{2\omega}$$
 average (half the max)

- transfer time depends on the rotational speed and on the amount of data transferred
- if k sectors are to be transferred and there are T sectors per track:

$$t_{transfer} = rac{k}{T\omega}$$
 k/T is fraction of total sectors

#### **Seek Time**

- seek time depends on the speed of the arm on which the read/write heads are mounted.
- a simple linear seek time model:
  - $t_{maxseek}$  is the time required to move the read/write heads from the innermost cylinder to the outermost cylinder
  - C is the total number of cylinders
- if k is the required seek distance (k > 0):

$$t_{seek}(k) = \frac{k}{C} t_{maxseek}$$

CS350 Operating Systems Fall 2015

I/O 20

## **Performance Implications of Disk Characteristics**

- larger transfers to/from a disk device are *more efficient* than smaller ones. That is, the cost (time) per byte is smaller for larger transfers. (Why?)

  don't have to seek again
- sequential I/O is faster than non-sequential I/O
  - sequential I/O operations eliminate the need for (most) seeks
  - disks use other techniques, like *track buffering*, to reduce the cost of sequential I/O even more
     once you get to the track you need read more sectors than specified, and keep in a buffer somewhere so we can get that

in a buffer somewhere so we can get that data later faster than doing another seek

small increase in transfer time, saves on potential later seek time and rotational latency

spatial locality - the idea that we tend to look up things close to each other pretty often

I/O 21

### **Disk Head Scheduling**

- goal: reduce seek times by controlling the order in which requests are serviced
- disk head scheduling may be performed by the controller, by the operating system, or both
- for disk head scheduling to be effective, there must be a queue of outstanding disk requests (otherwise there is nothing to reorder)
- an on-line approach is required: the disk request queue is not static

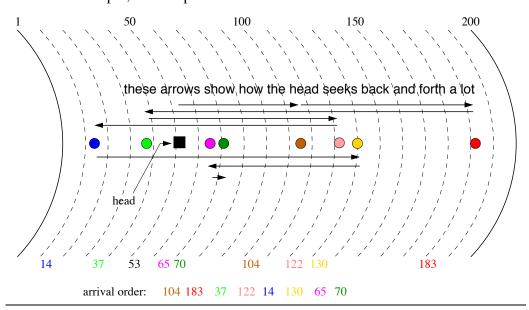
CS350 Operating Systems Fall 2015

I/O 22

### **FCFS Disk Head Scheduling**

pro: fair (no starvation) — con: slow/long jobs prevent short jobs from going quickly

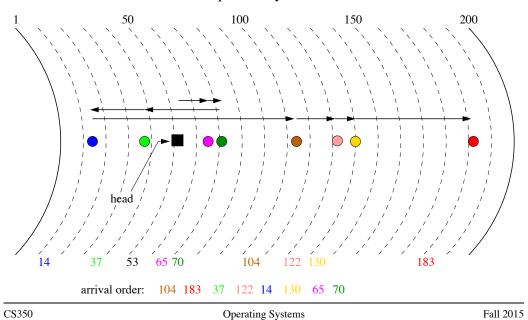
- handle requests in the order in which they arrive
- fair and simple, but no optimization of seek times



I/O 23

## **Shortest Seek Time First (SSTF)**

- choose closest request (a greedy approach)
- seek times are reduced, but requests may starve

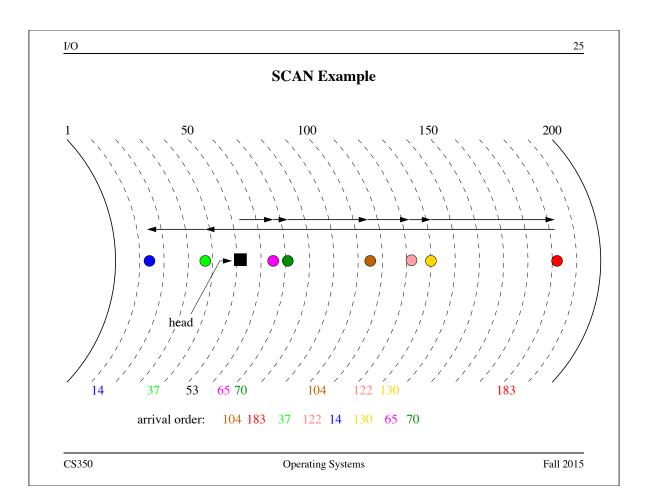


I/O 24

## **Elevator Algorithms (SCAN)**

- Under SCAN, aka the elevator algorithm, the disk head moves in one direction until there are no more requests in front of it, then reverses direction.
- there are many variations on this idea
- SCAN reduces seek times (relative to FCFS), while avoiding starvation

like a standard elevator



#### chat about SSDs

they way they work is instead of having a physical disk that spins, it uses transistors and is like DRAM on the machine — access latency is equal anywhere, no seeking

fun story - Bernard saw this old OS textbook with tape drives

- like audio tape, video tape "be kind, rewind"
- you fast forward, read data takes tens of seconds
- the textbook was talking about how to be more efficient in how you lay out memory on the tape, he laughed about how outdated and ridiculous and useless it was, and then realized that hard drives and what we talked about on the last few slides will also soon be outdated

DRAM - use transistors - periodically reads and rewrites data to memory to ensure it is still there and hasn't degraded

SSDs use floating gate transistors (you don't need to know this) creates a quantum cage that keeps the charge indefinitely - even if you turn it off and don't refresh

1 to 0 is really easy, 0 to 1 is really hard and requires high voltage and is hard to do to only one bit, you can only update a region — so to make changes you read off the SSD, set them all to 1, then write back by changing things to 0s — this will make small changes big "write amplification"

every time you write, it creates a new segment (updates a clear section), and cleans up the old one when the disk is idle (resets it without time crunch)

(none of this is exam material)

### **Job Scheduling Model**

- problem scenario: a set of *jobs* needs to be executed using a single server, on which only one job at a time may run
- for the ith job, we have an arrival time  $a_i$  and a run time  $r_i$
- after the *i*th job has run on the server for total time  $r_i$ , it finishes and leaves the system
- a job *scheduler* decides which job should be running on the server at each point in time
- let  $s_i$  ( $s_i \ge a_i$ ) represent the time at which the *i*th job first runs, and let  $f_i$  represent the time at which the *i*th job finishes
  - the turnaround time of the ith job is  $f_i a_i$
  - the response time of the ith job is  $s_i a_i$

CS350 Operating Systems Fall 2015

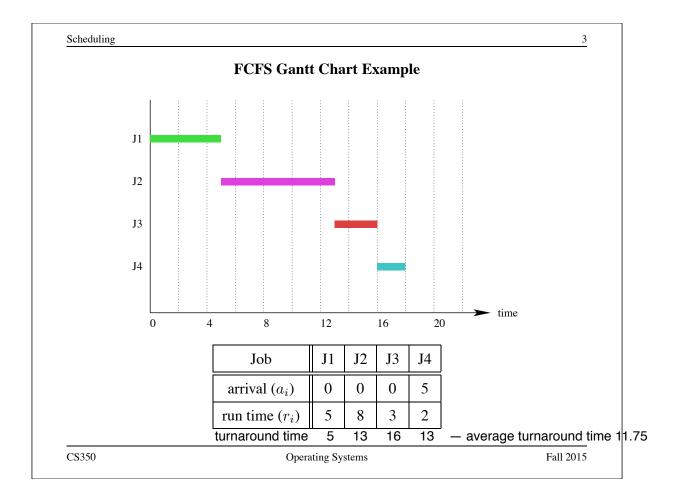
Scheduling 2

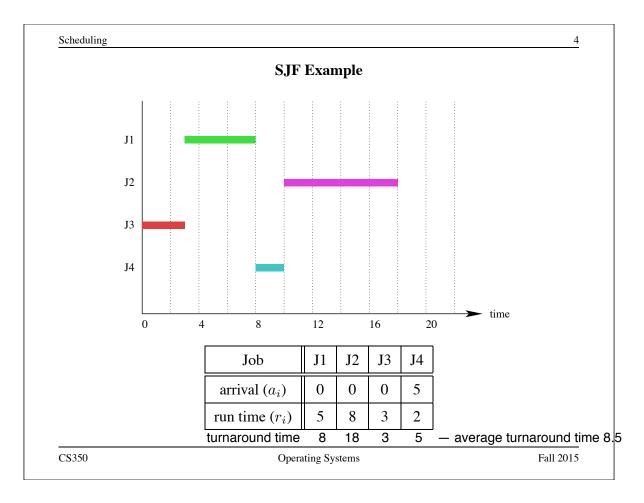
### **Basic Non-Preemptive Schedulers: FCFS and SJF**

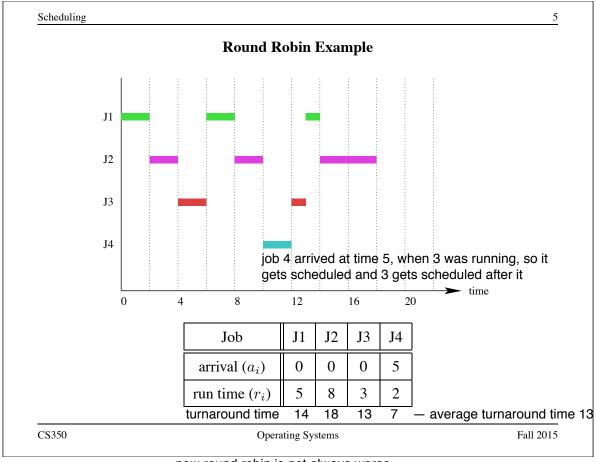
- FCFS: runs jobs in arrival time order.
  - simple, avoids starvation
  - pre-emptive variant: round-robin a job only runs for so long (quantum) and then we switch
- SJF: shortest job first run jobs in increasing order of  $r_i$ 
  - minimizes average turnaround time
  - long jobs may starve
  - pre-emptive variant: SRTF (shortest remaining time first)
     if a running job has more time remaining than an incoming job's time,
     we stop it and let the smallest job go

example: grocery store

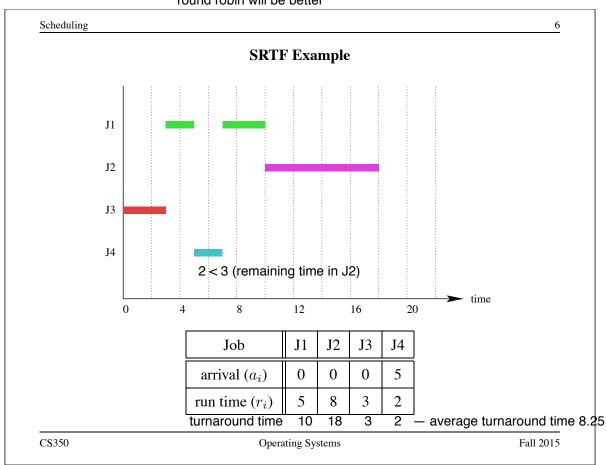
only one place to check out we have a stick of gum, and someone is buying a tonnnn of groceries do we ask to go first? do we let them go before us? just leave?







now round robin is not always worse if a really long process is scheduled first, and short ones come after, round robin will be better



### **CPU Scheduling**

- CPU scheduling is job scheduling where:
  - the server is a CPU (or a single core of a multi-core CPU)
  - the jobs are *ready threads*

arriving - transitions to ready starts - transitions to running

- \* a thread "arrives" when it becomes ready, i.e., when it is first created, or when it wakes up from sleep
- \* the run-time of the thread is the amount of time that it will run before it either finishes or blocks
- thread run times are typically not known in advance by the scheduler you could run it for a bit and try to guess how long it'll take
- typical scheduler objectives
  - responsiveness low *response time* for some or all threads show up 2s after you type it
  - "fair" sharing of the CPU
  - efficiency there is a cost to switching there's a lot of overhead - it takes time to switch so that will become slow

CS350 Operating Systems Fall 2015

Scheduling 8

#### **Prioritization**

- CPU schedulers are often expected to consider process or thread priorities
- priorities may be
  - specified by the application (e.g., Linux setpriority/sched\_setscheduler)

can set its own priority
needs special permissions, so random
applications can't just say they're important

- chosen by the scheduler
- some combination of these
- two approaches to scheduling with priorites
  - 1. schedule the highest priority thread
  - 2. weighted fair sharing
    - let  $p_i$  be the priority of the *i*th thread
    - try to give each thread a "share" of the CPU in proportion to its priority:

$$\frac{p_i}{\sum_j p_j} \tag{1}$$

### **Multi-level Feedback Queues**

- objective: good responsiveness for *interactive* processes or else user gets mad
  - threads of interactive processes block frequently, have short run times
- idea: gradually diminish priority of threads with long run times and infrequent blocking
  - if a thread blocks before its quantum is used up, raise its priority
  - if a thread uses its entire quantum, *lower* its priority

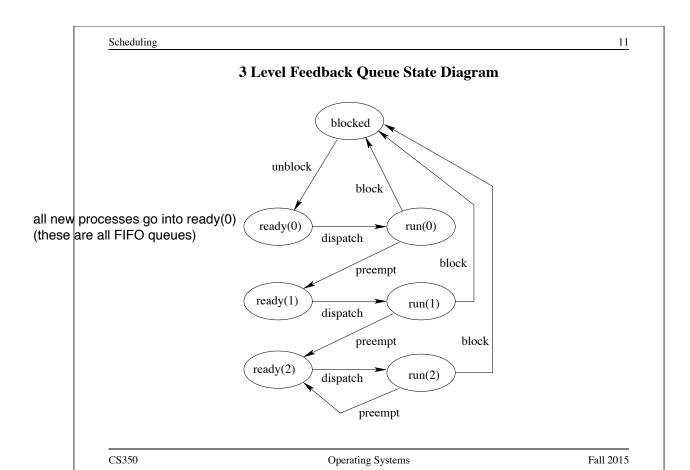
CS350 Operating Systems Fall 2015

Scheduling 10

### **Multi-level Feedback Queues (Algorithm)**

- scheduler maintains several round-robin ready queues
  - highest priority threads in queue  $Q_0$ , lower priority in  $Q_1$ , still lower in  $Q_2$ , and so on.
- scheduler always chooses thread from the lowest non-empty queue
- threads in queue  $Q_i$  use quantum  $q_i$ , and  $q_i \leq q_j$  if i < j
- newly ready threads go into ready queue  $Q_0$
- a level i thread that is preempted goes into queue  $Q_{i+1}$

This basic algorithm may *starve* threads in lower queues. Various enhancements can avoid this, e.g, periodically migrate all threads into  $Q_0$ .

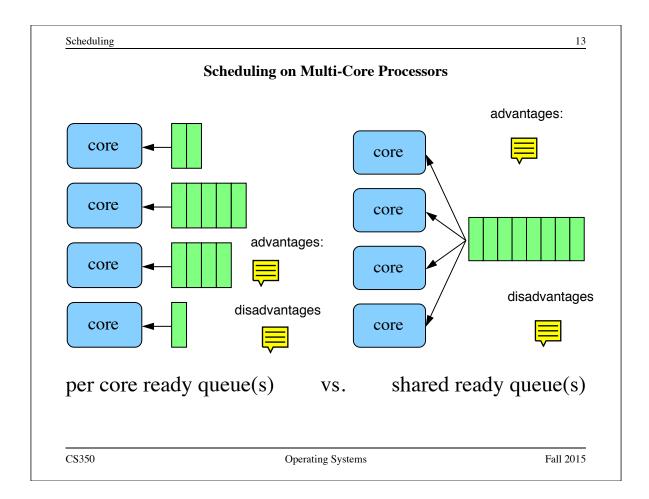


**Nov 12** 

### Linux Completely Fair Scheduler (CFS) - Key Ideas

\*\*\*completely\*\*\* (a strong word!)

- "Completely Fair Scheduling" a weighted fair sharing approach almost exactly like round robin
- suppose that  $c_i$  is the actual amount of time that the scheduler has allowed the *i*th thread to run.
- on an *ideally shared* processor, we would expect  $c_0 \frac{\sum_j p_j}{p_0} = c_1 \frac{\sum_j p_j}{p_1} = \cdots$
- CFS calls  $c_i \frac{\sum_j p_j}{p_i}$  the *virtual runtime* of the *i*th thread, and tracks it for each thread
- CFS chooses the thread with the lowest virtual runtime, and runs it until some other thread's virtual runtime is lower (subject to a minimum runtime quantum)
  - virtual runtime advances more slowly for higher priority threads, so they get longer time slices
  - all ready threads run regularly, so good responsiveness



### **Scalability and Cache Affinity**

- Contention and Scalability
  - access to shared ready queue is a critical section, mutual exclusion needed
  - as number of cores grows, contention for ready queue becomes a problem
  - per core design scales to a larger number of cores
- CPU cache affinity
  - as thread runs, data it accesses is loaded into CPU cache(s)
  - moving the thread to another core means data must be reloaded into that core's caches
  - as thread runs, it acquires an *affinity* for one core because of the cached data
  - per core design benefits from affinity by keeping threads on the same core
  - shared queue design does not

# **Load Balancing**

- in per-core design, queues may have different lengths
- this results in *load imbalance* across the cores
  - cores may be idle while others are busy
  - threads on lightly loaded cores get more CPU time than threads on heavily loaded cores
- not an issue in shared queue design
- per-core designs typically need some mechanism for *thread migration* to address load imbalances
  - migration means moving threads from heavily loaded cores to lightly loaded cores

| CS350     | O                 | Fall 2015 |
|-----------|-------------------|-----------|
| (.5.2.20) | Operating Systems | Fall ZULD |
|           |                   |           |

### Files and File Systems

abstraction for managing data

- files: persistent, named data objects
  - data consists of a sequence of numbered bytes
  - file may change size over time
  - file has associated meta-data
    - examples: owner, access controls, file type, creation and access
       timestamps
       name is actually meta-data of the directory
- file system: a collection of files which share a common name space
  - allows files to be created, destroyed, renamed, . . .

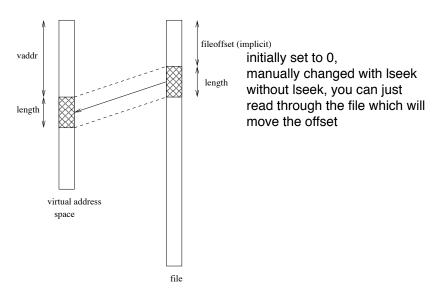
CS350 Operating Systems Fall 2015

File Systems 2

#### **File Interface**

- open, close
  - open returns a file identifier (or handle or descriptor), which is used in subsequent operations to identify the file. (Why is this done?)
- read, write, seek
  - read copies data from a file into a virtual address space
  - write copies data from a virtual address space into a file
  - seek enables non-sequential reading/writing
- get/set file meta-data, e.g., Unix fstat, chmod stats on size, permissions

### File Read



read(fileID, vaddr, length)

CS350 Operating Systems Fall 2015

File Systems 4

### **File Position**

- each file descriptor (open file) has an associated file position
- read and write operations
  - start from the current file position
  - update the current file position
- this makes sequential file I/O easy for an application to request
- for non-sequential (random) file I/O, use:
  - a seek operation (lseek) to adjust file position before reading or writing
  - a positioned read or write operation, e.g., Unix pread, pwrite:
     pread(fileId, vaddr, length, filePosition)

### **Sequential File Reading Example (Unix)**

Read the first 100 \* 512 bytes of a file, 512 bytes at a time.

CS350 Operating Systems Fall 2015

File Systems 6

### File Reading Example Using Seek (Unix)

```
char buf[512];
int i;
int f = open("myfile",O_RDONLY);
for(i=1; i<=100; i++) {
  lseek(f,(100-i)*512,SEEK_SET); reading 512 bytes at a time backwards so we have to seek
  read(f,(void *)buf,512);
}
close(f);</pre>
```

Read the first 100 \* 512 bytes of a file, 512 bytes at a time, in reverse order.

### File Reading Example Using Positioned Read

```
char buf[512];
int i;
int f = open("myfile",O_RDONLY);
for(i=0; i<100; i+=2) {
   pread(f,(void *)buf,512,i*512);
}
close(f);</pre>
read uses implicit offset pread specifies position
```

Read every second 512 byte chunk of a file, until 50 have been read.

warning - these code snippets don't have any checks always check return values! (reading, malloc, etc)

CS350 Operating Systems Fall 2015

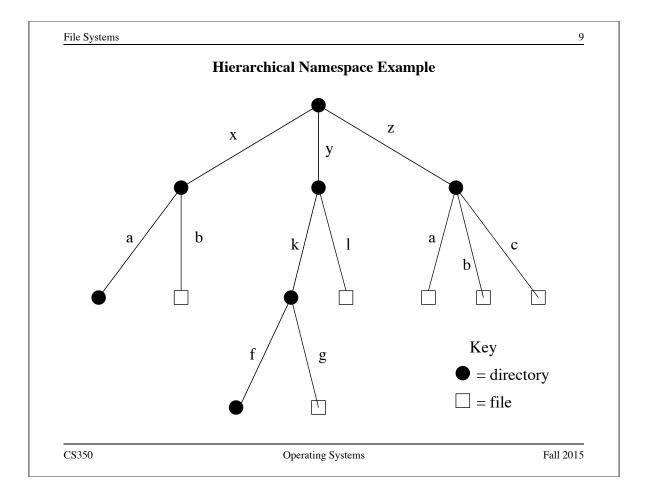
File Systems 8

#### **Directories and File Names**

- A directory maps file names (strings) to i-numbers
  - an i-number is a unique (within a file system) identifier for a file or directory
  - given an i-number, the file system can find the data and meta-data for the file
- Directories provide a way for applications to group related files
- Since directories can be nested, a filesystem's directories can be viewed as a tree, with a single *root* directory.
- In a directory tree, files are leaves
- Files may be identified by *pathnames*, which describe a path through the directory tree from the root directory to the file, e.g.:

/home/user/courses/cs350/notes/filesys.pdf

- Directories also have pathnames
- Applications refer to files using pathnames, not i-numbers



#### **Hard Links**

- a hard link is an association between a name (string) and an i-number
  - each entry in a directory is a hard link
- when a file is created, so is a hard link to that file
  - open(/a/b/c,O\_CREAT|O\_TRUNC)
  - this creates a new file if a file called /a/b/c does not already exist
  - it also creates a hard link to the file in the directory /a/b
- Once a file is created, *additional* hard links can be made to it.
  - example: link(/x/b,/y/k/h) creates a new hard link h in directory /y/k. The link refers to the i-number of file /x/b, which must exist.
- linking to an existing file creates a new pathname for that file
  - each file has a unique i-number, but may have multiple pathnames
- Not possible to link to a directory (to avoid cycles)

### **Unlinking and Referential Integrity**

- hard links can be removed:
  - unlink(/x/b)
- the file system ensures that hard links have *referential integrity*, which means that if the link exists, the file that it refers to also exists.
  - When a hard link is created, it refers to an existing file.
  - There is no system call to delete a file. Instead, a file is deleted when its last hard link is removed.

CS350 Operating Systems Fall 2015

File Systems 12

### Symbolic Links

- a *symbolic link*, or *soft link*, is an association between a name (string) and a pathname.
  - symlink(/z/a,/y/k/m) creates a symbolic link m in directory /y/k.
     The symbolic link refers to the pathname /z/a.
- If an application attempts to open /y/k/m, the file system will
  - 1. recognize /y/k/m as a symbolic link, and
  - 2. attempt to open /z/a instead
- referential integrity is *not* preserved for symbolic links
  - in the example above, /z/a need not exist!

### **UNIX/Linux Link Example (1 of 3)**

```
% cat > file1
This is file1.
<cntl-d>
% ls -li
685844 -rw------ 1 user group 15 2008-08-20 file1
% ln file1 link1
% ln -s file1 sym1
% ln not-here link2
ln: not-here: No such file or directory
% ln -s not-here sym2
```

Files, hard links, and soft/symbolic links.

CS350 Operating Systems Fall 2015

File Systems 14

### **UNIX/Linux Link Example (2 of 3)**

```
% ls -li
685844 -rw------ 2 user group 15 2008-08-20 file1
685844 -rw----- 2 user group 15 2008-08-20 link1
685845 lrwxrwxrwx 1 user group 5 2008-08-20 sym1 -> file1
685846 lrwxrwxrwx 1 user group 8 2008-08-20 sym2 -> not-here
% cat file1
This is file1.
% cat link1
This is file1.
% cat sym1
This is file1.
% cat sym2
cat: sym2: No such file or directory
% /bin/rm file1
```

Accessing and manipulating files, hard links, and soft/symbolic links.

### **UNIX/Linux Link Example (3 of 3)**

```
% ls -li
685844 -rw----- 1 user group 15 2008-08-20 link1
685845 lrwxrwxrwx 1 user group 5 2008-08-20 sym1 -> file1
685846 lrwxrwxrwx 1 user group 8 2008-08-20 sym2 -> not-here
% cat link1
This is file1.
% cat sym1
cat: sym1: No such file or directory
% cat > file1
This is a brand new file1.
<cntl-d>
% ls -li
685847 -rw----- 1 user group 27 2008-08-20 file1
685844 -rw----- 1 user group 15 2008-08-20 link1
685845 lrwxrwxrwx 1 user group 5 2008-08-20 sym1 -> file1
685846 lrwxrwxrwx 1 user group 8 2008-08-20 sym2 -> not-here
% cat link1
This is file1.
% cat sym1
This is a brand new file1.
```

Different behaviour for hard links and soft/symbolic links.

CS350 Operating Systems Fall 2015

File Systems 16

#### **Multiple File Systems**

- it is not uncommon for a system to have multiple file systems
- some kind of global file namespace is required
- two examples:

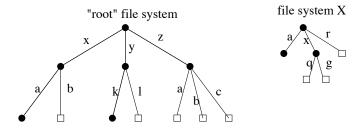
**DOS/Windows:** use two-part file names: file system name, pathname within file system

- example: C:\user\cs350\schedule.txt

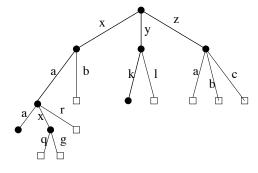
**Unix:** create single hierarchical namespace that combines the namespaces of two file systems

- Unix mount system call does this
- mounting does *not* make two file systems into one file system
  - it merely creates a single, hierarchical namespace that combines the namespaces of two file systems
  - the new namespace is temporary it exists only until the file system is unmounted

### Unix mount Example



result of mount (file system X, /x/a)



CS350 Operating Systems Fall 2015

File Systems 18

### **Links and Multiple File Systems**

- hard links cannot cross file system boundaries
  - each hard link maps a name to an i-number, which is unique only within a file system
- for example, even after the mount operation illustrated on the previous slide, link(/x/a/x/g,/z/d) would result in an error, because the new link, which is in the root file system refers to an object in file system X
- soft links do not have this limitation
- for example, after the mount operation illustrated on the previous slide:
  - symlink(/x/a/x/g,/z/d) would succeed
  - open(/z/d) would succeed, with the effect of opening /z/a/x/g.
- even if the symlink operation were to occur before the mount command, it would succeed

### **File System Implementation**

- what needs to be stored persistently?
  - file data
  - file meta-data
  - directories and links
  - file system meta-data
- non-persistent information
  - open files per process
  - file position for each open file
  - cached copies of persistent data

CS350 Operating Systems Fall 2015

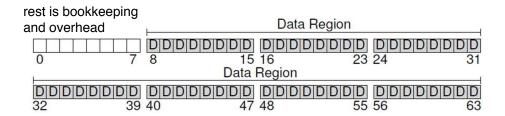
File Systems 20

#### File System Example

- Use an extremely small disk as an example:
  - 256 KB disk!
  - Most disks have a sector size of 512 bytes
    - \* Memory is usually byte addressable
    - \* Disk is usually "sector addressable"
  - 512 total sectors on this disk
- Group every 8 consecutive sectors into a block
  - Better spatial locality (fewer seeks)
     therefore better performance
  - Reduces the number of block pointers (we'll see what this means soon)
  - 4 KB block is a convenient size for demand paging
  - 64 total blocks on this disk

### **VSFS: Very Simple File System (1 of 5)**

• Most of the blocks should be for storing user data (last 56 blocks)



CS350 Operating Systems Fall 2015

File Systems 22

### **VSFS: Very Simple File System (2 of 5)**

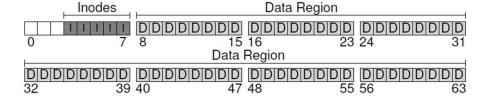
any nonzero sized file has at least

Need some way to map files to data blocks one data block

each file has its own i-node

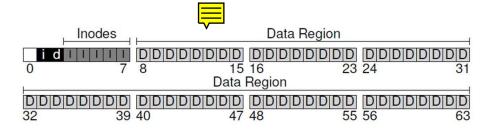


- Create an array of i-nodes, where each i-node contains the meta-data for a file
  - The index into the array is the file's index number (i-number) i-number index tells us which i-node we want
- Assume each i-node is 256 bytes, and we dedicate 5 blocks for i-nodes
  - This allows for 80 total i-nodes/files



### **VSFS: Very Simple File System (3 of 5)**

- We also need to know which i-nodes and blocks are unused
- Many ways of doing this:
  - In VSFS, we use a bitmap for each i-node bitmap each bit represents an i-node 0 means available, 1 means used
  - Can also use a free list instead of a bitmap
- A block size of 4 KB means we can track 32K i-nodes and 32K blocks
  - This is far more than we actually need



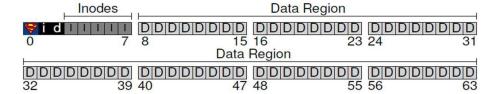
CS350 Operating Systems Fall 2015

File Systems 24

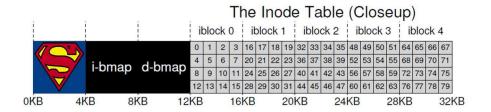
### **VSFS: Very Simple File System (4 of 5)**

- Reserve the first block as the **superblock**
- A superblock contains meta-information about the entire file system
  - e.g., how many i-nodes and blocks are in the system, where the i-node table begins, etc.

#### file system meta-data



### **VSFS: Very Simple File System (5 of 5)**



CS350 Operating Systems Fall 2015

File Systems 26

#### i-nodes

 An i-node is a *fixed size* index structure that holds both file meta-data and a small number of pointers to data blocks
 a directory is a special file that maps file names to inumbers

• i-node fields include:

file type

- file permissions

- file length

- number of file blocks

- time of last file access

i-node (what's stored - diagram on the board)

[ ] attributes
[ ] more on
[ ] direct pointers pointers/indirect on
[ ] single indirect next slides

] double indirect

- time of last i-node update, last file update

 number of hard links to this file it's possible to have multiple file names pointing to same file

- direct data block pointers

- single, double, and triple indirect data block pointers

when the #of hardlinks is decremented to 0, we free this i-node and it's blocks

### VSFS: i-node

- Assume disk blocks can be referenced based on a 4 byte address
  - 2<sup>32</sup> blocks, 4 KB blocks
  - Maximum disk size is 16 TB
- In VSFS, an i-node is 256 bytes
  - Assume there is enough room for 12 direct pointers to blocks
  - Each pointer points to a different block for storing user data
  - Pointers are ordered: first pointer points to the first block in the file, etc.
- What is the maximum file size if we only have direct pointers?
  - -12\*4 KB = 48 KB
- Great for small files (which are common)
- Not so great if you want to store big files
   we'd have to break it up into a bunch of smaller files
   good for small files because all we have to do is read inode, and we get a block of data

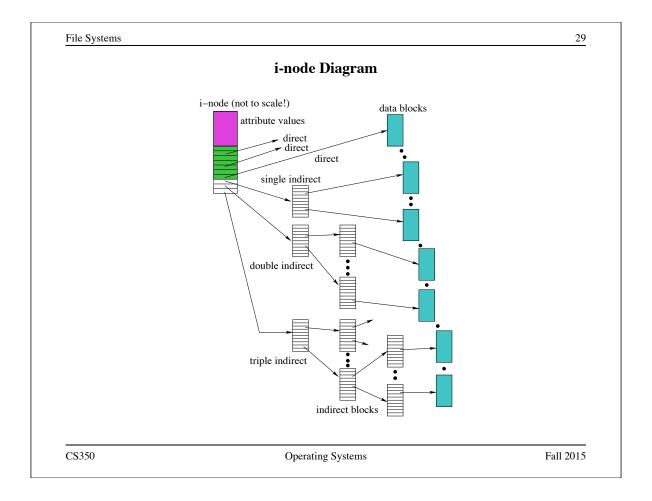
CS350 Operating Systems Fall 2015

File Systems 28

#### **VSFS: Indirect Blocks**

helps us have big files

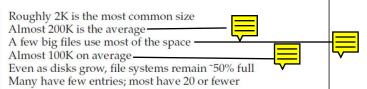
- In addition to 12 direct pointers, we can also introduce an **indirect pointer** 
  - An indirect pointer points to a block full of direct pointers
- 4 KB block of direct pointers = 1024 pointers
  - Maximum file size is: (12 + 1024) \* 4 KB = 4144 KB
- Better, but still not enough
- Add a double indirect pointer
  - Points to a 4 KB block of indirect pointers
  - (12 + 1024 + 1024 \* 1024) \* 4 KB
  - Just over 4 GB in size (is this enough?)



### **File System Design**

- File system parameters:
  - How many i-nodes should a file system have?
  - How many direct and indirect blocks should an i-node have?
  - What is the "right" block size?
- For a general purpose file system, design it to be efficient for the common case

Most files are small Average file size is growing Most bytes are stored in large files File systems contains lots of files File systems are roughly half full Directories are typically small





#### **Directories**

- Implemented as a special type of file.
- Directory file contains directory entries, each consisting of
  - a file name (component of a path name) and the corresponding i-number

names must be unique

| name   | i-number |
|--------|----------|
|        | 5        |
| ••     | 2        |
| foo    | 12       |
| bar    | 13       |
| foobar | 24       |

even root will have .. either the value will be missing, or point to itself, or whatever

- Directory files can be read by application programs (e.g., 1s) lists file names
- Directory files are only updated by the kernel, in response to file system operations, e.g, create file, create link
- Application programs cannot write directly to directory files. (Why not?)



CS350 Operating Systems Fall 2015

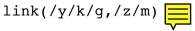
it's really inefficient if we have a lot of files inside a directory we'd have to scan through everyything - storing it in a tree (maybe btree) would be better

File Systems 32

### **Implementing Hard Links**

- hard links are simply directory entries
- hard miks are simply directory entires

• for example, consider:



• to implement this:

when we map a file name to an i-number that's a hardlink to the i-number

when hardlink count goes to 0, nothing points to it so we can free the file and reuse the data block(s) and i-node

- to implement this.
  - 1. find out the internal file identifier for /y/k/g
  - 2. create a new entry in directory / z
    - file name in new entry is m
    - file identifier (i-number) in the new entry is the one discovered in step 1

### **Implementing Soft Links**

• soft links can be implemented as a special type of file

• for example, consider: can read and write to m, and will modify g so far, same as hardlink

symlink(/y/k/g,/z/m)

BUT if we delete g, and try to access through m, we get an error

• to implement this:

- create a new symlink file

- add a new entry in directory / z
  - \* file name in new entry is m
  - \* i-number in the new entry is the i-number of the new symlink file
- store the pathname string "/y/k/g" as the contents of the new symlink file

hardlink has only one i-node and both files map to same i-number

symlink - two inodes - one for file, and one for symlink but the data block for symlink m is just the absolute path /y/k/g

CS350 Operating Systems Fall 2015

you can't have hard links across multiple file systems - each file system has different i-numbers symlinks can - because they just use file path

File Systems 34

#### Free Space Management

- Use the bitmaps to find a free i-node and free blocks
  - Each bit represents the availability of an i-node or block
- There are often many free blocks to choose from
  - To improve spatial locality and reduce fragementation, a file system may want to select a free block that is followed by a sequence of other free blocks

e.g. we allocate first file in first block

next file - put somewhere a bit further down (not second data block) so the first file has some room to grow

there are bad cases for any choice of next block (resulting in fragmentation)

## Reading From a File (/foo/bar)

- First read the root i-node
  - At "well known" position (i-node 2) root is often in i-node 2
  - i-node 1 is usually for tracking bad locks typo this should say 'bad blocks'

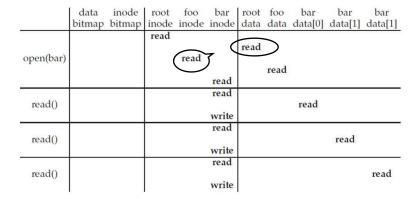
|           | data  | inode<br>bitmap |      | foo  |       |      |      |        | bar<br>data[1] | bar<br>data[1] |
|-----------|-------|-----------------|------|------|-------|------|------|--------|----------------|----------------|
|           | оттар | оннар           | read |      | mode  | read | data | unujoj | data[1]        | cata[1]        |
| open(bar) |       |                 |      | read | read  |      | read |        |                |                |
|           |       |                 |      |      | read  |      |      | 671    |                |                |
| read()    |       |                 |      |      | write |      |      | read   |                |                |
| read()    |       |                 |      |      | read  |      |      |        | read           |                |
|           |       |                 |      |      | write |      |      |        | read           |                |
| read()    |       |                 |      |      | read  |      |      |        |                | and d          |
|           |       |                 |      |      | write |      |      |        |                | read           |

CS350 Operating Systems Fall 2015

File Systems 36

## Reading From a File (/foo/bar)

- Read the directory information from root
  - Find the i-number for foo
  - Read the foo i-node



## Reading From a File (/foo/bar)

- Read the directory information from foo
  - Find the i-number for bar
  - Read the bar i-node

|           | data   | inode  | root  | foo   |       | root       |         | bar     | bar     | bar     |
|-----------|--------|--------|-------|-------|-------|------------|---------|---------|---------|---------|
|           | bitmap | bitmap | inode | inode | inode | data       | data    | data[0] | data[1] | data[1] |
| *         |        |        | read  |       |       | 25         |         |         |         |         |
| 191 01    |        |        |       |       |       | read       |         |         |         |         |
| open(bar) |        |        |       | read  |       | l .        | _       | 7       |         |         |
|           |        |        |       |       |       | <b>人</b> ( | read    | /       |         |         |
|           |        |        |       | (     | read  | <u> </u>   | $\cong$ |         |         |         |
| 10        |        |        |       |       | read  |            |         |         |         |         |
| read()    |        |        |       |       |       |            |         | read    |         |         |
|           |        |        |       |       | write |            |         |         |         |         |
|           | ľ      |        |       |       | read  |            |         |         |         |         |
| read()    |        |        |       |       | 21    |            |         |         | read    |         |
|           |        |        |       |       | write |            |         |         |         |         |
|           |        |        |       |       | read  |            |         |         |         |         |
| read()    |        |        |       |       |       |            |         |         |         | read    |
|           |        |        |       |       | write |            |         |         |         |         |

CS350 Operating Systems Fall 2015

File Systems

## Reading From a File (/foo/bar)

we can't release a file when hardlink count goes to 0 what if it's open??

38

• Permission check (is the user allowed to read this file?)

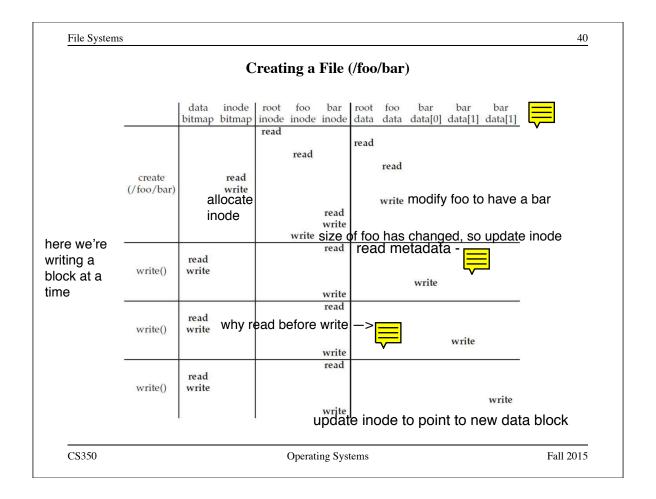
- Allocate a file descriptor in the per-process descriptor table
- Increment the counter for this i-number in the global open file table

|            |           | data   | inode  | root  | too   | bar   | root | too  | bar     | bar     | bar     |
|------------|-----------|--------|--------|-------|-------|-------|------|------|---------|---------|---------|
|            | 52        | bitmap | bitmap | inode | inode | inode | data | data | data[0] | data[1] | data[1] |
|            |           |        |        | read  |       |       |      |      |         |         |         |
|            | 1000      |        |        |       |       |       | read |      |         |         |         |
|            | open(bar) |        |        |       | read  |       |      |      |         |         |         |
|            |           |        |        |       |       |       |      | read |         |         |         |
|            | _         |        |        |       |       | read  |      |      |         |         |         |
|            | read()    |        |        |       |       | read  |      |      | 161     |         |         |
| once it's  |           |        |        |       |       |       |      |      | read    |         |         |
| open, we   | -         |        |        |       |       | write |      |      |         |         |         |
| read three |           |        |        |       |       | read  |      |      |         |         |         |
|            | read()    |        |        |       |       | 21    |      |      |         | read    |         |
| blocks     |           |        |        |       |       | write |      |      |         |         |         |
|            |           |        |        |       |       | read  |      |      |         |         |         |
|            | read()    |        |        |       |       |       |      |      |         |         | read    |
|            |           |        |        | l     |       | write |      |      |         |         |         |

### Reading From a File (/foo/bar)

- Find the block using a direct/indirect pointer and read the data
- Update the i-node with a new access time
- Update the file position in the per-process descriptor table
- Closing a file deallocates the file descriptor and decrements the counter for this i-number in the global open file table

|           | data   | inode  |       |       | bar   |      |      |         | bar     | bar     |
|-----------|--------|--------|-------|-------|-------|------|------|---------|---------|---------|
|           | bitmap | bitmap | inode | inode | inode | data | data | data[0] | data[1] | data[1] |
| 2         |        |        | read  | read  |       | read |      |         |         |         |
| open(bar) |        |        |       | read  | read  |      | read |         |         |         |
|           |        |        |       |       | read  |      |      |         |         |         |
| read()    |        |        |       |       | read  |      |      | read    |         |         |
|           |        |        |       |       | write |      |      |         |         |         |
| read()    |        |        |       |       | read  |      |      |         |         |         |
|           |        |        |       |       | write |      |      |         | read    |         |
| read()    |        |        |       |       | read  |      |      |         |         |         |
|           |        |        |       |       | write |      |      |         |         | read    |



### **In-Memory (Non-Persistent) Structures**

- per process
  - descriptor table

\* which file descriptors does this process have open?

file descriptor is the pointer we use to read/write

- \* to which file does each open descriptor refer?
- \* what is the current file position for each descriptor? on reboot, these are cleared - this is fine because processes also disappear
- system wide
  - open file table
    - \* which files are currently open (by any process)?
  - i-node cache
    - \* in-memory copies of recently-used i-nodes
- delete file still open:
- still in memory, but no one else can open it
- once it's closed, then it is removed from memory

- block cache
  - \* in-memory copies of data blocks and indirect blocks for efficiency to not have to read from disk as much

caches are cleared on reboot - this is fine since it's stored elsewhere, will slow down for a bit to access them again

CS350 Operating Systems Fall 2015

if you turn off the computer by immediately disconnecting power, you might be in the middle of a file system operation, which can mess up your file data

42 File Systems

### Chaining

- VSFS uses a per-file index (direct and indirect pointers) to access blocks
- Two alternative approaches:
  - Chaining:



- \* Each block includes a pointer to the next block
- External chaining:
  - \* The chain is kept as an external structure
  - \* Microsoft's File Allocation Table (FAT) uses external chaining
    - a separate data structure stores pointers between data blocks

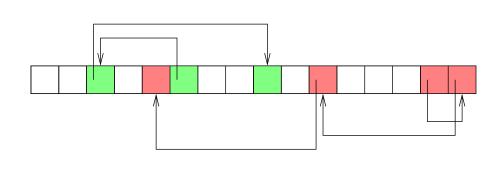
#### better because:

Fall 2015 CS350 Operating Systems

## Chaining

- Directory table contains the name of the file, and each file's starting block
- Acceptable for sequential access, very slow for random access (why?)





CS350 Operating Systems Fall 2015

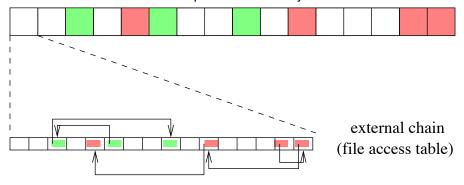
File Systems 44

## **External Chaining**

• Introduces a special file access table that specifies all of the file chains

chain can be stored in one or two data blocks - so we can load it all into \*memory\* which is faster to access

so it's faster to traverse a linked list in memory instead of going through blocks in persistent memory



### **Log-Structured File System (LFS)**

- LFS is built on a very different set of assumptions:
  - Memory sizes are growing
    - \* Most reads will be served from cache
    - \* Reads therefore are fast and do not require any seeks
  - Large gap between random I/O and sequential I/O performance
- Main idea: Make all writes sequential writes

problem: when we update the inode to point to a new data block, we have to seek all the way back to the front of the blocks

idea: make a new version of the inode right where we are on the disk after the new data block

CS350 Operating Systems Fall 2015

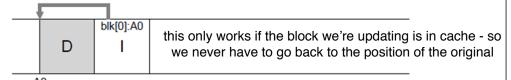
File Systems 46

### **Log-Structured File System**

• Write data block D at address A0



• Write i-node I for the file and point it to block D

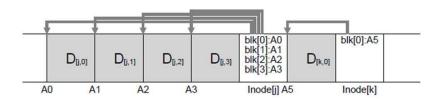


to read, we move backwards - this always gives us the latest version of the inode but the latest version of inodes should always be in cache

- The next block write will be placed after i-node I
- To overwrite data, just write a **new** block and i-node at the next position

## **Log-Structured File System**

- Write buffering is used to collect multiple writes to the same file together
  - Reduces the number of i-nodes that need to be written



problem: if we have infinitely large disk, this would be fine - but we'll run out to clear old versions, delete all inodes and data blocks that are not most recent (garbage collection)

and then move all the remaining blocks to the front (compaction/vacuuming)

CS350 Operating Systems Fall 2015

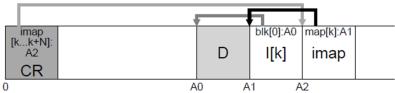
File Systems 48

### **Log-Structured File System**

- Problem:
  - How do we find the i-node for a file on a cache miss?



- i-node map (imap) takes an i-node number as an input, and returns the disk
   address of the most recent version of the i-node helps us rely not so heavily on cache
- Write a portion of the i-node map every time an i-node is updated a block's worth, including the block we just updated



- How do we find the i-node maps?
  - Maintain a periodically updated checkpoint region with pointers to the latest imaps
  - Checkpoint region always updated during graceful shutdown

### **Problems Caused by Failures**

- a single logical file system operation may require several disk I/O operations
- example: deleting a file
  - remove entry from directory
  - remove file index (i-node) from i-node table
  - mark file's data blocks free in free space index
- what if, because of a failure, some but not all of these changes are reflected on the disk?
  - system failure will destroy in-memory file system structures
  - persistent structures should be *crash consistent*, i.e., should be consistent when system restarts after a failure

CS350 Operating Systems Fall 2015

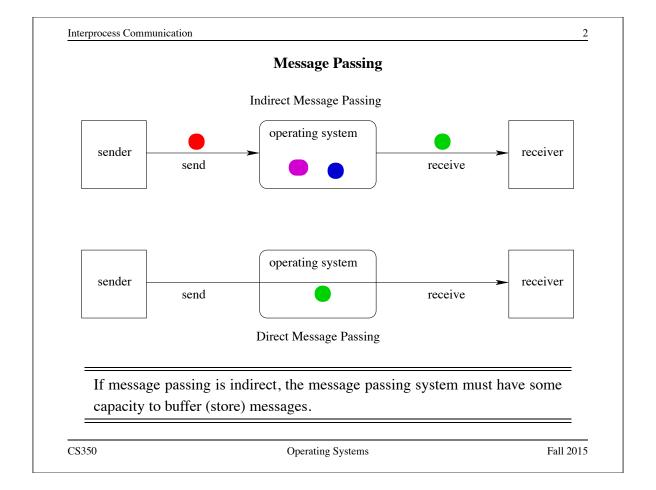
File Systems 50

#### **Fault Tolerance**

- special-purpose consistency checkers (e.g., Unix fsck in Berkeley FFS, Linux ext2)
  - runs after a crash, before normal operations resume
  - find and attempt to repair inconsistent file system data structures, e.g.:
    - \* file with no directory entry
    - \* free space that is not marked as free
- journaling (e.g., Veritas, NTFS, Linux ext3)
  - record file system meta-data changes in a journal (log), so that sequences of changes can be written to disk in a single operation
  - after changes have been journaled, update the disk data structures (write-ahead logging)
  - after a failure, redo journaled updates in case they were not done before the failure

# **Interprocess Communication Mechanisms**

- shared storage
  - shared virtual memory
  - shared files
- · message-based
  - sockets
  - pipes
  - signals
  - **–** ...



### **Properties of Message Passing Mechanisms**

## **Directionality:**

- simplex (one-way), duplex (two-way)
- half-duplex (two-way, but only one way at a time)

### **Message Boundaries:**

datagram model: message boundaries

stream model: no boundaries

**Connections:** need to connect before communicating?

- in connection-oriented models, recipient is specified at time of connection, not by individual send operations. All messages sent over a connection have the same recipient.
- in connectionless models, recipient is specified as a parameter to each send operation.

## **Reliability:**

• can messages get lost? reordered? damaged?

CS350 Operating Systems Fall 2015

Interprocess Communication

4

#### **Sockets**

- a socket is a communication *end-point*
- if two processes are to communicate, each process must create its own socket
- two common types of sockets

**stream sockets:** support connection-oriented, reliable, duplex communication under the stream model (no message boundaries)

**datagram sockets:** support connectionless, best-effort (unreliable), duplex communication under the datagram model (message boundaries)

• both types of sockets also support a variety of address domains, e.g.,

**Unix domain:** useful for communication between processes running on the same machine

**INET domain:** useful for communication between process running on different machines that can communicate using IP protocols.

### **Using Datagram Sockets (Receiver)**

```
s = socket(addressType, SOCK_DGRAM);
bind(s,address);
recvfrom(s,buf,bufLength,sourceAddress);
...
close(s);
```

- socket creates a socket
- bind assigns an address to the socket
- recvfrom receives a message from the socket
  - buf is a buffer to hold the incoming message
  - sourceAddress is a buffer to hold the address of the message sender
- both buf and sourceAddress are filled by the recvfrom call

CS350 Operating Systems Fall 2015

Interprocess Communication

6

### **Using Datagram Sockets (Sender)**

```
s = socket(addressType, SOCK_DGRAM);
sendto(s,buf,msgLength,targetAddress)
...
close(s);
```

- socket creates a socket
- sendto sends a message using the socket
  - buf is a buffer that contains the message to be sent
  - msgLength indicates the length of the message in the buffer
  - targetAddress is the address of the socket to which the message is to be delivered

### **More on Datagram Sockets**

- sendto and recvfrom calls may block
  - recvfrom blocks if there are no messages to be received from the specified socket
  - sendto blocks if the system has no more room to buffer undelivered messages
- datagram socket communications are (in general) unreliable
  - messages (datagrams) may be lost
  - messages may be reordered
- The sending process must know the address of the receive process's socket.

CS350 Operating Systems Fall 2015

Interprocess Communication

8

#### **Using Stream Sockets (Passive Process)**

```
s = socket(addressType, SOCK_STREAM);
bind(s,address);
listen(s,backlog);
ns = accept(s,sourceAddress);
recv(ns,buf,bufLength);
send(ns,buf,bufLength);
...
close(ns); // close accepted connection
close(s); // don't accept more connections
```

- listen specifies the number of connection requests for this socket that will be queued by the kernel
- accept accepts a connection request and creates a new socket (ns)
- recv receives up to bufLength bytes of data from the connection
- send sends bufLength bytes of data over the connection.

### **Notes on Using Stream Sockets (Passive Process)**

- accept creates a new socket (ns) for the new connection
- sourceAddress is an address buffer. accept fills it with the address of the socket that has made the connection request
- additional connection requests can be accepted using more accept calls on the original socket (s)
- accept blocks if there are no pending connection requests
- connection is duplex (both send and recv can be used)

CS350 Operating Systems Fall 2015

Interprocess Communication

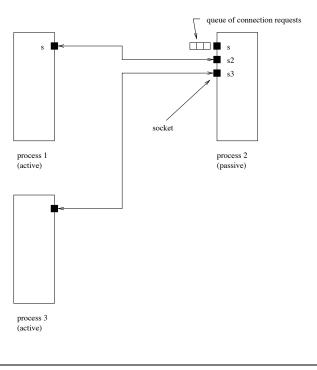
10

#### **Using Stream Sockets (Active Process)**

```
s = socket(addressType, SOCK_STREAM);
connect(s,targetAddress);
send(s,buf,bufLength);
recv(s,buf,bufLength);
...
close(s);
```

- connect sends a connection request to the socket with the specified address
  - connect blocks until the connection request has been accepted
- active process may (optionally) bind an address to the socket (using bind)
   before connecting. This is the address that will be returned by the accept call in the passive process
- if the active process does not choose an address, the system will choose one

### **Illustration of Stream Socket Connections**



CS350 Operating Systems Fall 2015

Interprocess Communication

12

### **Pipes**

- pipes are communication objects (not end-points)
- pipes use the stream model and are connection-oriented and reliable
- some pipes are simplex, some are duplex
- pipes use an implicit addressing mechanism that limits their use to communication between *related* processes, typically a child process and its parent
- a pipe() system call creates a pipe and returns two descriptors, one for each end of the pipe
  - for a simplex pipe, one descriptor is for reading, the other is for writing
  - for a duplex pipe, both descriptors can be used for reading and writing

## One-way Child/Parent Communication Using a Simplex Pipe

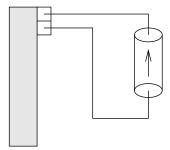
```
int fd[2];
char m[] = "message for parent";
char y[100];
pipe(fd); // create pipe
pid = fork(); // create child process
if (pid == 0) {
    // child executes this
    close(fd[0]); // close read end of pipe
    write(fd[1],m,19);
    ...
} else {
    // parent executes this
    close(fd[1]); // close write end of pipe
    read(fd[0],y,19);
    ...
}
```

CS350 Operating Systems Fall 2015

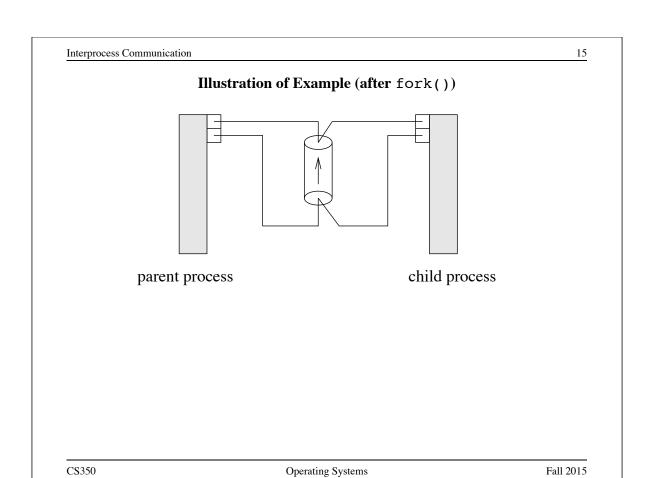
Interprocess Communication

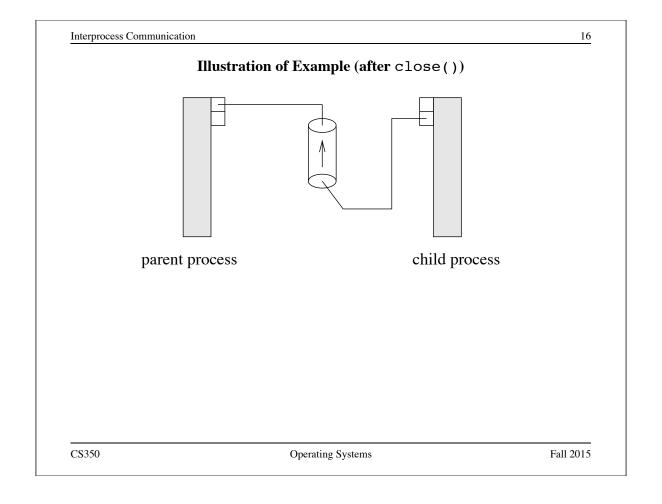
14

## Illustration of Example (after pipe())



parent process





### **Examples of Other Interprocess Communication Mechanisms**

### named pipe:

- similar to pipes, but with an associated name (usually a file name)
- name allows arbitrary processes to communicate by opening the same named pipe
- must be explicitly deleted, unlike an unnamed pipe

### message queue:

- like a named pipe, except that there are message boundaries
- msgsend call sends a message into the queue, msgrecv call receives the next message from the queue

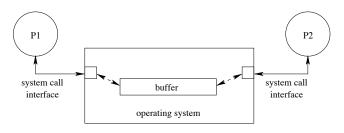
CS350 Operating Systems Fall 2015

Interprocess Communication

18

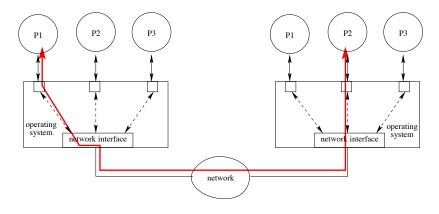
### **Implementing IPC**

- application processes use descriptors (identifiers) provided by the kernel to refer to specific sockets and pipes, as well as files and other objects
- kernel *descriptor tables* (or other similar mechanism) are used to associate descriptors with kernel data structures that implement IPC objects
- kernel provides bounded buffer space for data that has been sent using an IPC mechanism, but that has not yet been received
  - for IPC objects, like pipes, buffering is usually on a per object basis
  - IPC end points, like sockets, buffering is associated with each endpoint



### **Network Interprocess Communication**

- some sockets can be used to connect processes that are running on different machines
- the kernel:
  - controls access to network interfaces
  - multiplexes socket connections across the network



Operating Systems

Interprocess Communication

CS350

20

Fall 2015

### **Signals**

- signals permit asynchronous one-way communication
  - from a process to another process, or to a group of processes, via the kernel
  - from the kernel to a process, or to a group of processes
- there are many types of signals
- the arrival of a signal may cause the execution of a *signal handler* in the receiving process
- there may be a different handler for each type of signal

## **Examples of Signal Types**

| Signal  | Value    | Action | Comment                     |
|---------|----------|--------|-----------------------------|
|         |          |        |                             |
| SIGINT  | 2        | Term   | Interrupt from keyboard     |
| SIGILL  | 4        | Core   | Illegal Instruction         |
| SIGKILL | 9        | Term   | Kill signal                 |
| SIGCHLD | 20,17,18 | Ign    | Child stopped or terminated |
| SIGBUS  | 10,7,10  | Core   | Bus error                   |
| SIGXCPU | 24,24,30 | Core   | CPU time limit exceeded     |
| SIGSTOP | 17,19,23 | Stop   | Stop process                |

CS350 Operating Systems Fall 2015

Interprocess Communication

22

## **Signal Handling**

- operating system determines default signal handling for each new process
- example default actions:
  - ignore (do nothing)
  - kill (terminate the process)
  - stop (block the process)
- a running process can change the default for some types of signals
- signal-related system calls
  - calls to set non-default signal handlers, e.g., Unix signal, sigaction
  - calls to send signals, e.g., Unix kill